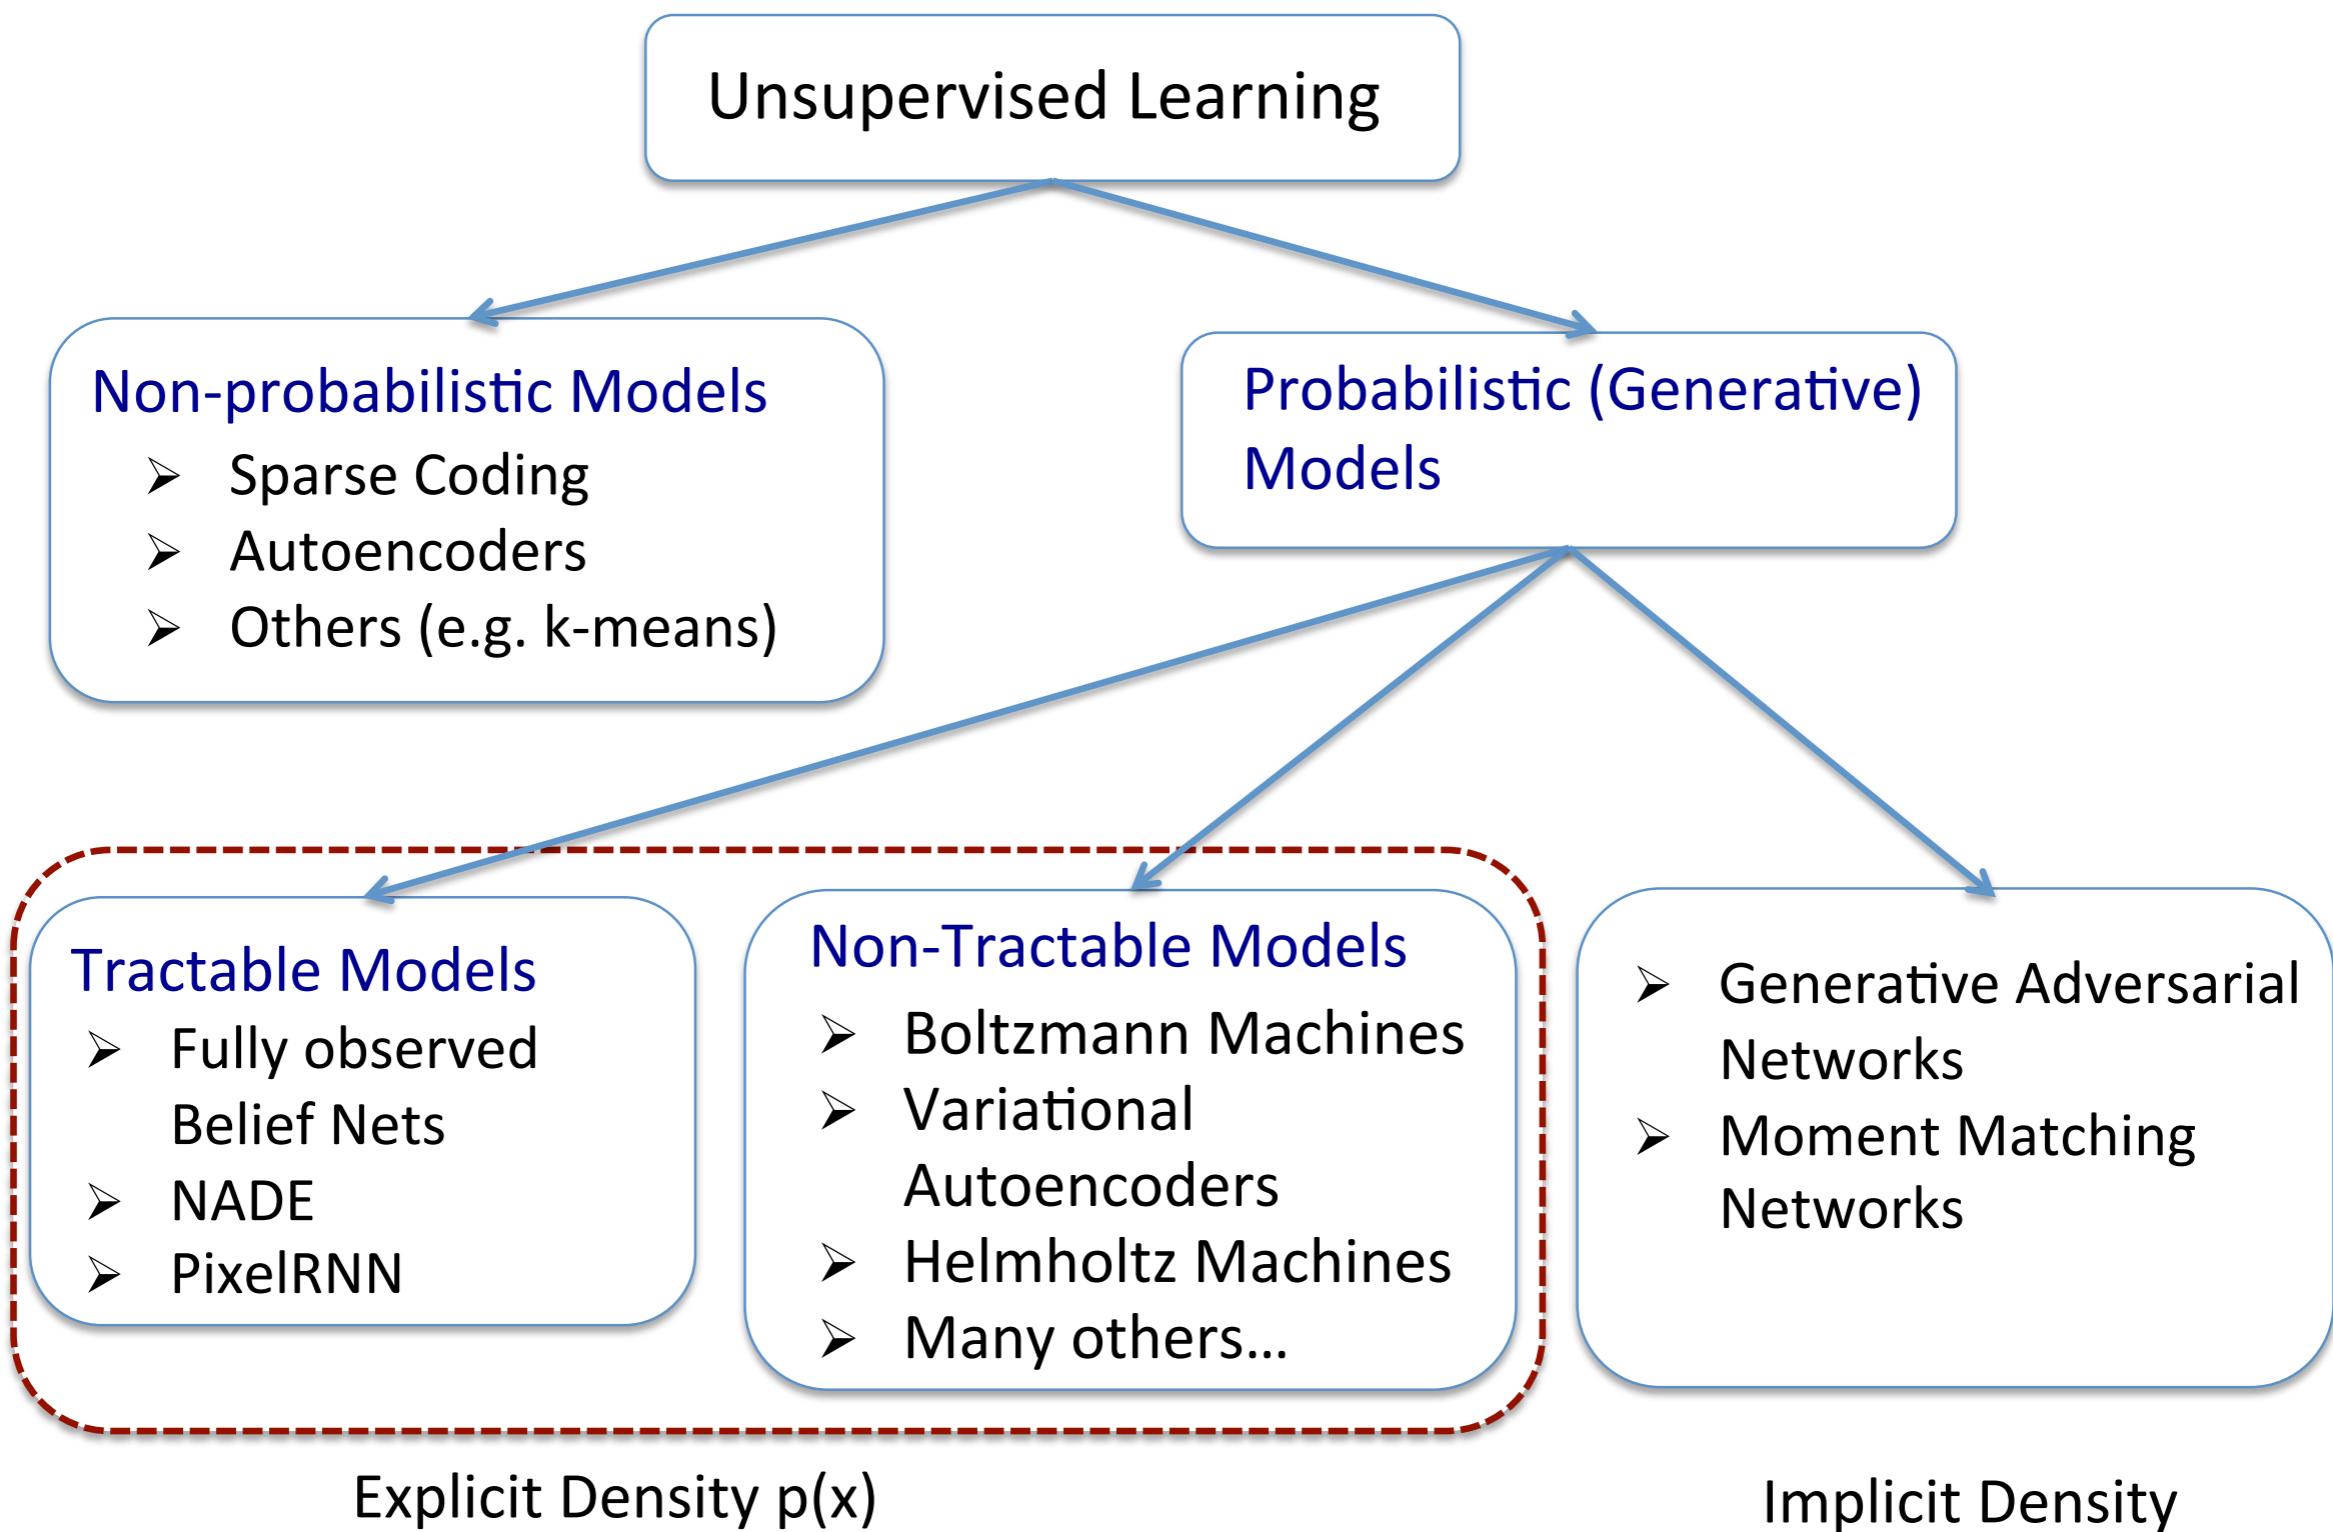


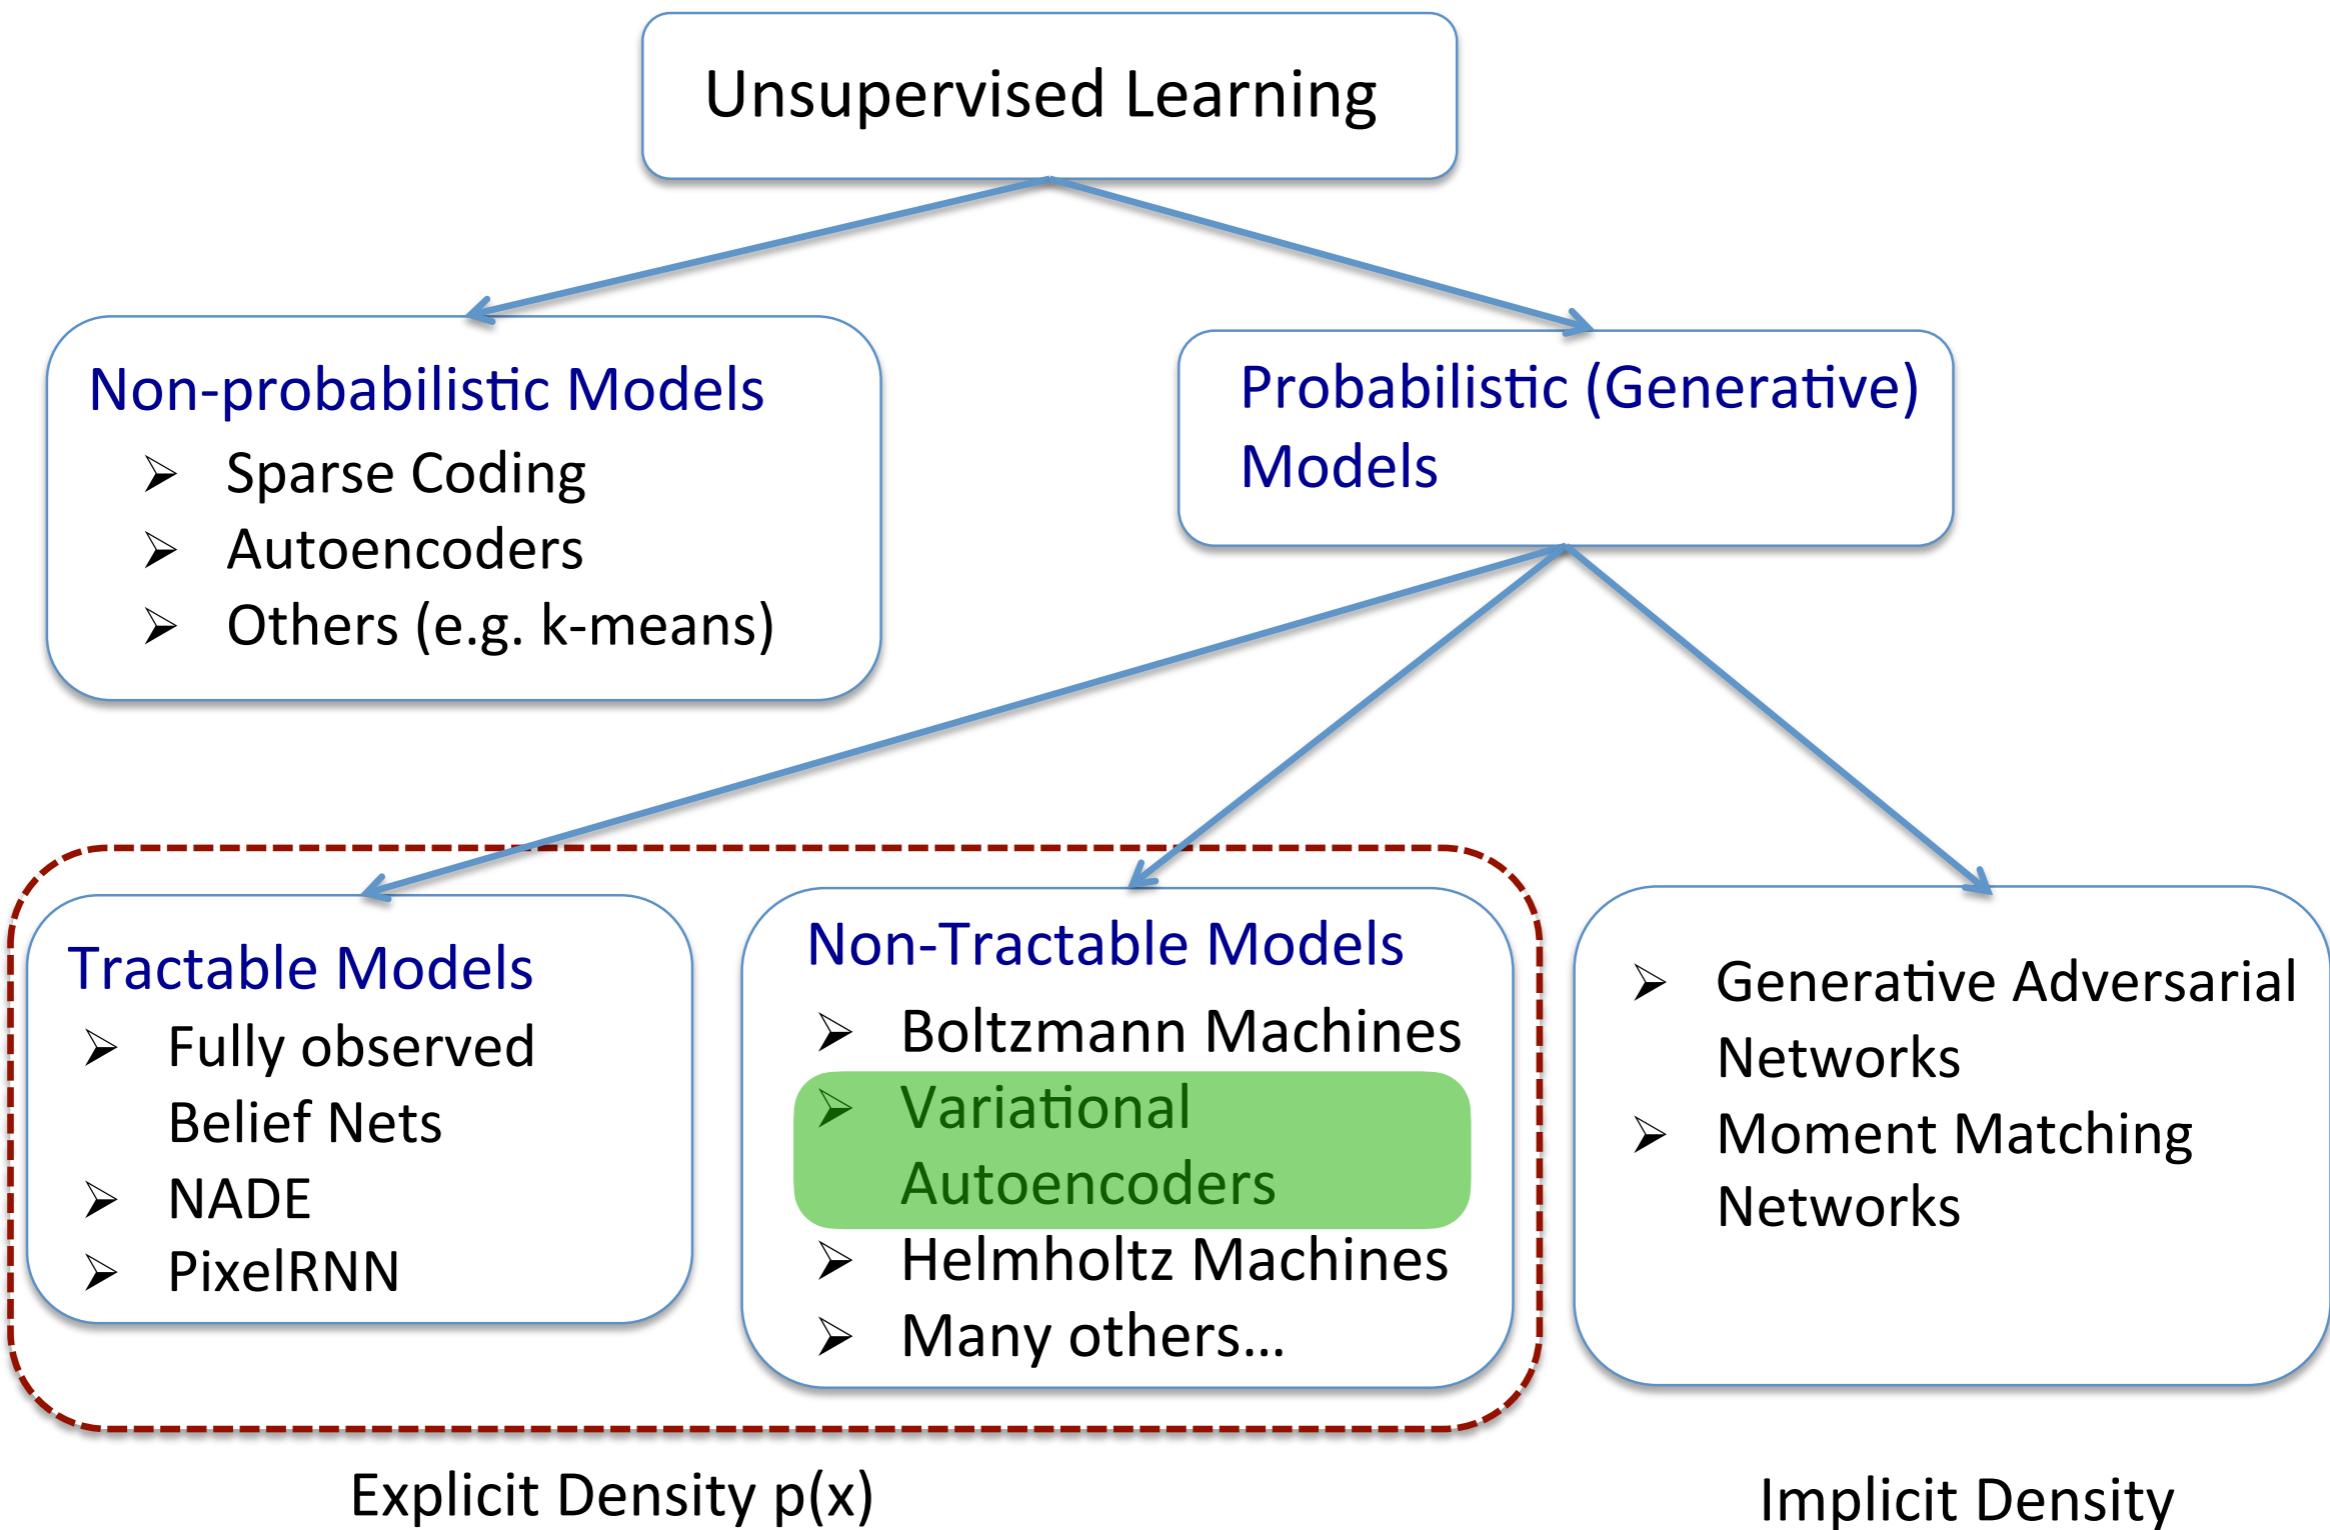
Deep Learning with Neural Networks

Introduction to Generative Adversarial Networks

Alejandro Lancho Serrano, alancho@ing.uc3m.es

Material original por: Pablo Martínez Olmos



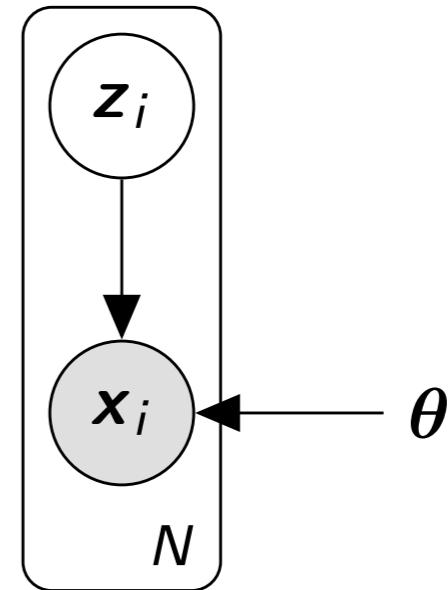


Variational Autoencoder (Kingma & Welling 2014)

Consider a set of i.i.d observations $\mathbf{x}^{(i)} \in \mathbb{R}^d$, $i = 1, \dots, N$. Let $\mathcal{D} = \{\mathbf{x}^{(i)}\}_{i=1}^N$.

Generative Model using a Latent Space

$$p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{x}|z)p(z)dz$$



- $z \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_k)$ ($k \leq d$) (low-dimensional embedding)
- $x|z \sim \mathcal{N}(\mu_{\theta}(z), \text{diag}(\sigma_{\theta}(z)))$
- $\mu_{\theta}(z)$ and $\log(\sigma_{\theta}(z))$ are the outputs of a NN $\mathbb{R}^k \rightarrow \mathbb{R}^d$ with parameter vector θ (Decoding network).

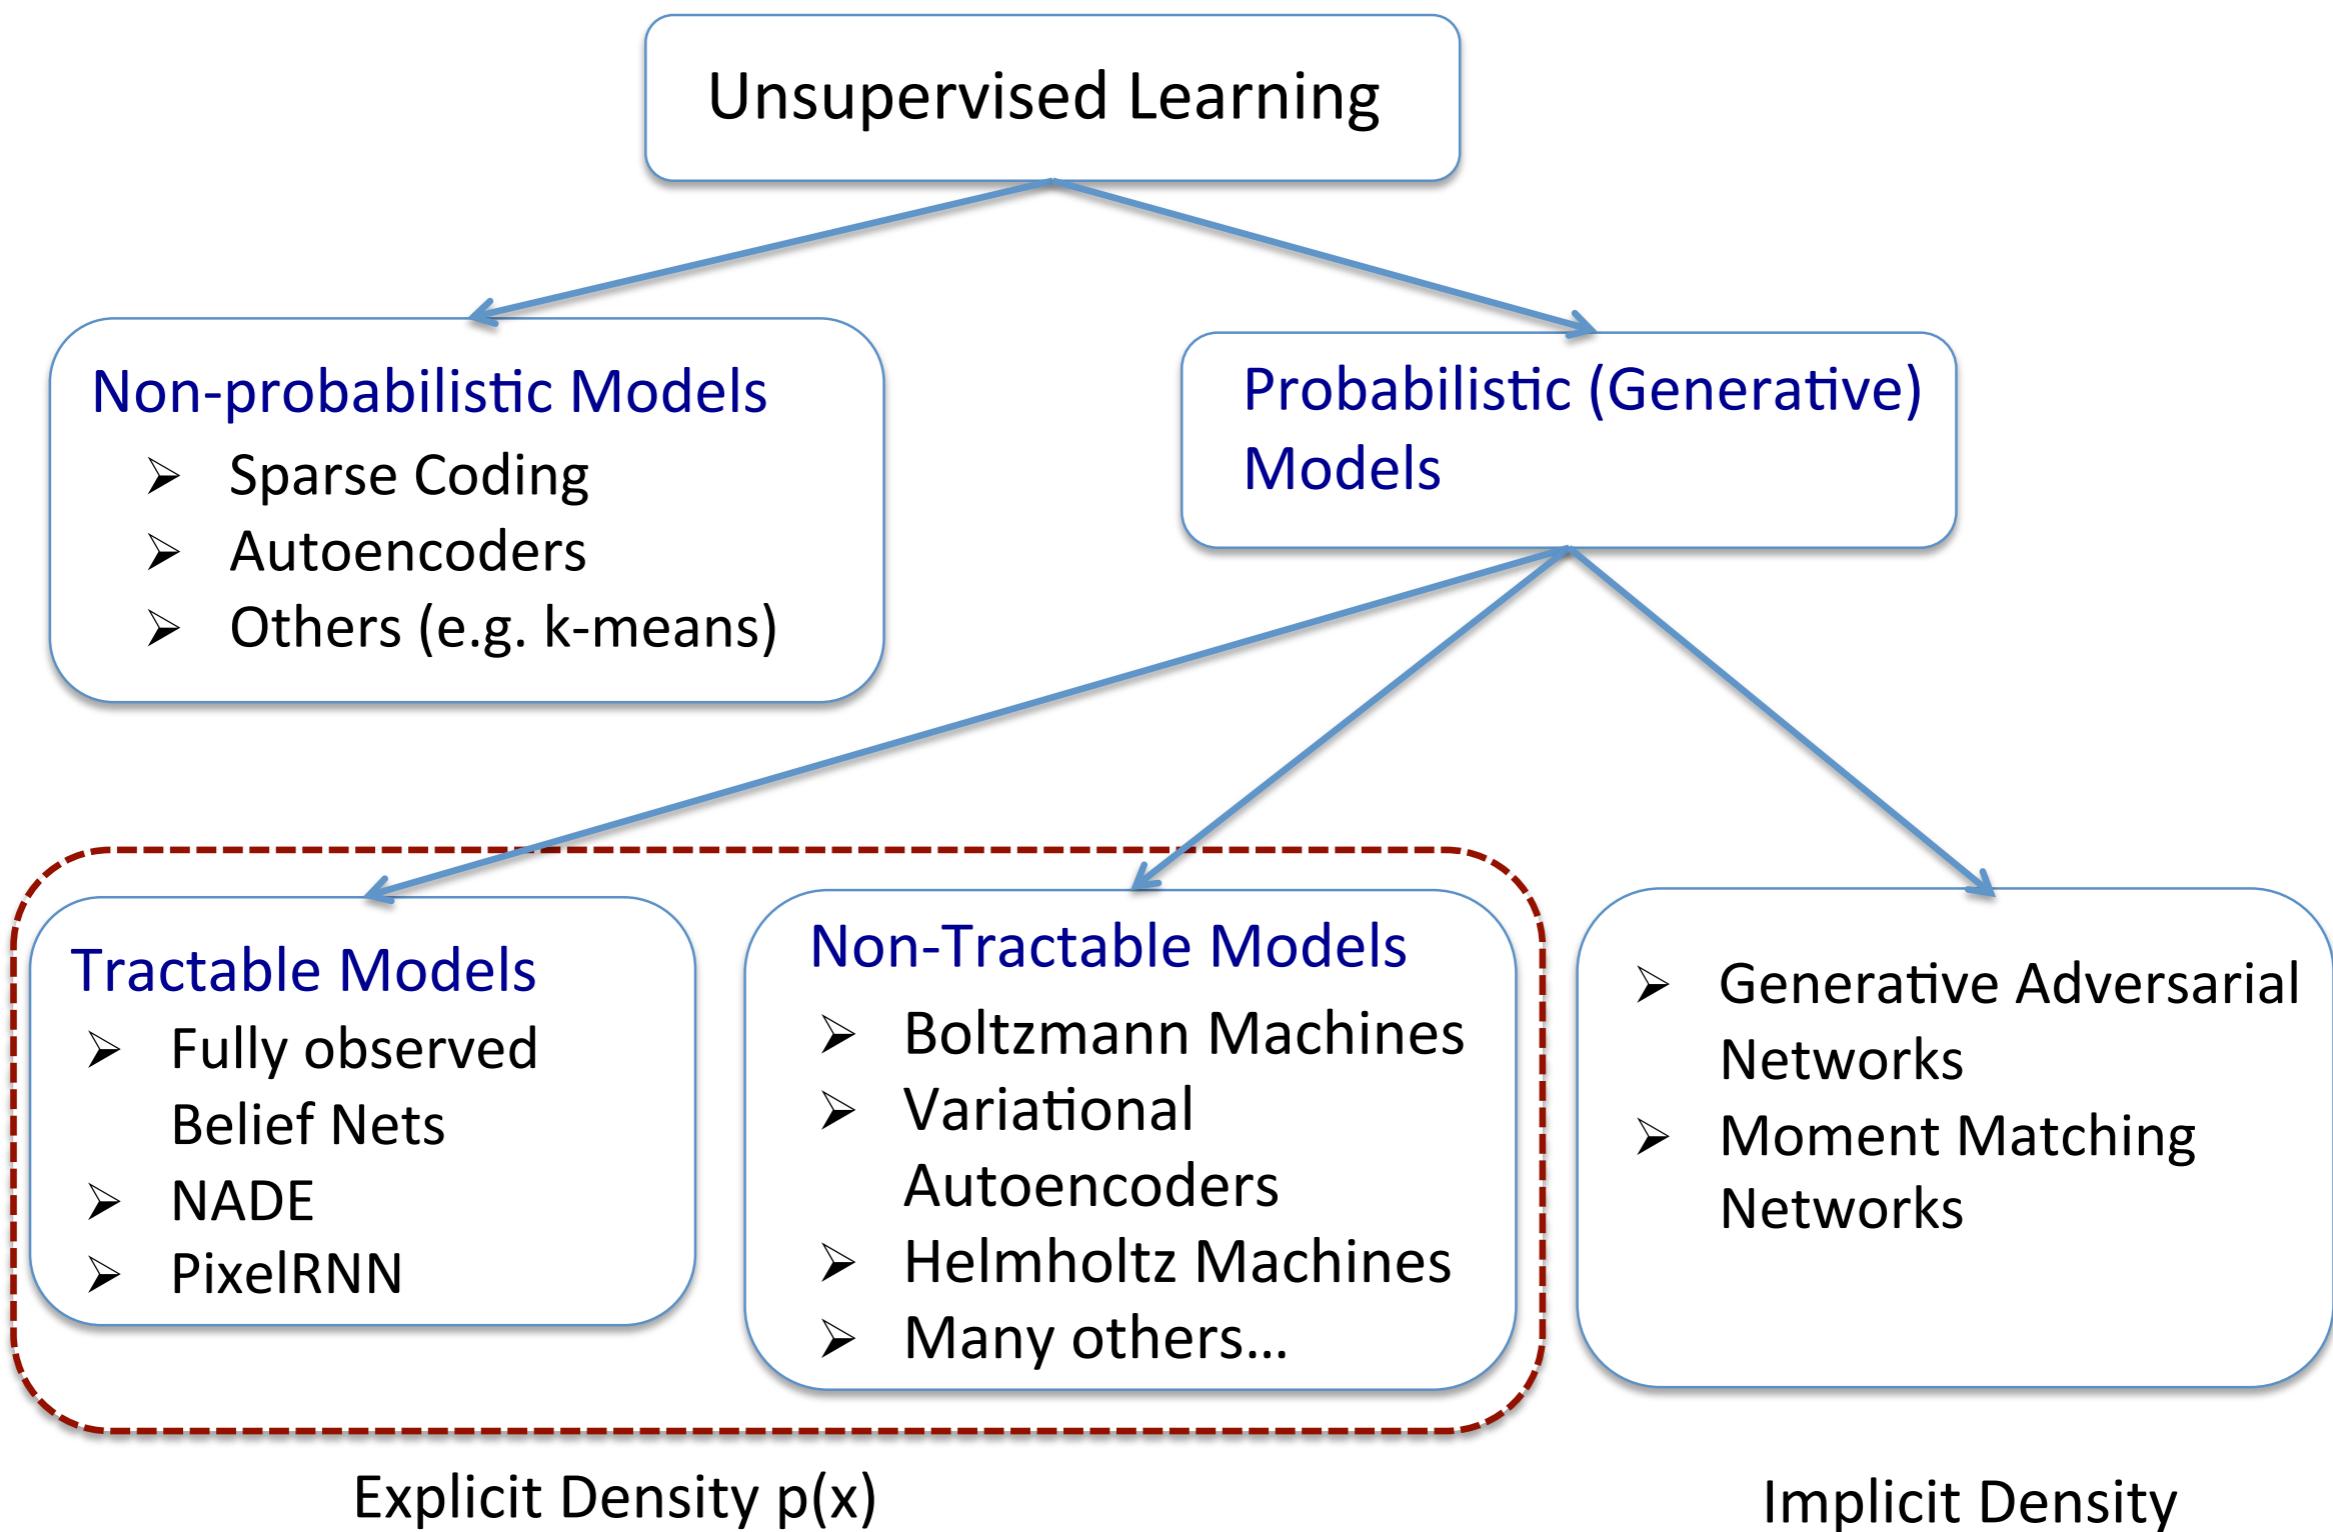
Variational Autoencoder Pros and Cons summary

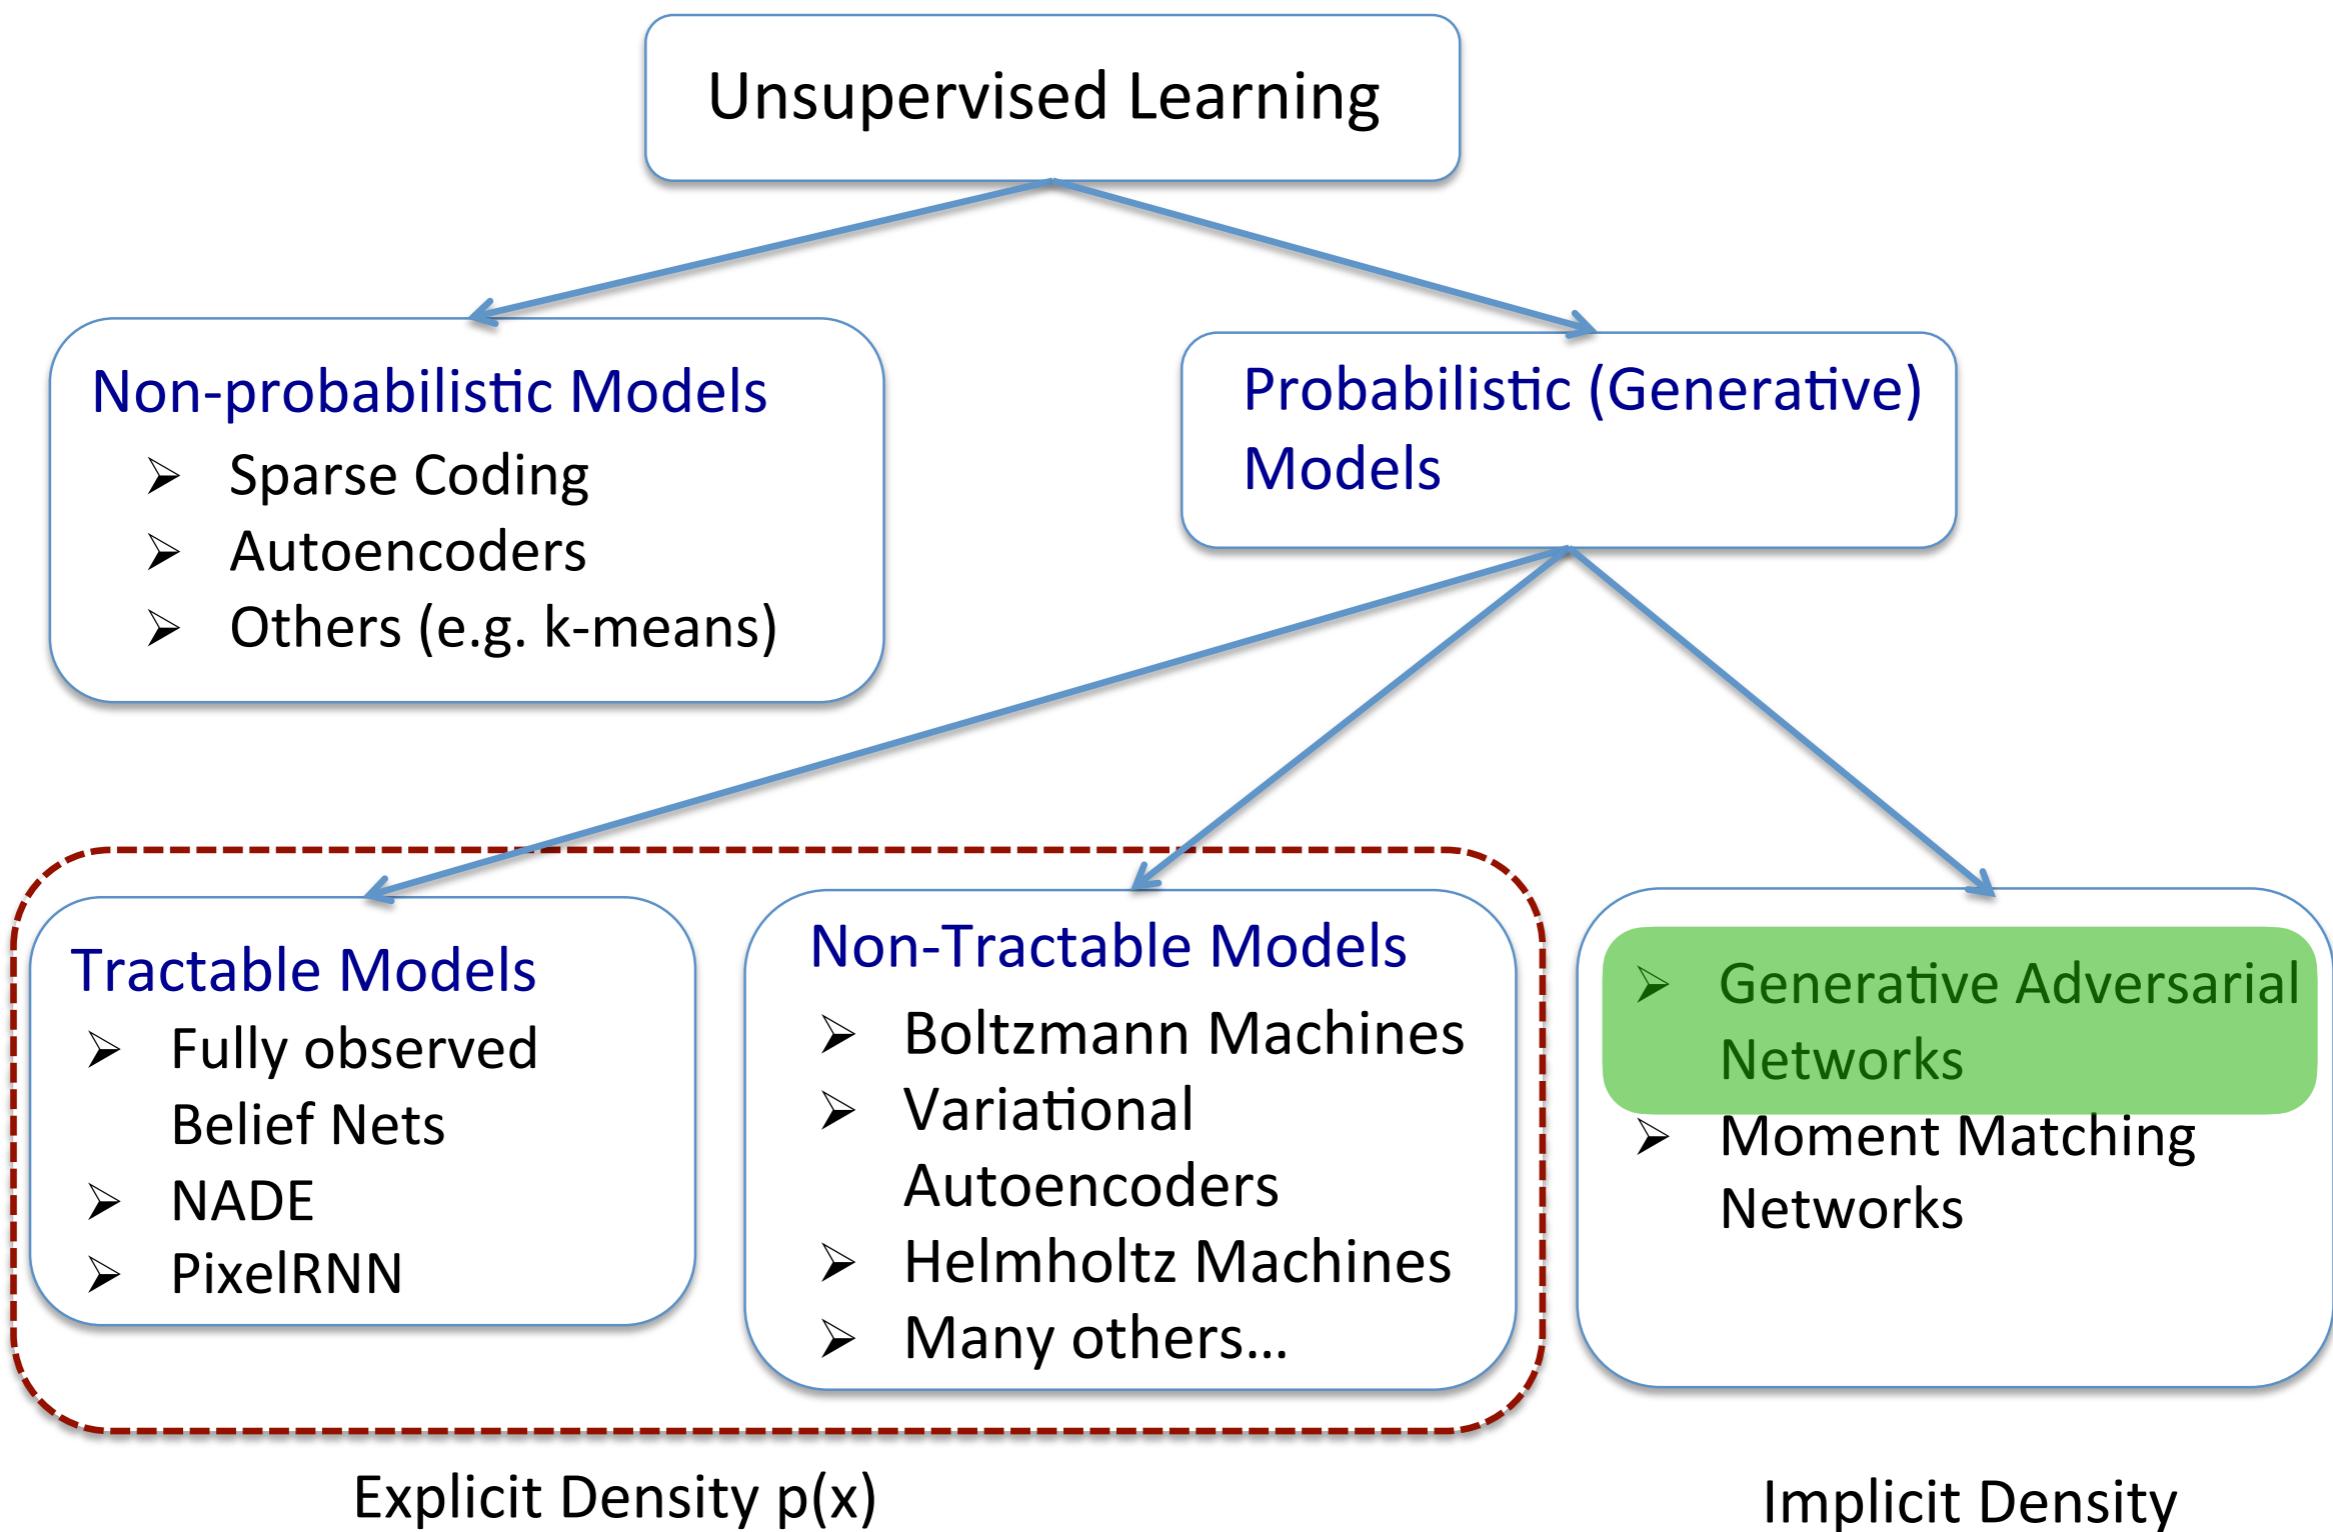
Pros:

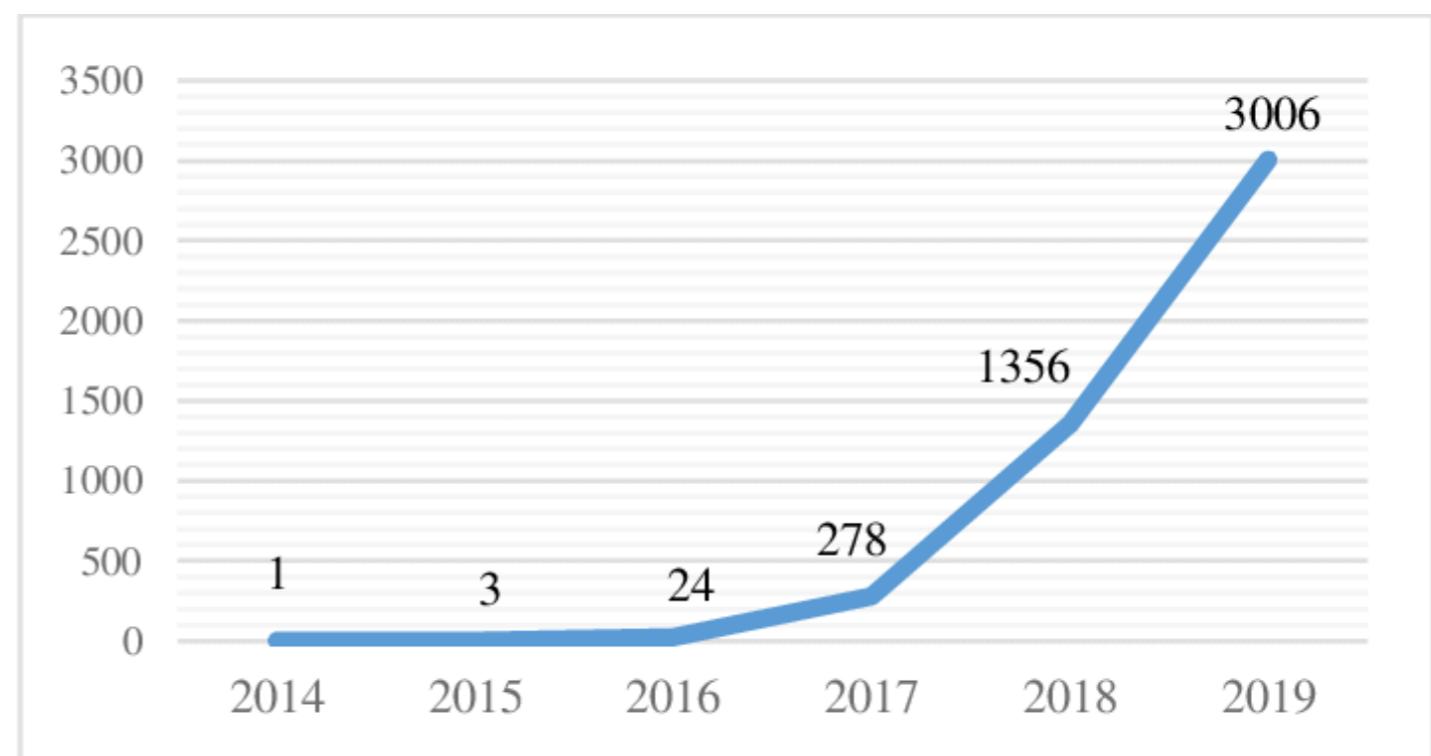
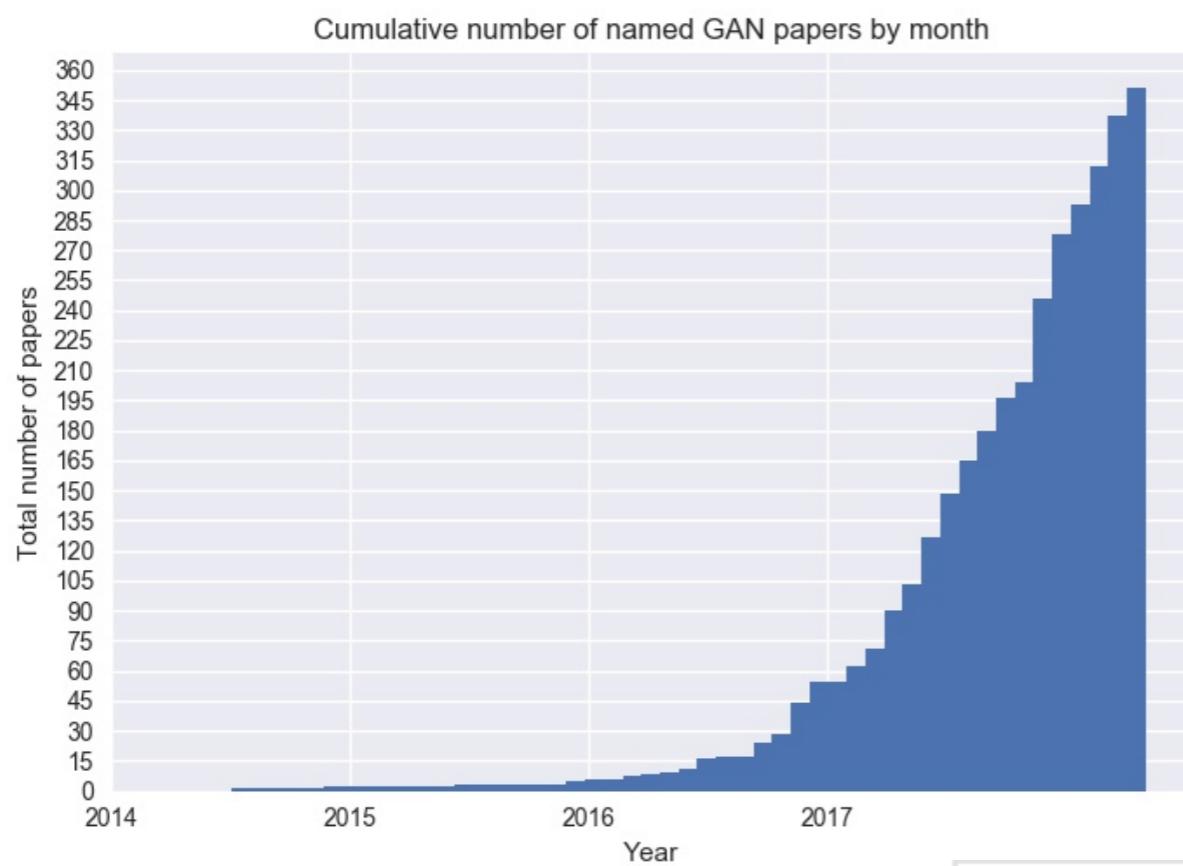
- **Stability** and **ease of training**
- **Probabilistic interpretation** through the probabilistic latent space
- Easy to get **diverse samples** and avoid mode collapse
- Good for **semi-supervised** and structured learning tasks
 - How? conditioning latent space
- Suitable for applications where **uncertainty quantification** is needed
 - Anomaly detection, decision making, etc.

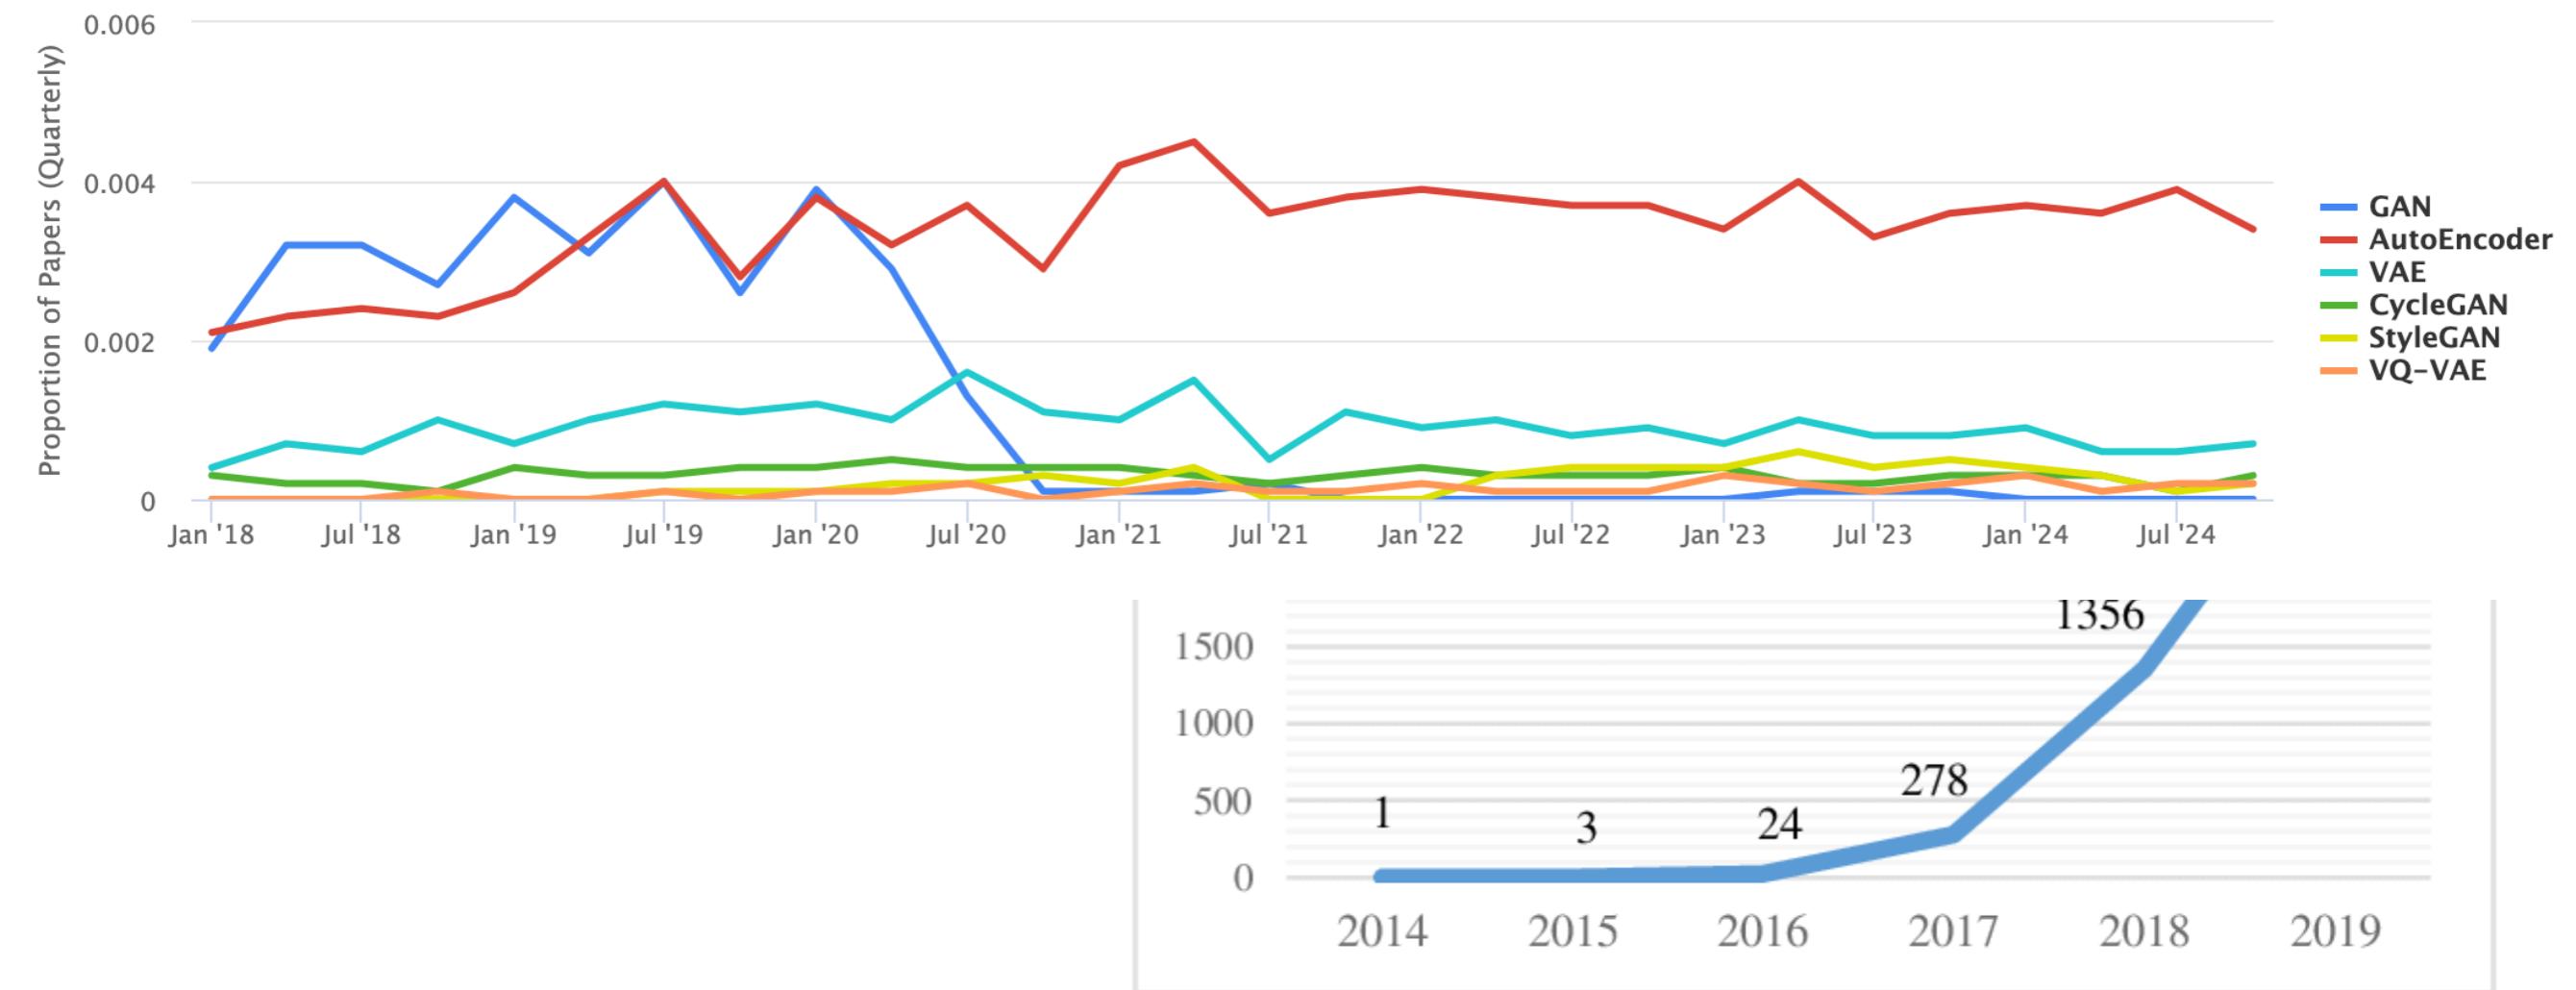
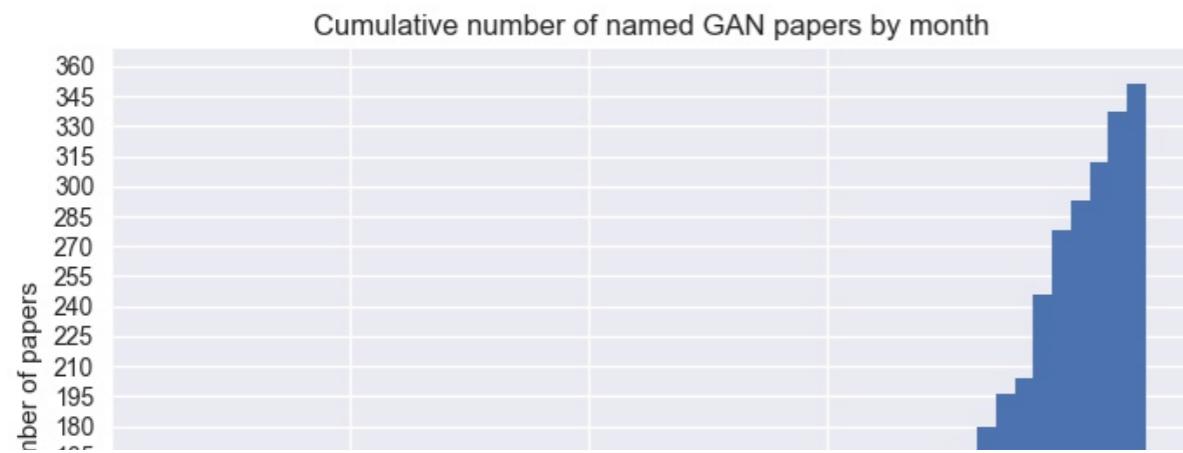
Cons:

- Prone to give **blurry samples** due to reconstruction error terms based on MSE
- **Limited in capturing complex distributions**
 - KL term pulls the latent space towards a Gaussian distributions
 - ▶ Limits generation quality in some cases

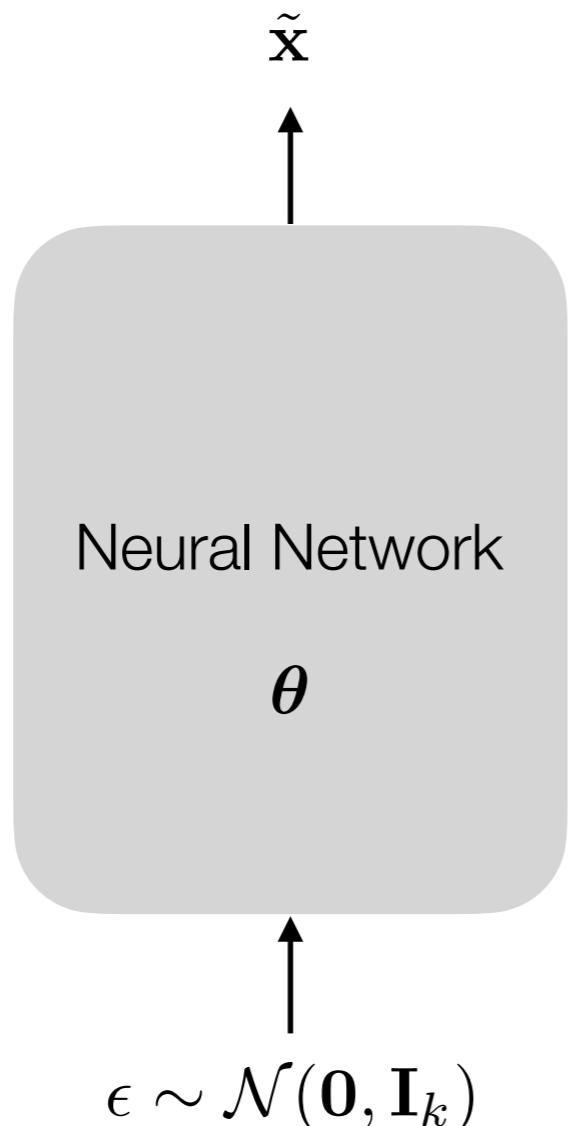






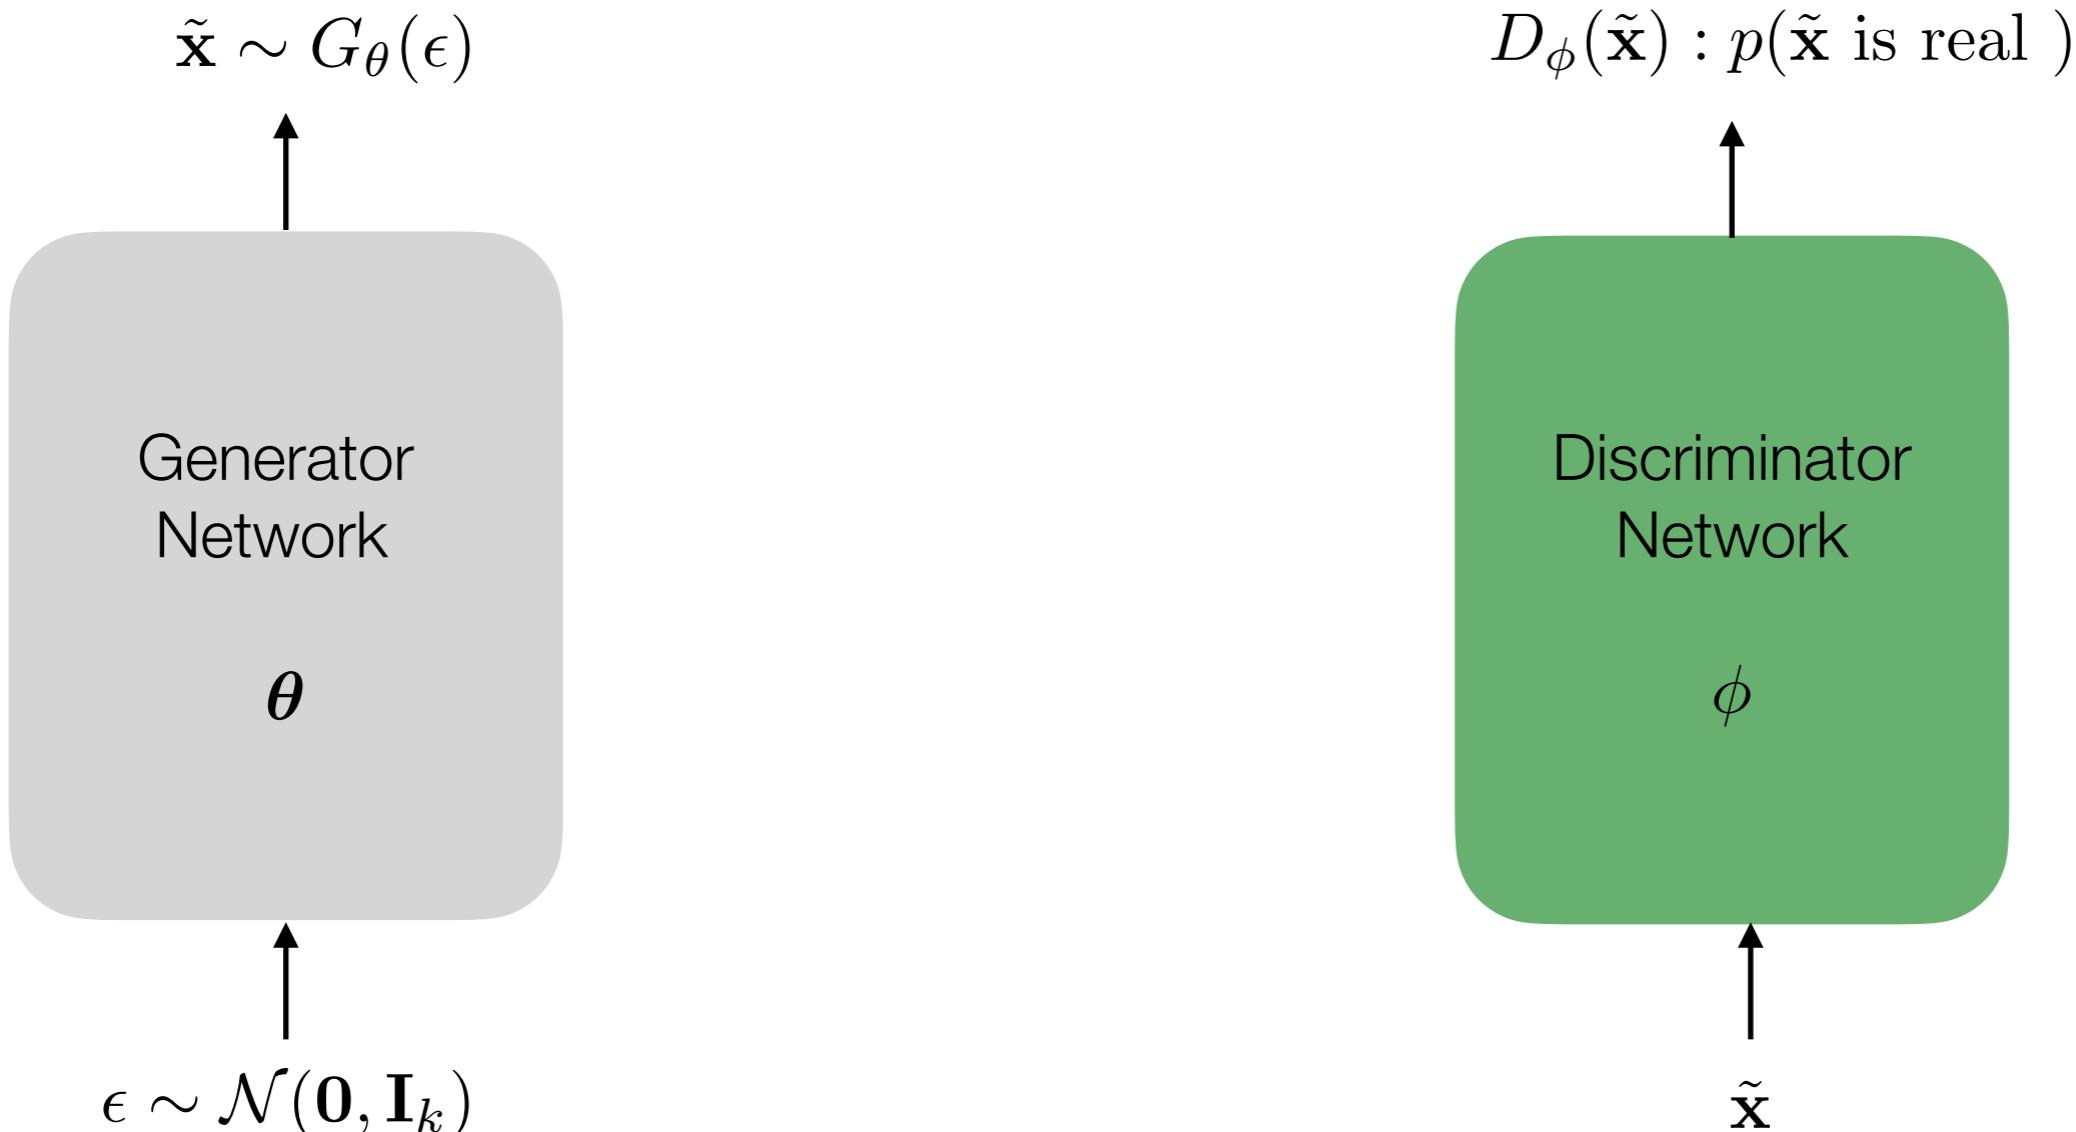


Implicit density estimation

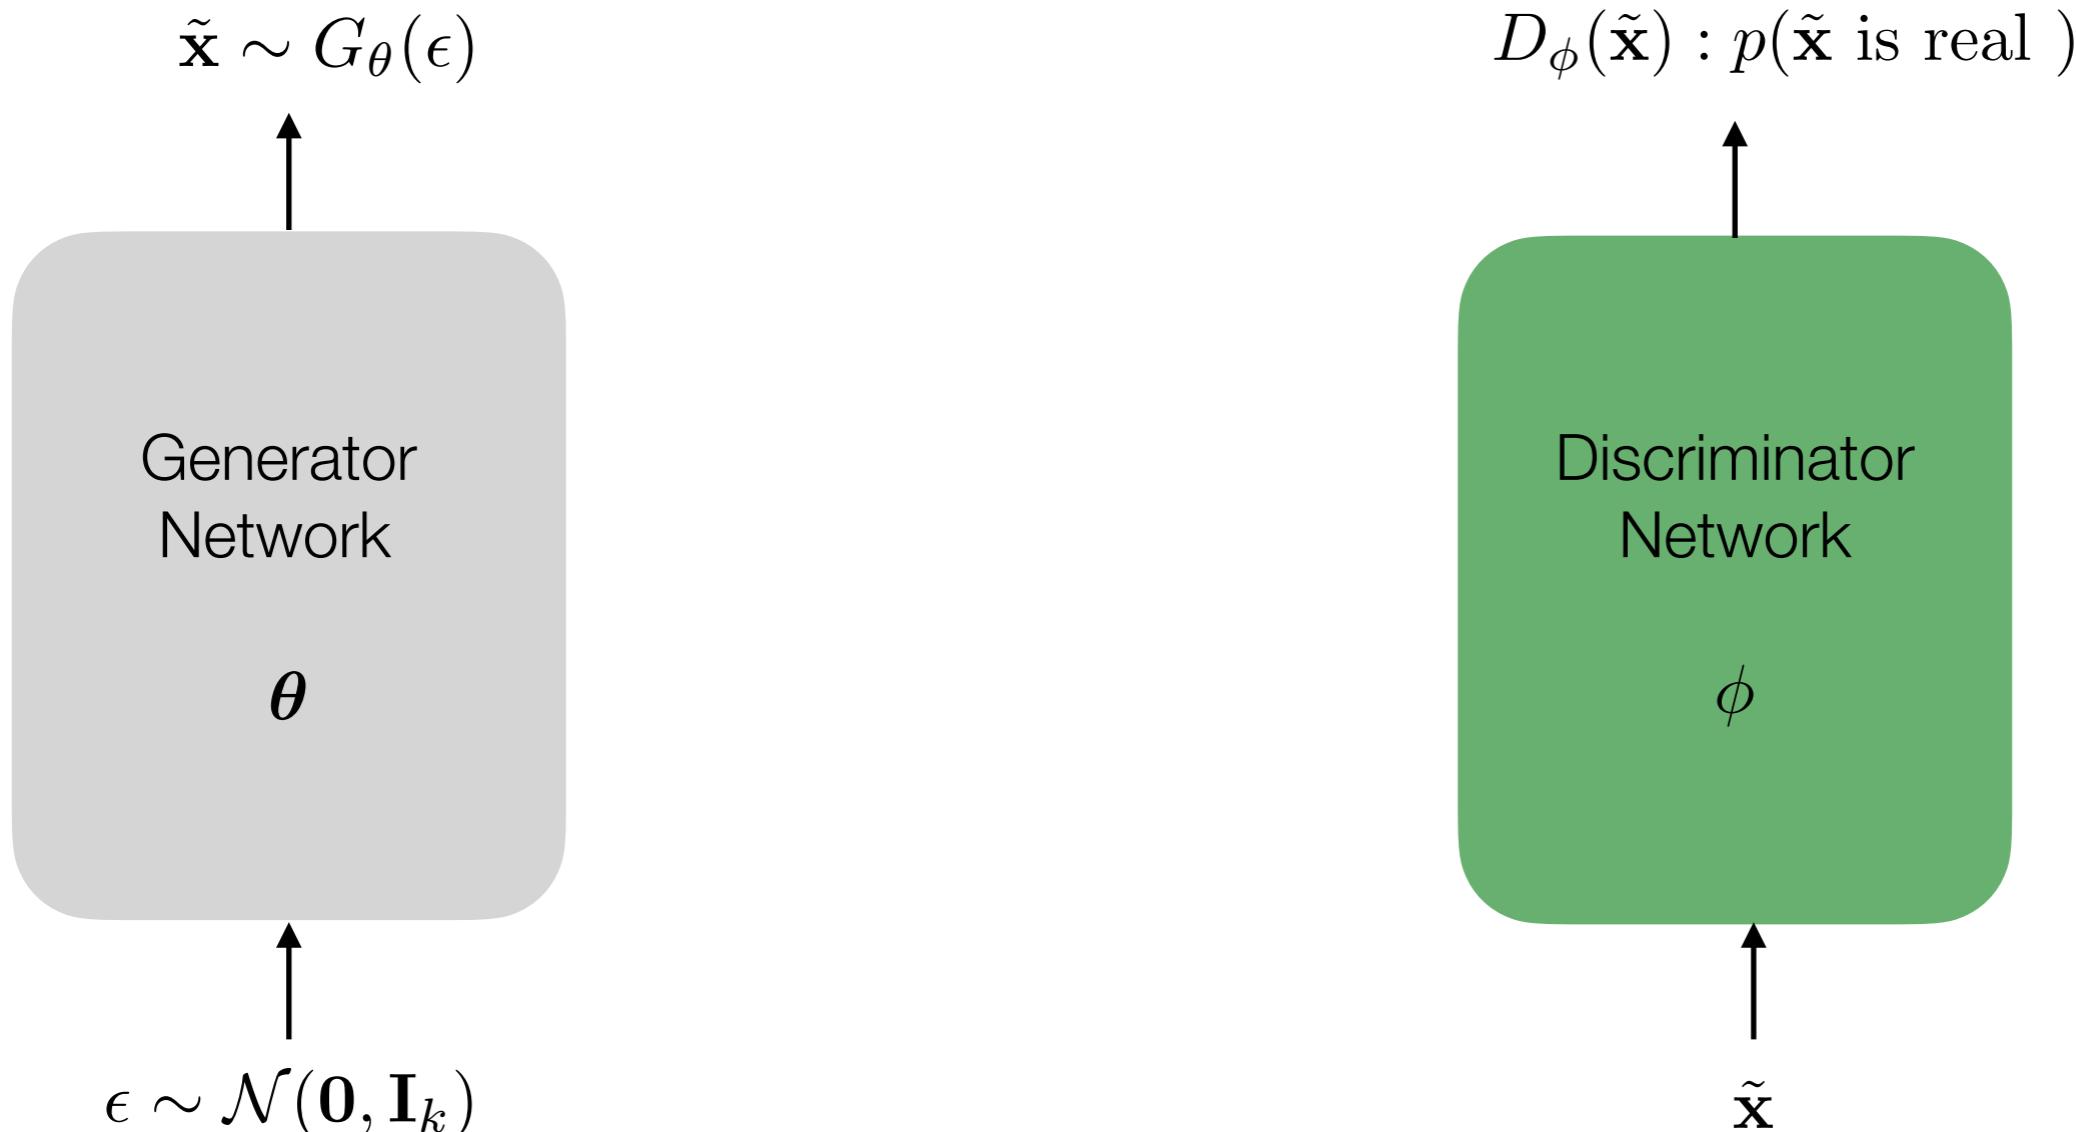


- **Unknown density** at the output
- **Sampling** is easy!

Generative Adversarial Network (Goodfellow 2014)



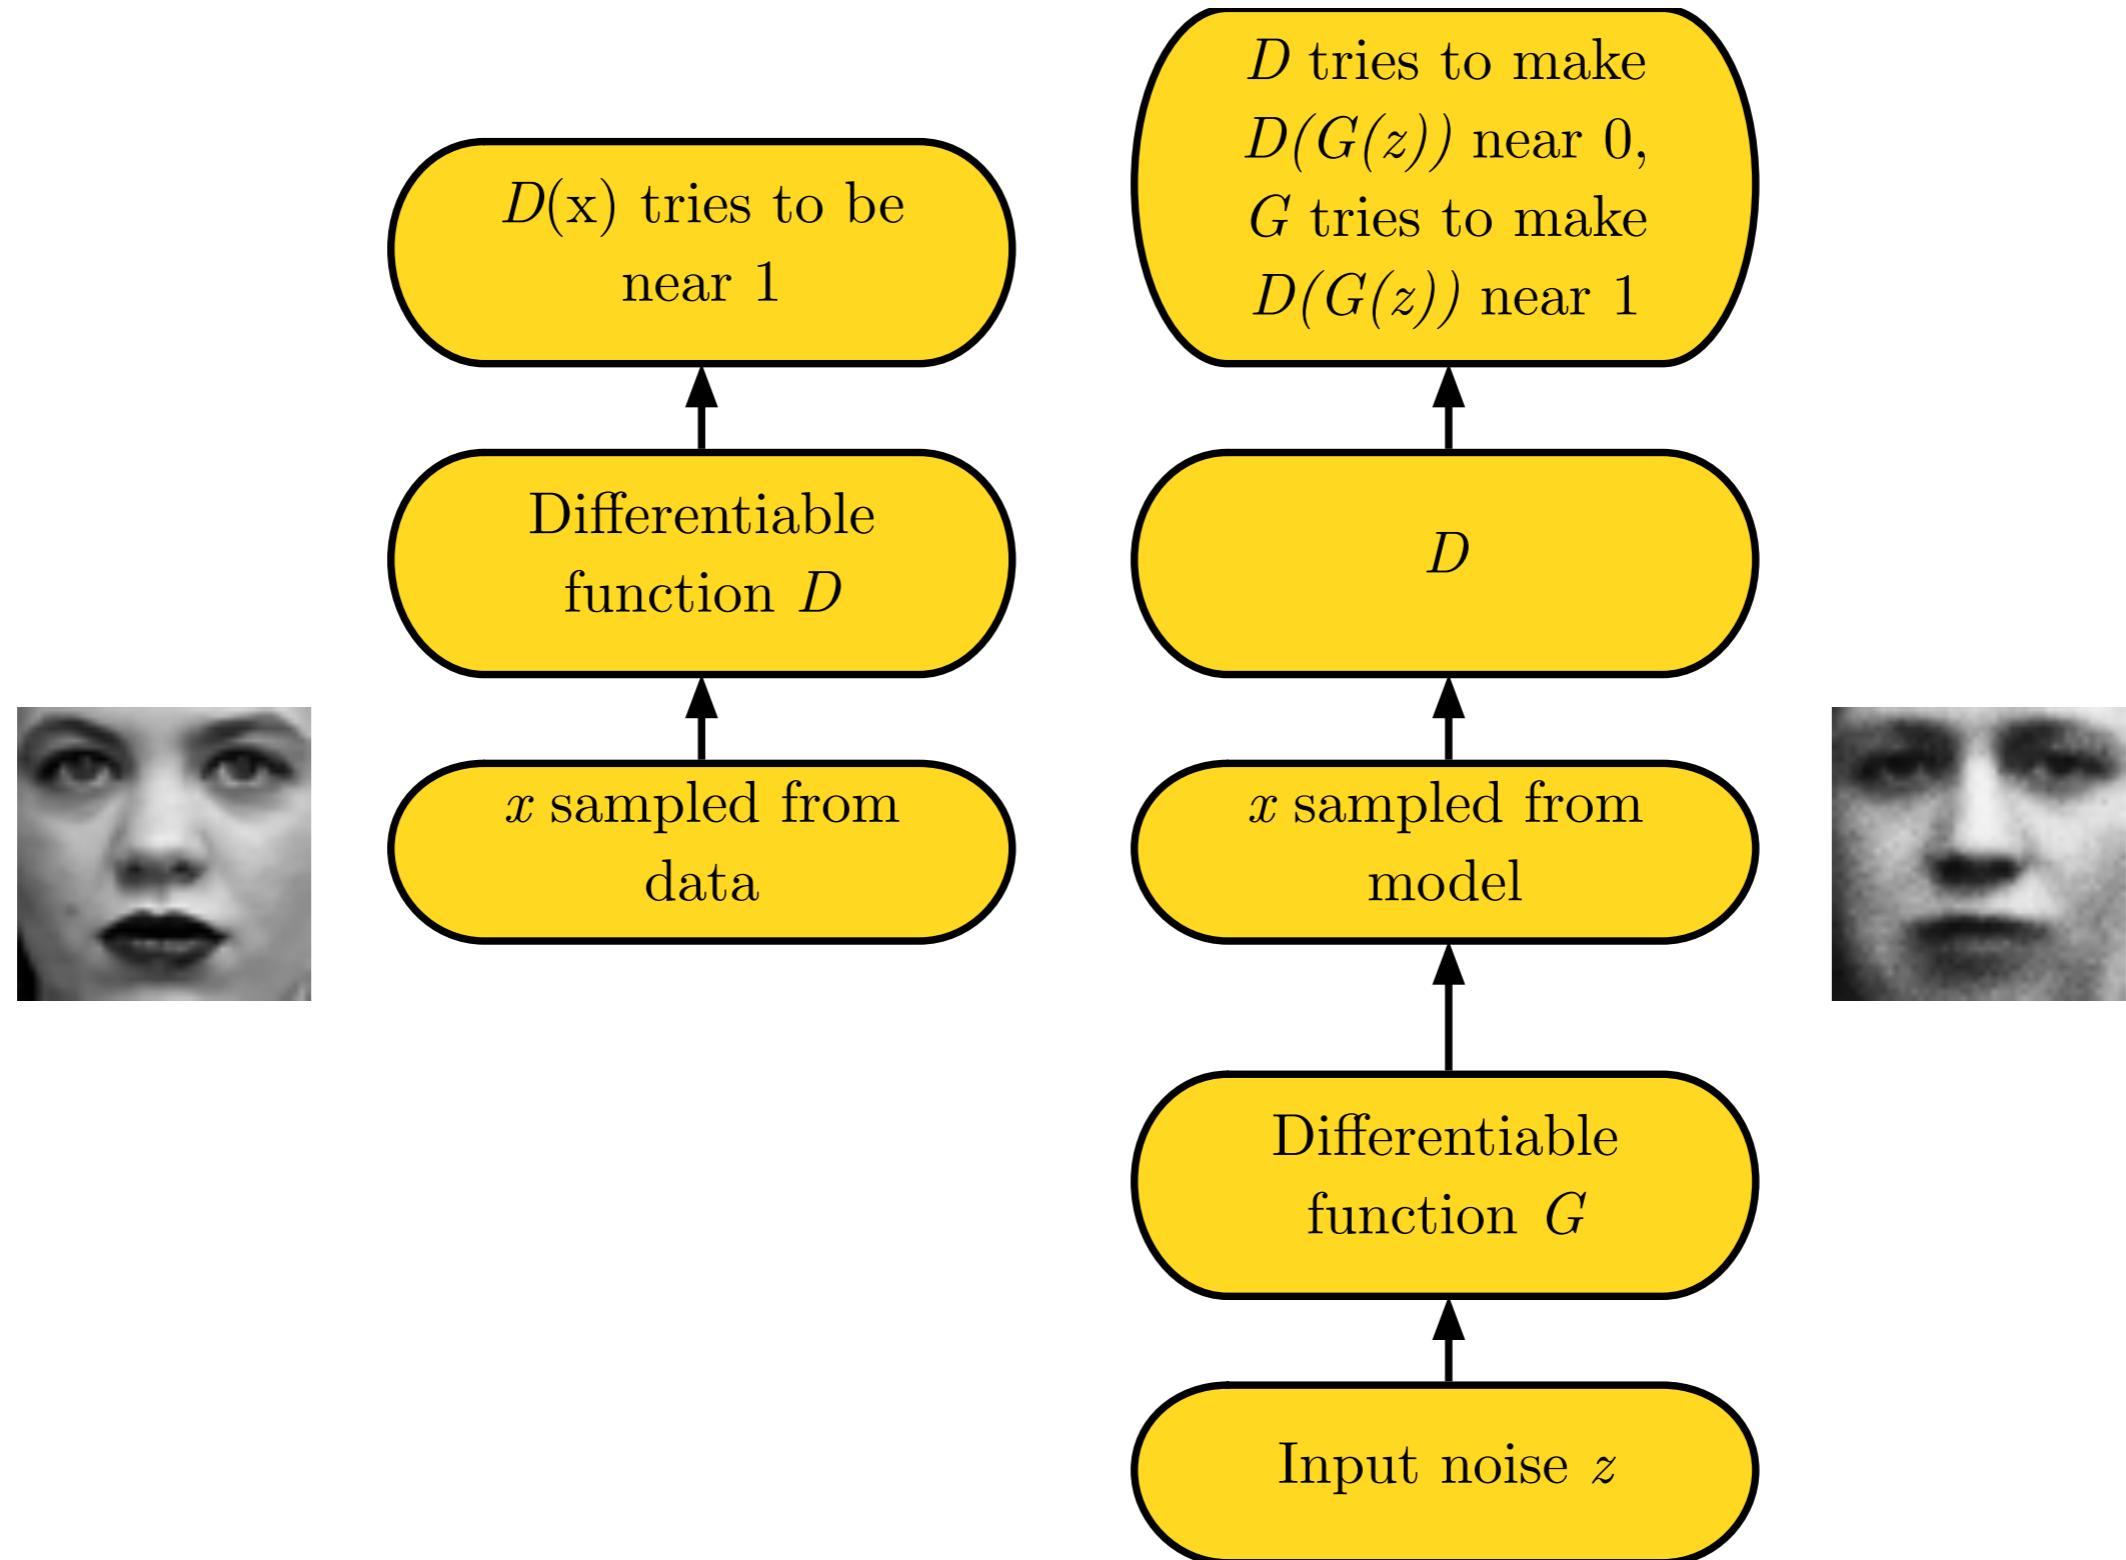
Generative Adversarial Network (Goodfellow 2014)



Binary Classifier with a binary cross entropy loss

$$D_\phi(\tilde{\mathbf{x}}) \in (0, 1)$$

Generative Adversarial Network (Goodfellow 2014)



Generative Adversarial Network (Goodfellow 2014)

- Use SGD-like algorithm of choice (Adam) on two mini batches simultaneously:
 - A minibatch of training examples
 - A minibatch of generated samples
- Minimax Game:

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z}} \log (1 - D(G(\mathbf{z})))$$
$$J^{(G)} = -J^{(D)}$$

- Equilibrium is a saddle point of the discriminator loss
- Generator minimizes the log-probability of the discriminator being correct

Additional reading: [LINK](#)

Generative Adversarial Network (Goodfellow 2014)

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Generative Adversarial Network (Goodfellow 2014)

Proposition 1. *For G fixed, the optimal discriminator D is*

$$D_G^*(\mathbf{x}) = \frac{p_{data}(\mathbf{x})}{p_{data}(\mathbf{x}) + p_g(\mathbf{x})} \quad (2)$$

obtained by $\frac{\delta}{\delta D(\mathbf{x})} J^{(D)} = 0$

Proposition 2. *If G and D have enough capacity, and at each step of Algorithm 1, the discriminator is allowed to reach its optimum given G , and p_g is updated so as to improve the criterion*

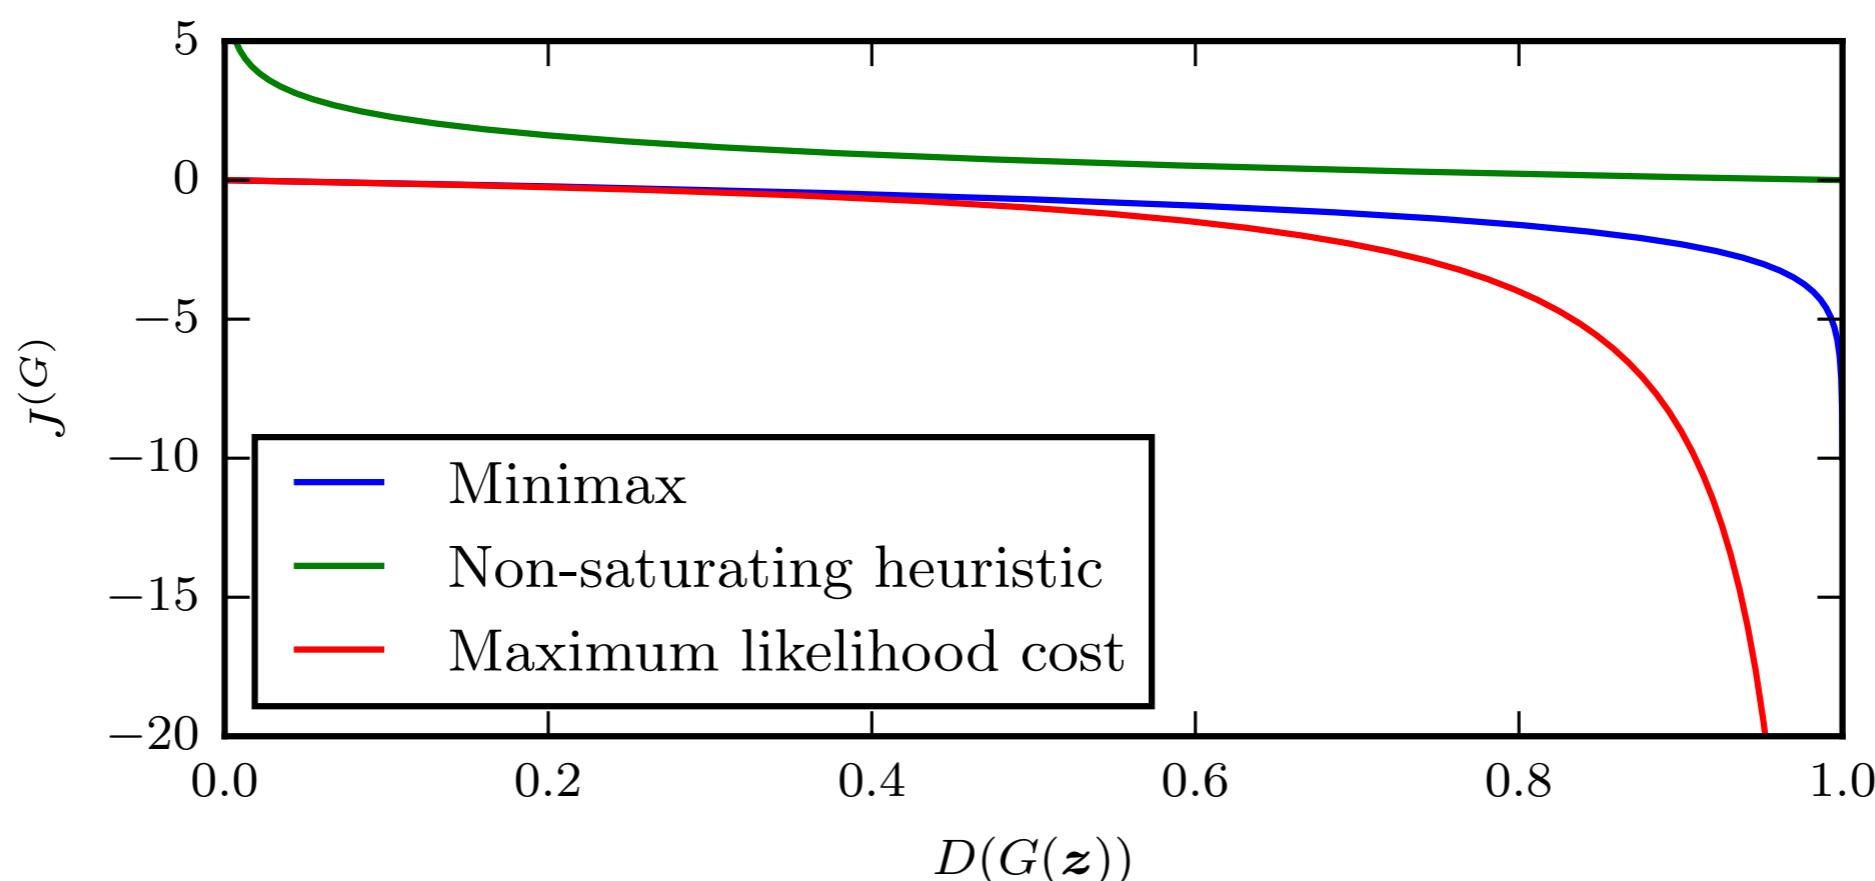
$$\mathbb{E}_{\mathbf{x} \sim p_{data}} [\log D_G^*(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim p_g} [\log(1 - D_G^*(\mathbf{x}))]$$

then p_g converges to p_{data}

Generative Adversarial Network (Goodfellow 2014)

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z}} \log (1 - D(G(\mathbf{z}))) \quad \text{Minimax}$$
$$J^{(G)} = -J^{(D)}$$

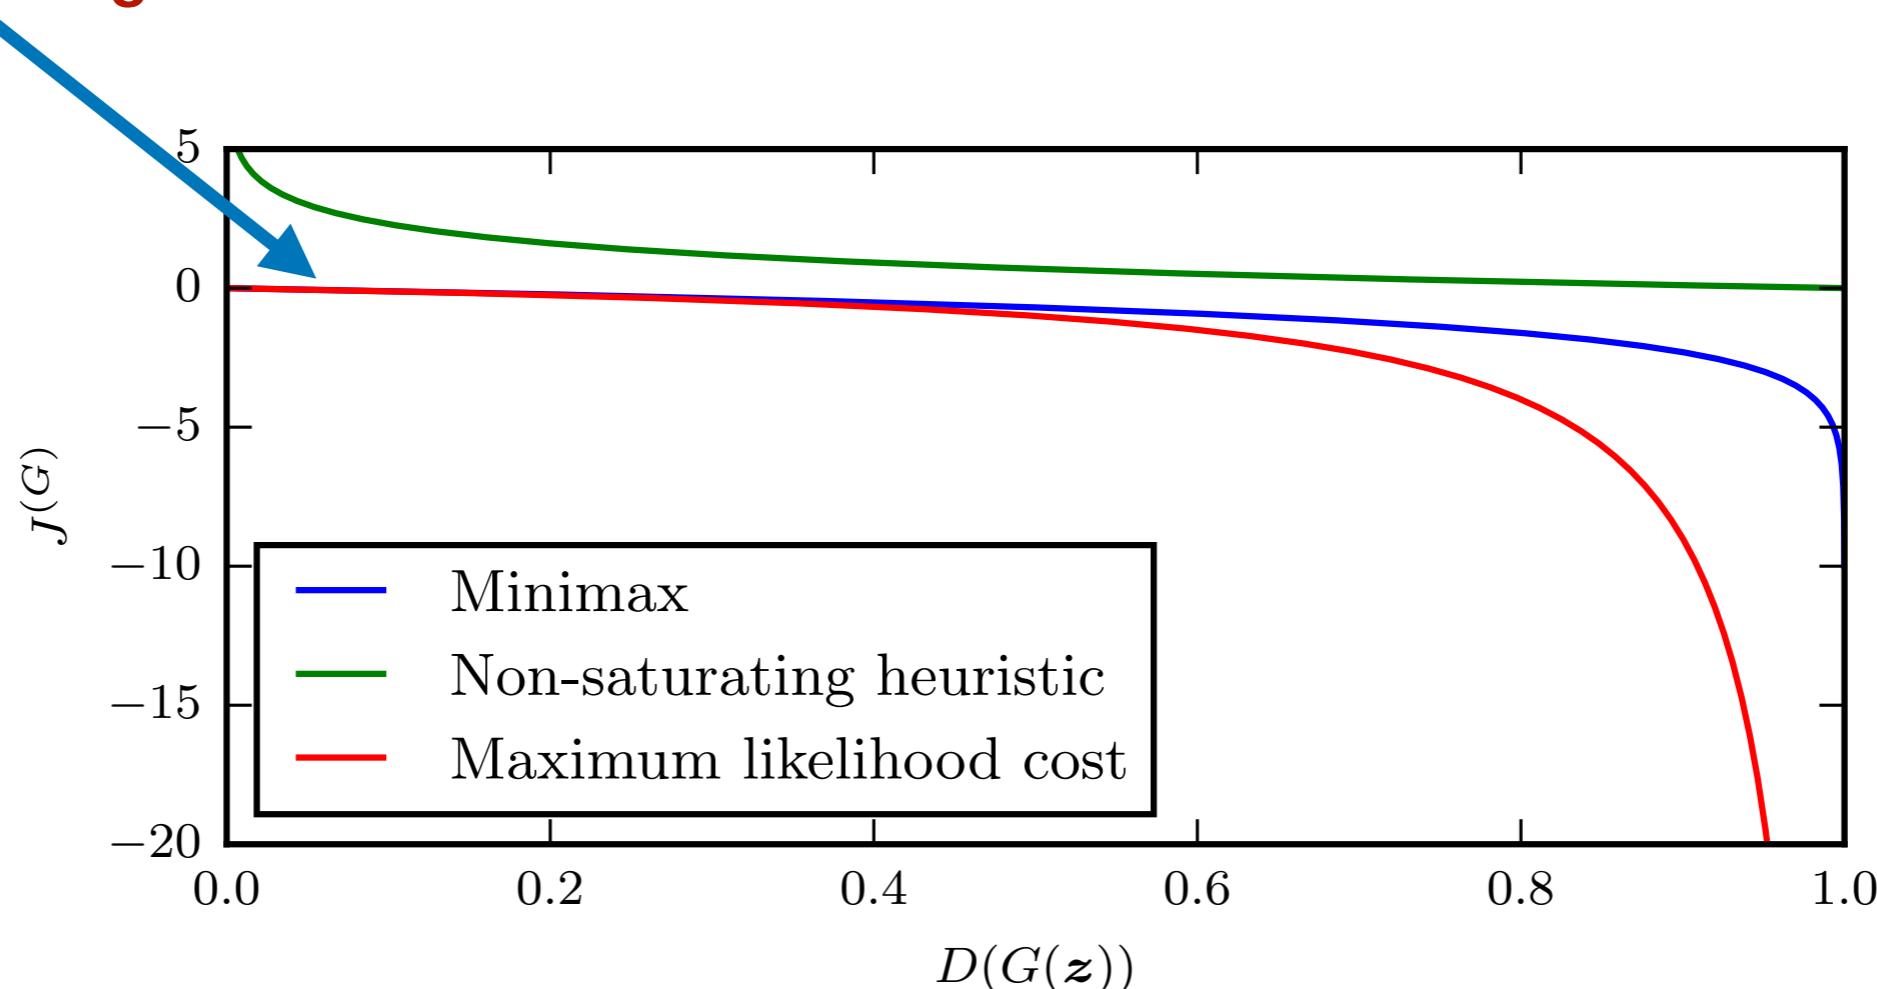
- When the discriminator successfully rejects generator samples with high confidence, **the generator's gradient vanishes**



Generative Adversarial Network (Goodfellow 2014)

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z}} \log (1 - D(G(\mathbf{z}))) \quad \text{Minimax}$$
$$J^{(G)} = -J^{(D)}$$

- When the discriminator successfully rejects generator samples with high confidence, **the generator's gradient vanishes**



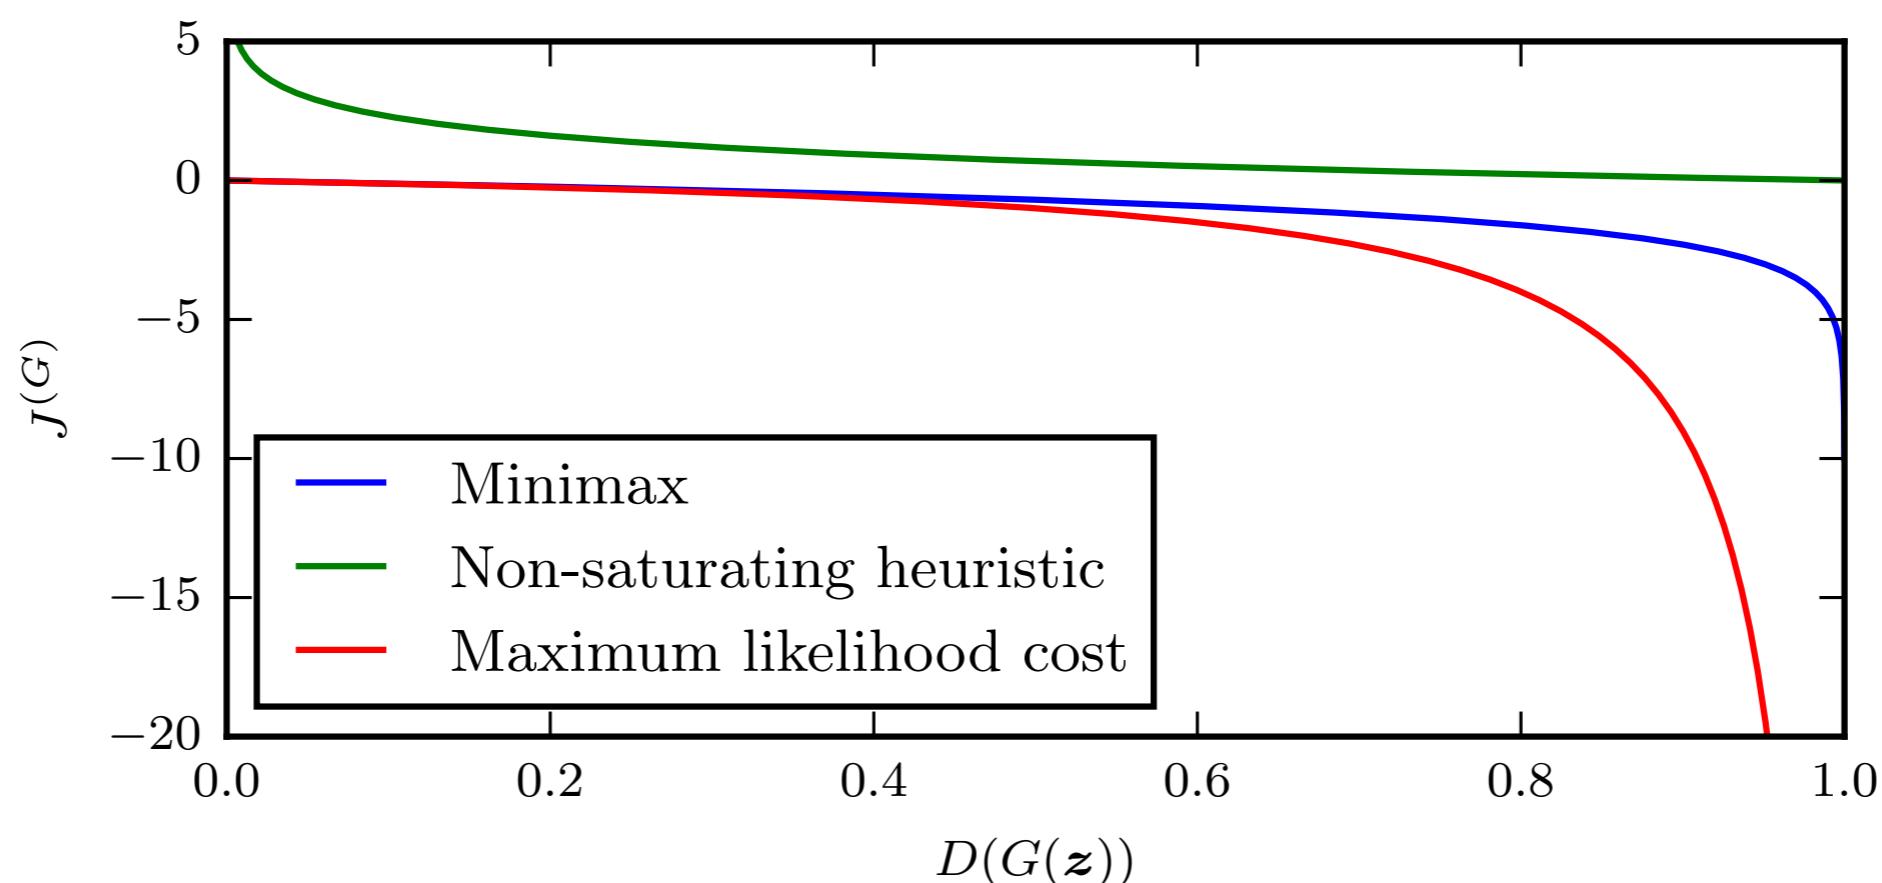
Generative Adversarial Network (Goodfellow 2014)

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) - \frac{1}{2} \mathbb{E}_z \log (1 - D(G(z)))$$

Non-saturating heuristic

$$J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log D(G(z))$$

- The generator maximizes the log probability of the discriminator **being mistaken**
- The sole motivation for this version of the game is to **ensure that each player has a strong gradient when that player is “losing” the game.**



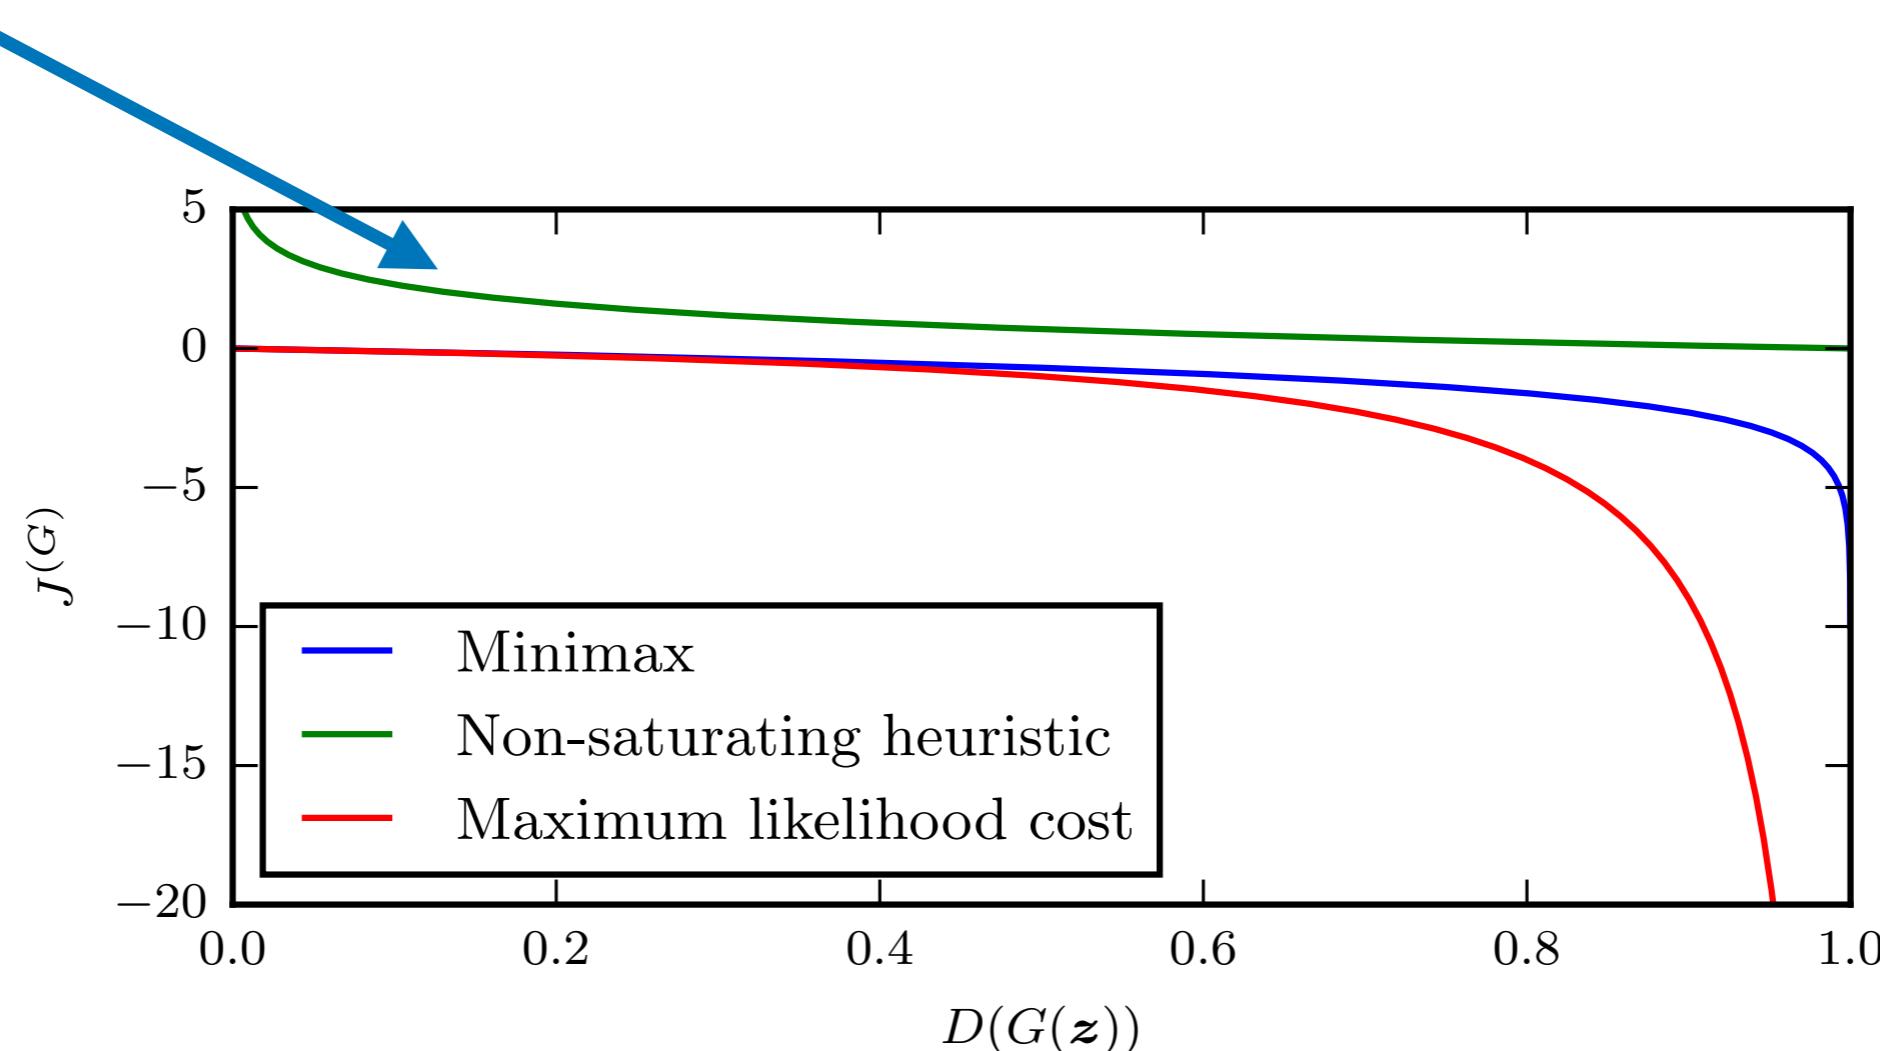
Generative Adversarial Network (Goodfellow 2014)

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) - \frac{1}{2} \mathbb{E}_z \log (1 - D(G(z)))$$

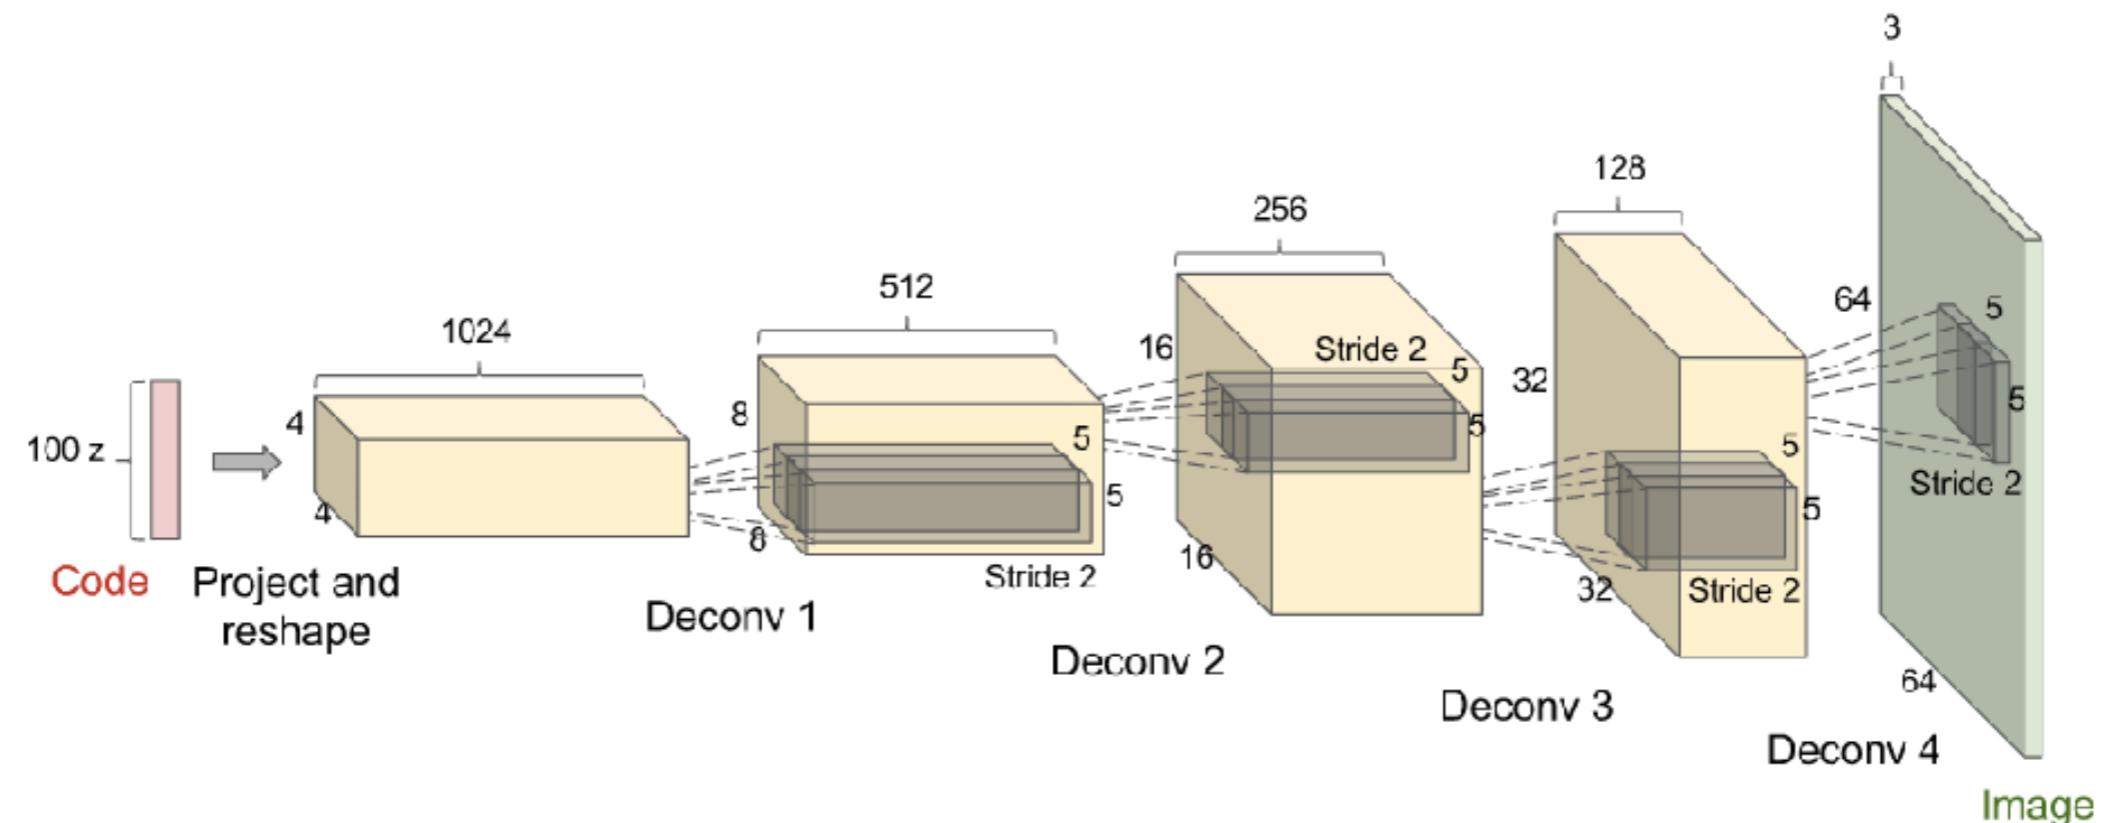
Non-saturating heuristic

$$J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log D(G(z))$$

- The generator maximizes the log probability of the discriminator **being mistaken**
- The sole motivation for this version of the game is to **ensure that each player has a strong gradient when that player is “losing” the game.**

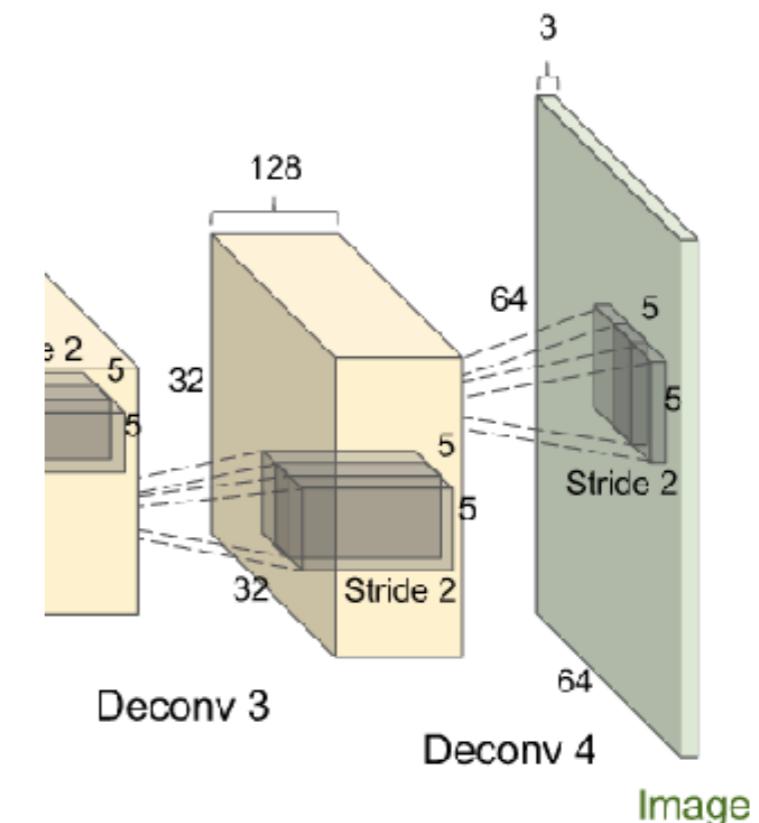
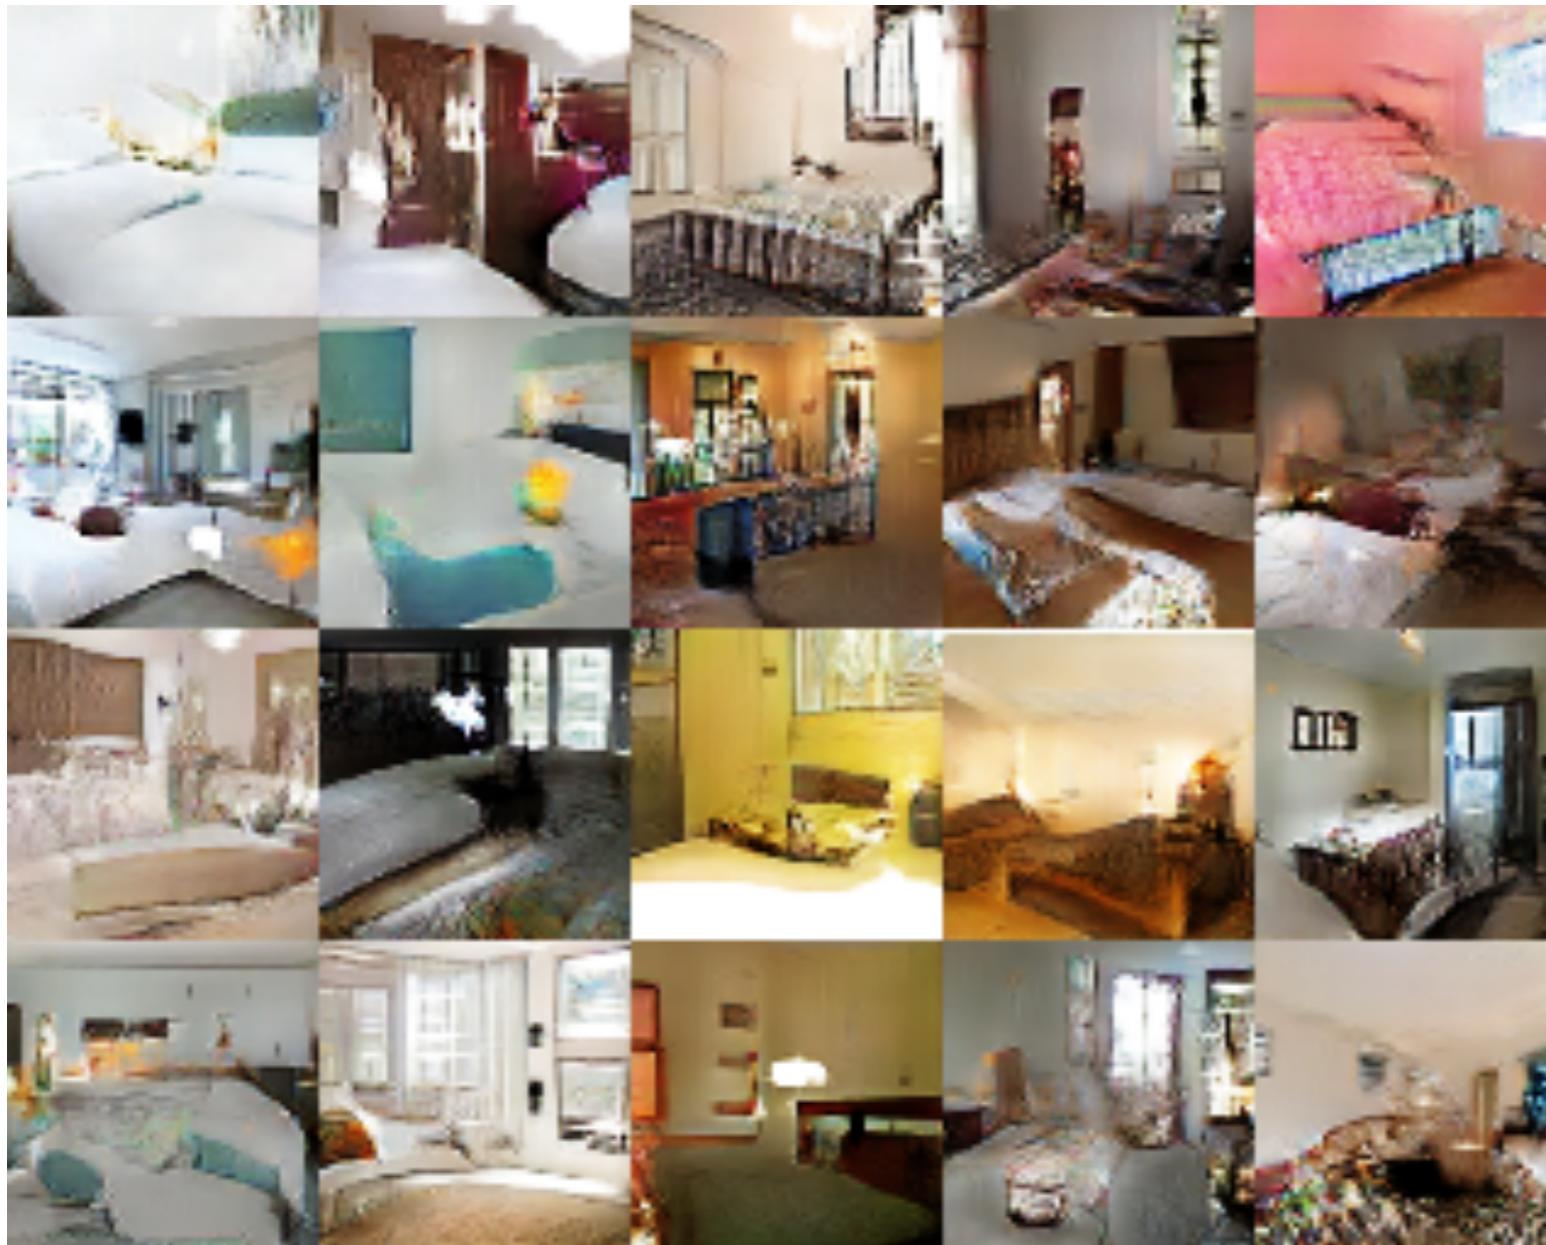


DCGAN Architecture (Radford 2015)



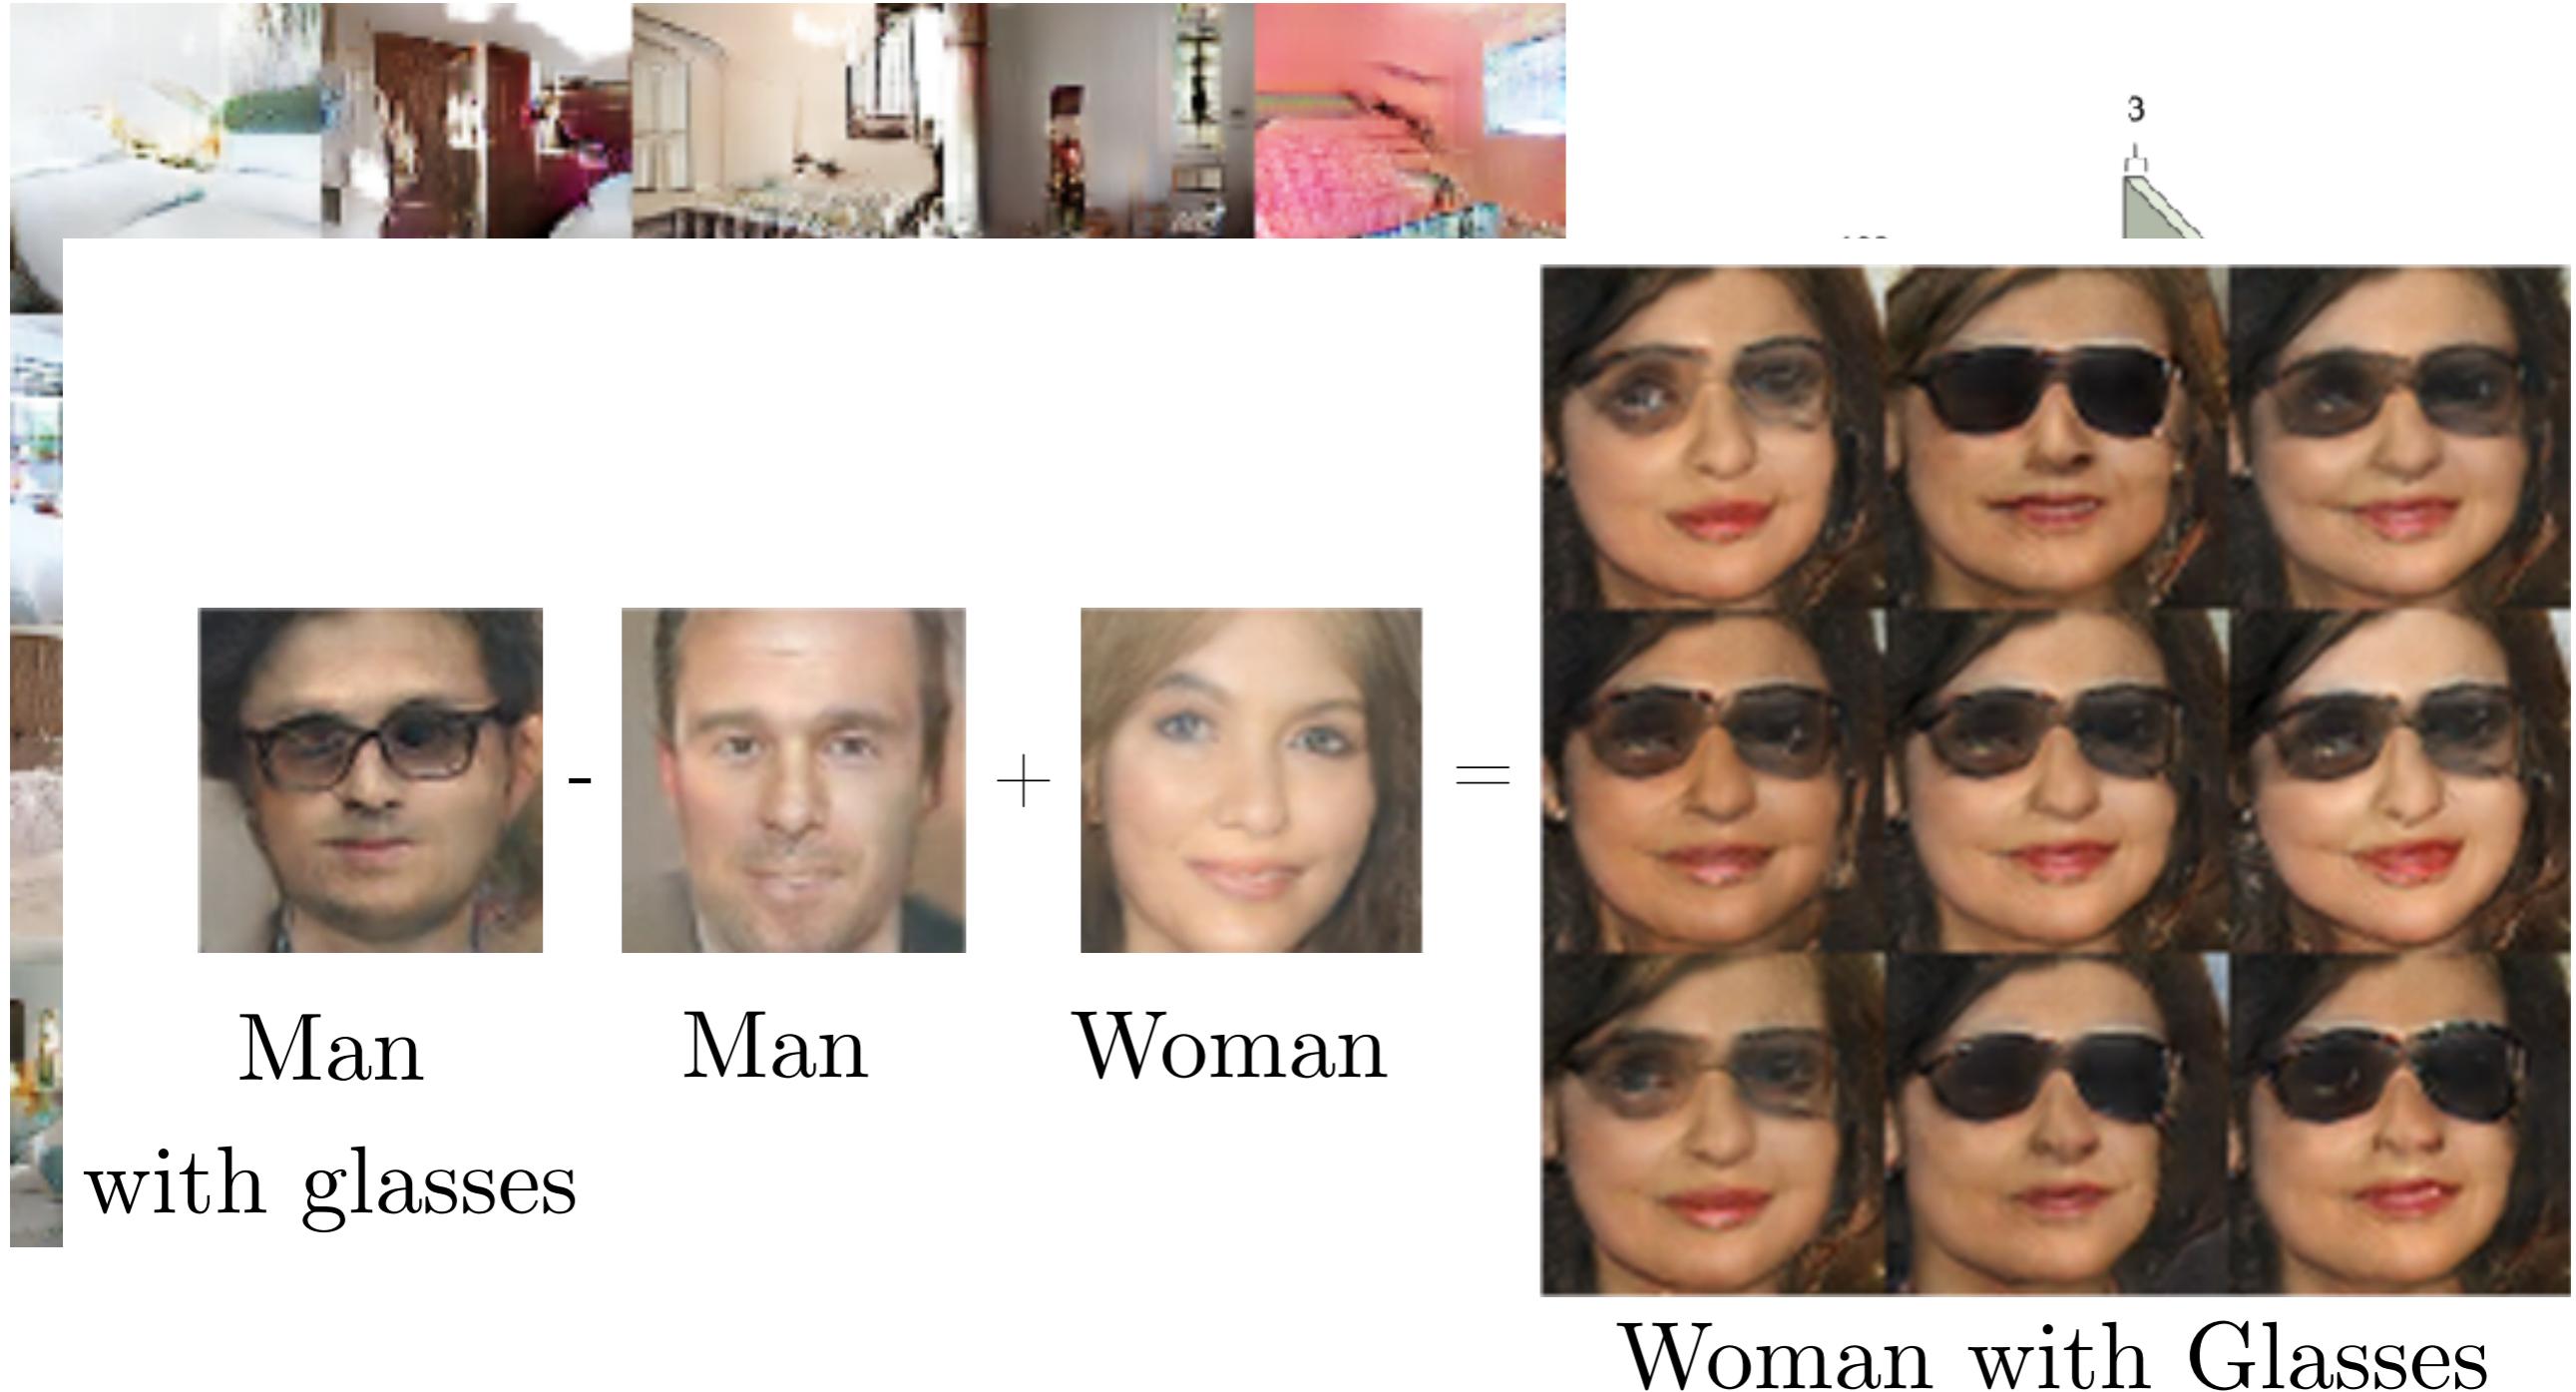
(Radford et al 2015)

DCGAN Architecture (Radford 2015)



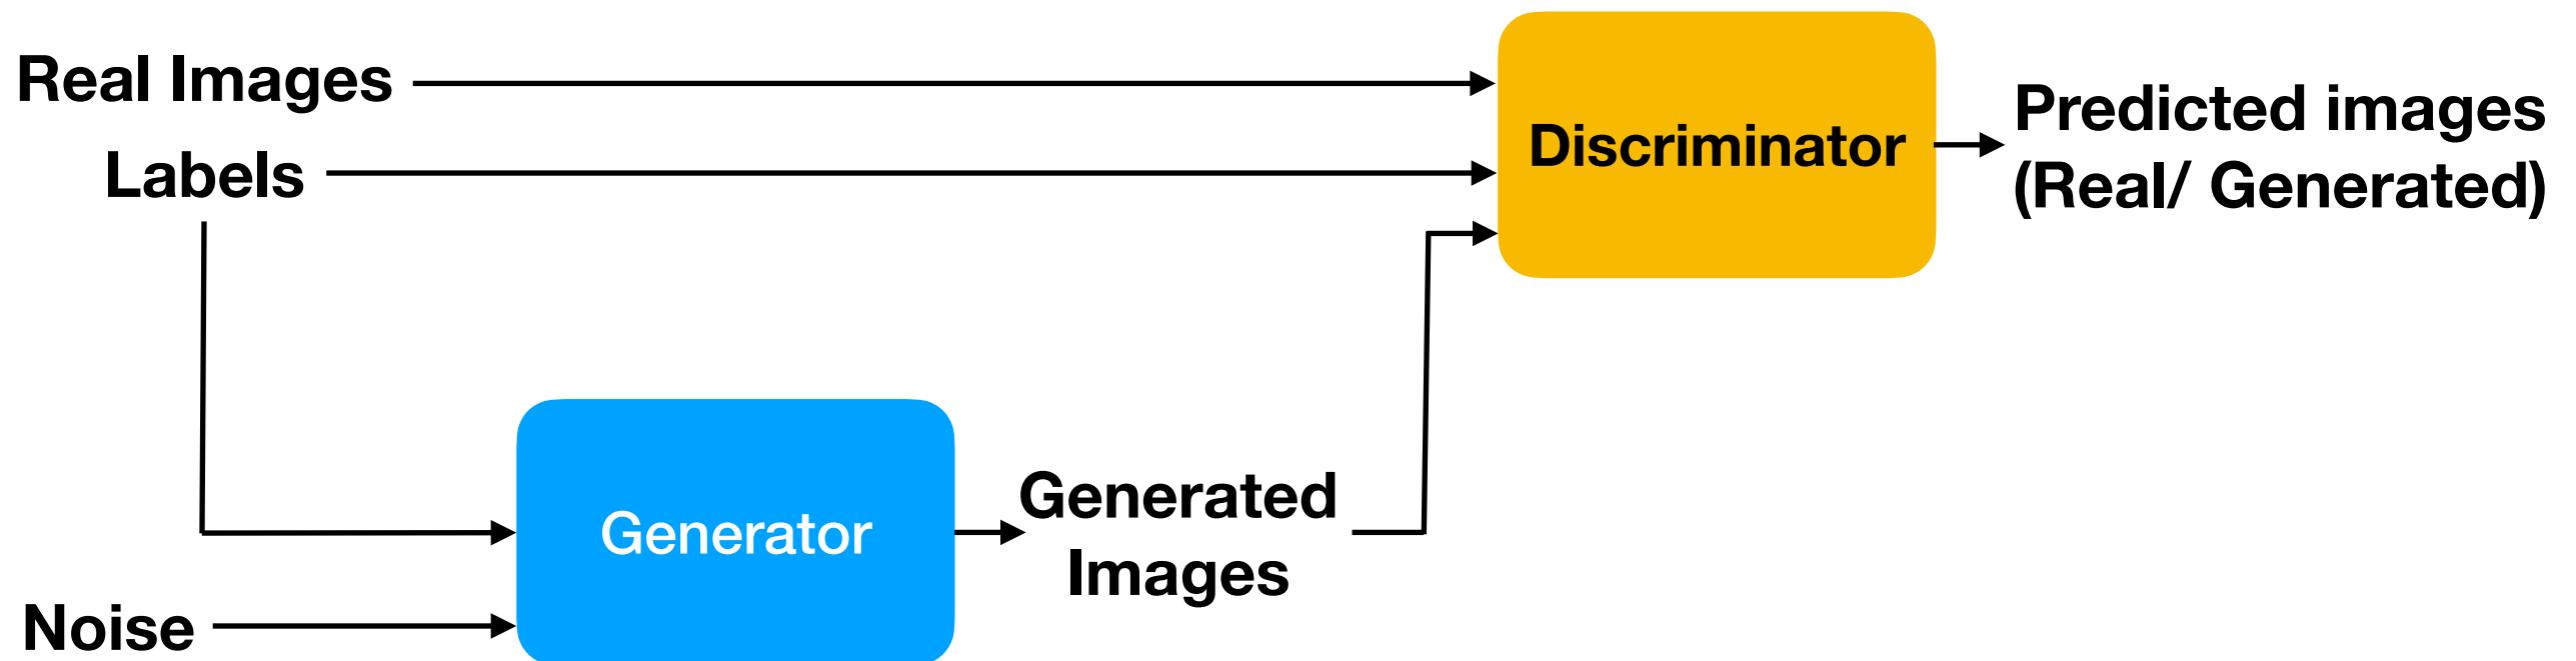
(Radford et al 2015)

DCGAN Architecture (Radford 2015)

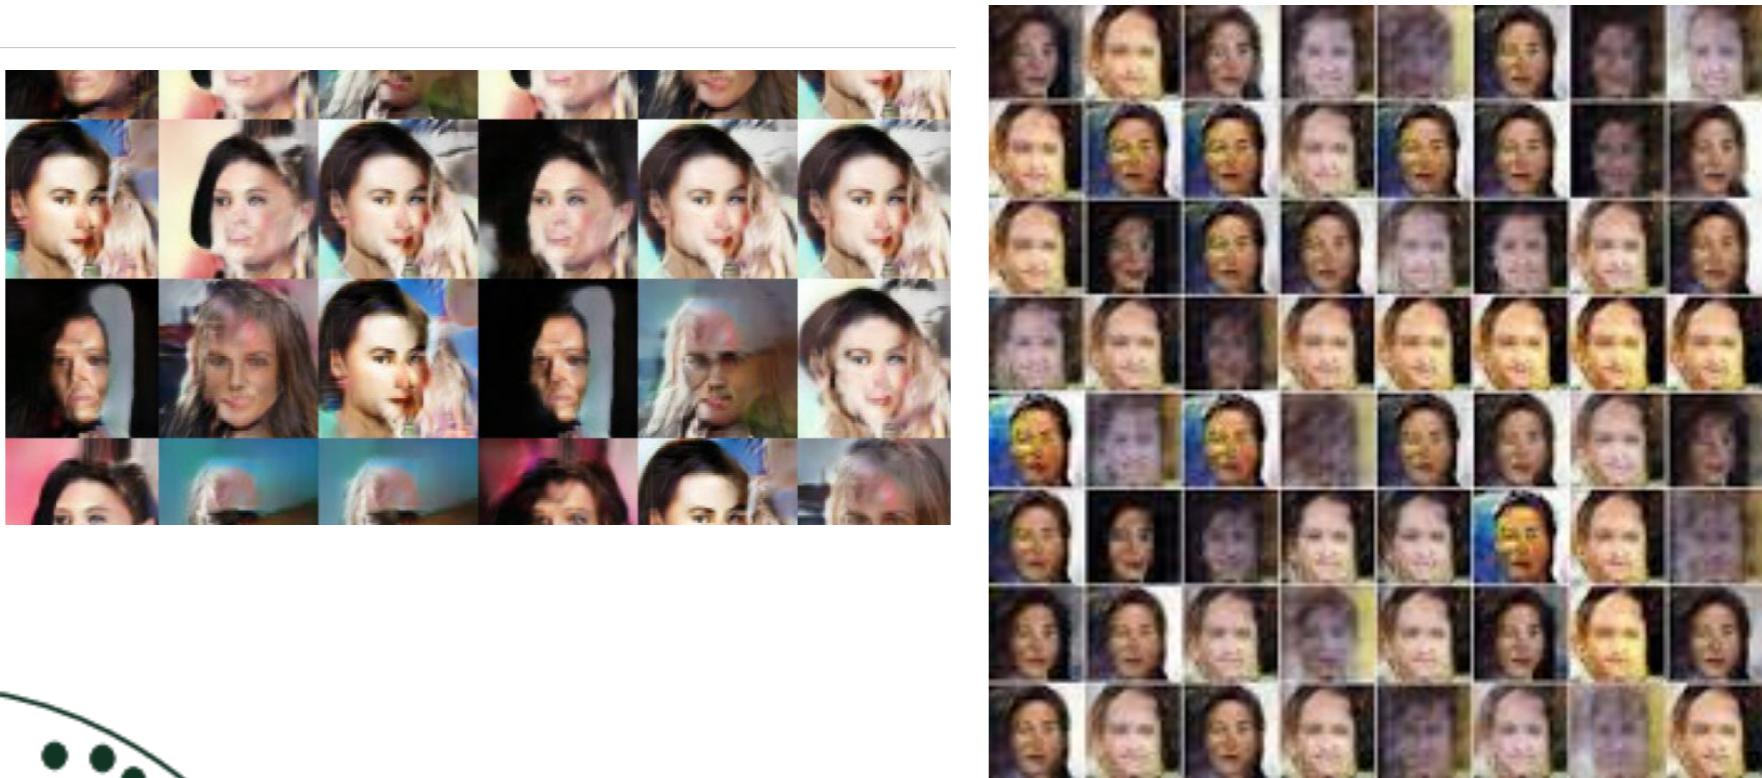
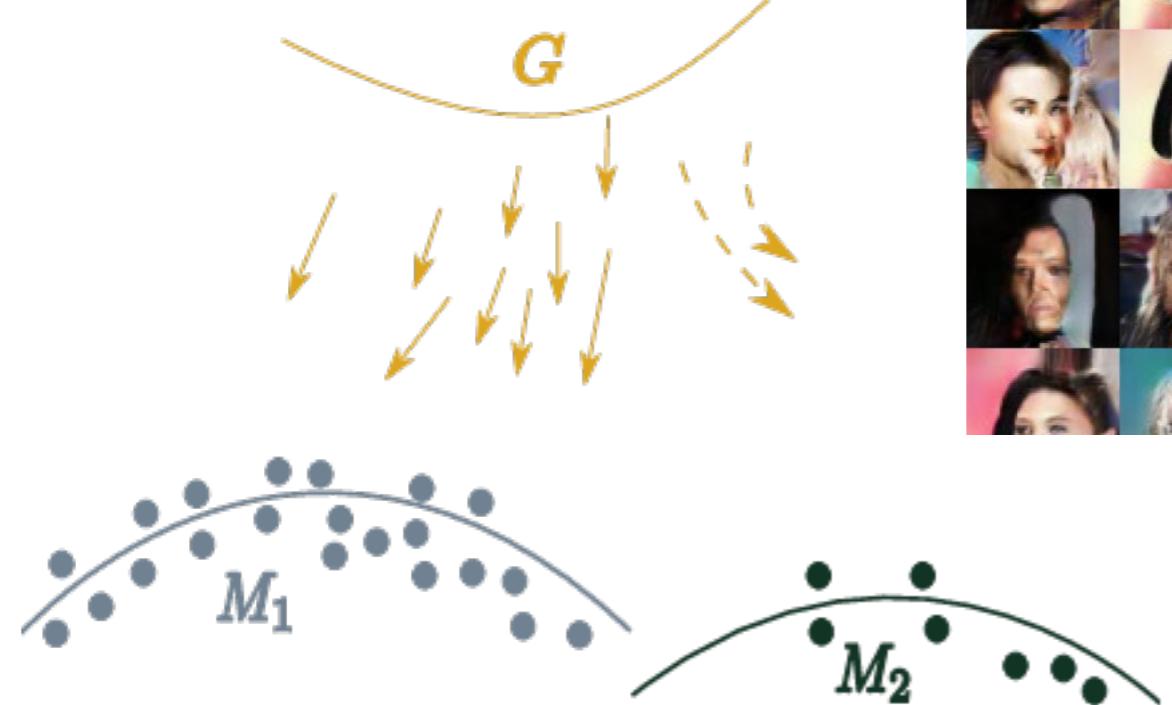


Conditional GAN (CGAN)

- In a **standard GAN** we have **no control on the class of output** we are producing
- **Conditional GAN: Add a y-label** to generator and discriminator
- Generator and discriminator **trained for the specific y-label**
- In **generation**, we can **pass the y-label** to generator to produce a class of ouput



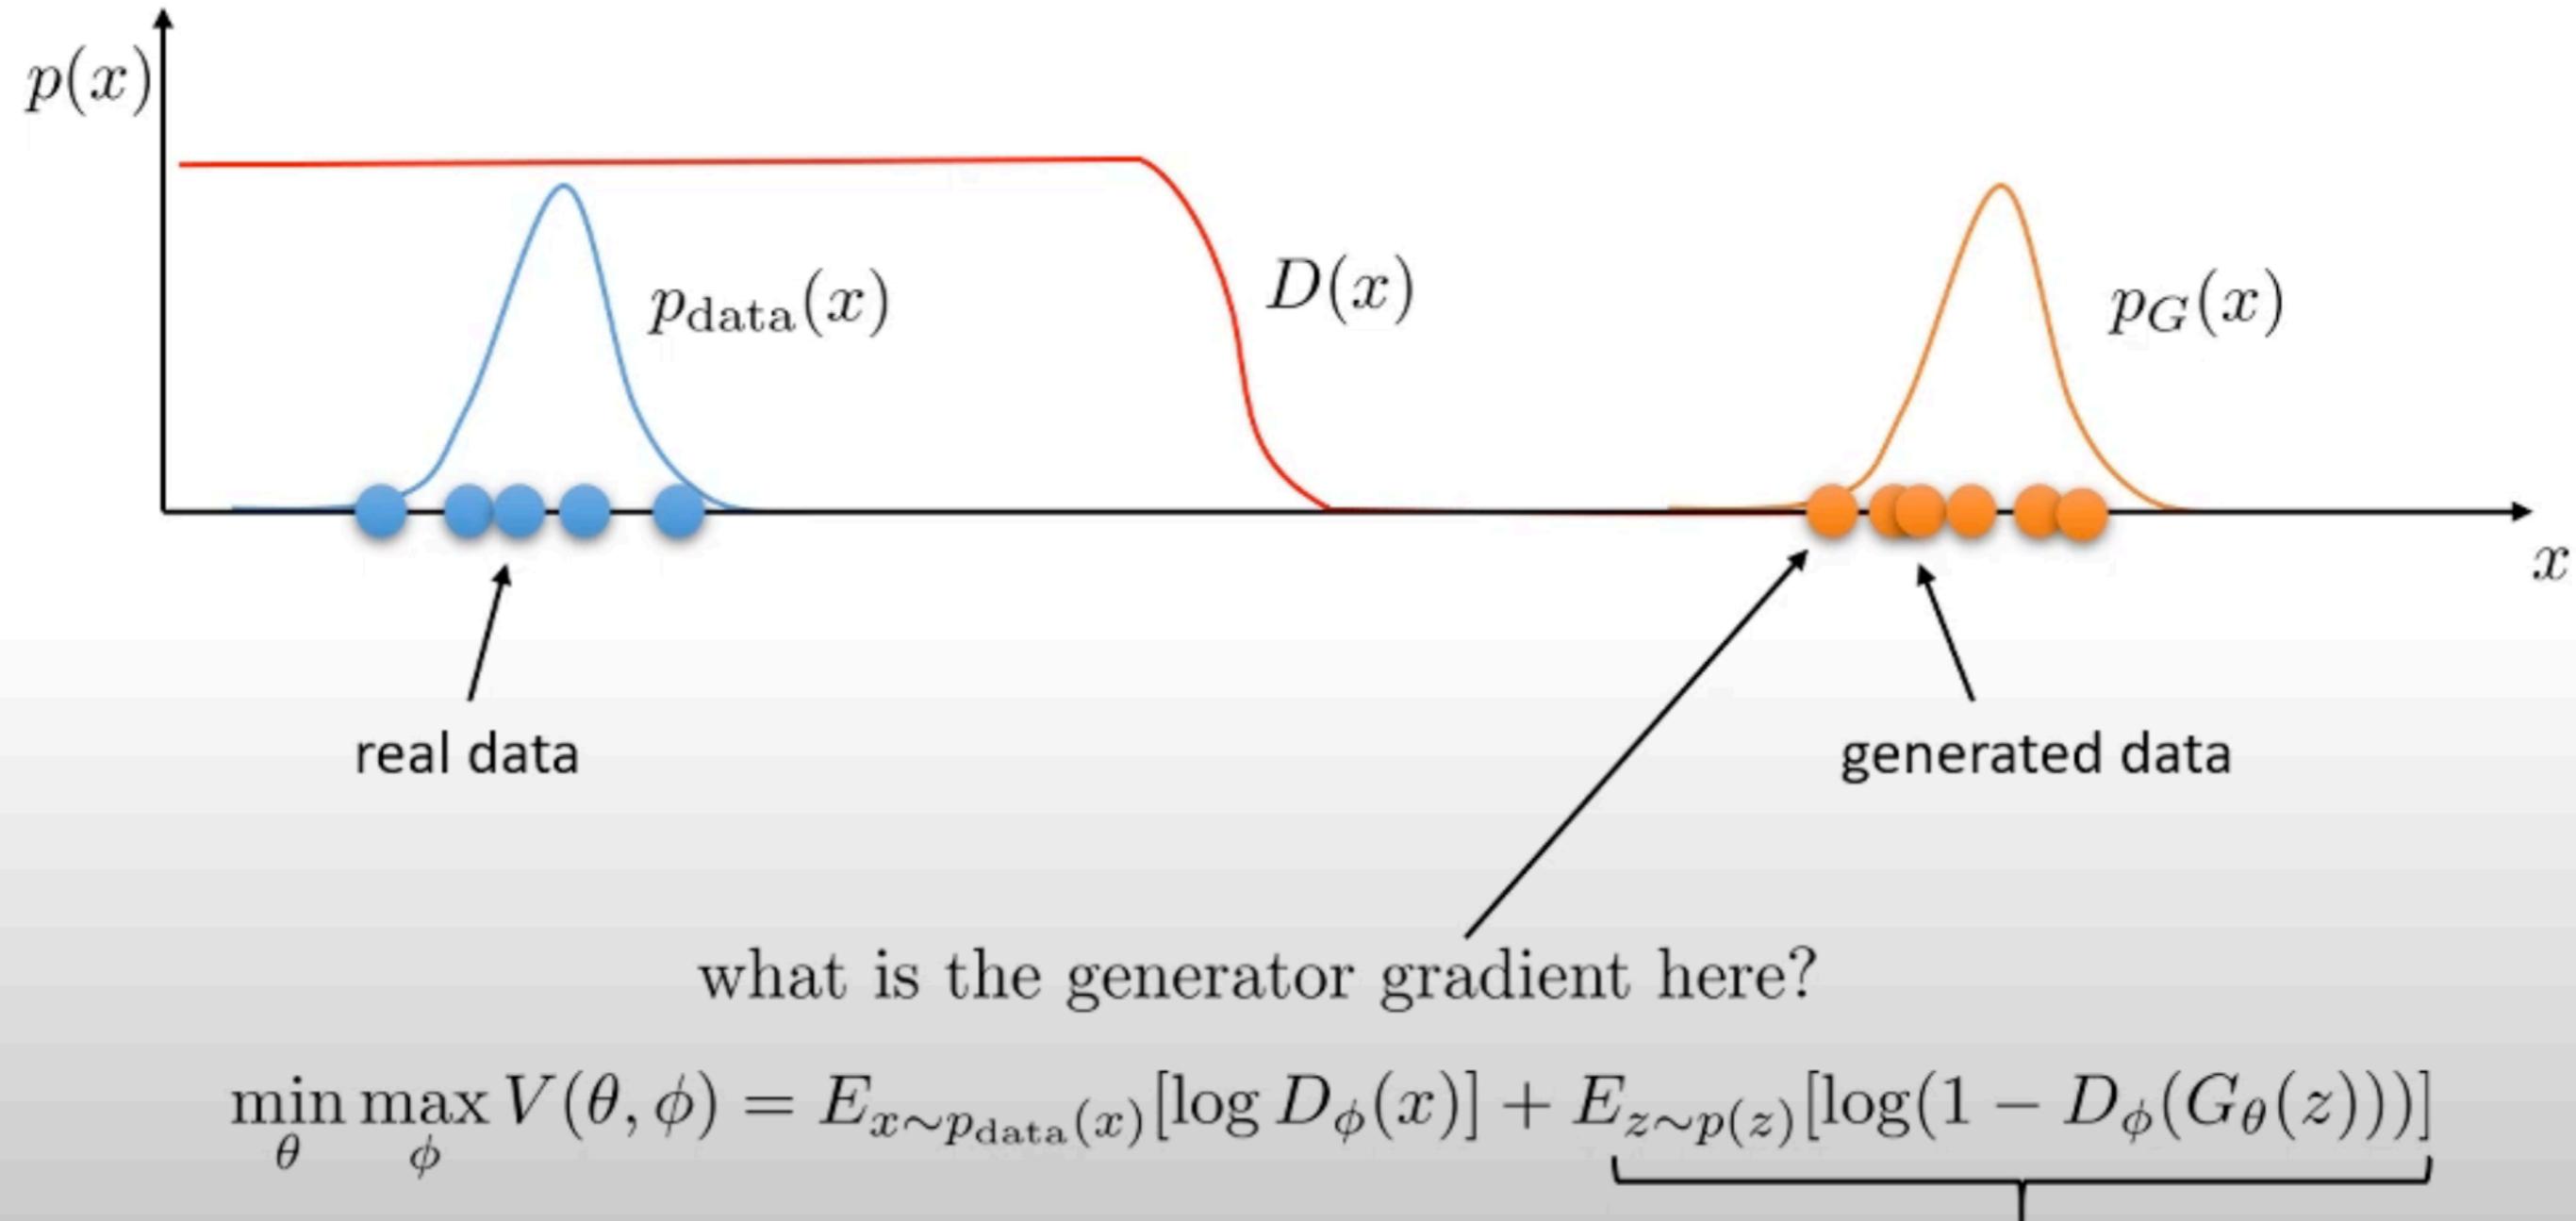
A problem of GANs: Mode Collapse



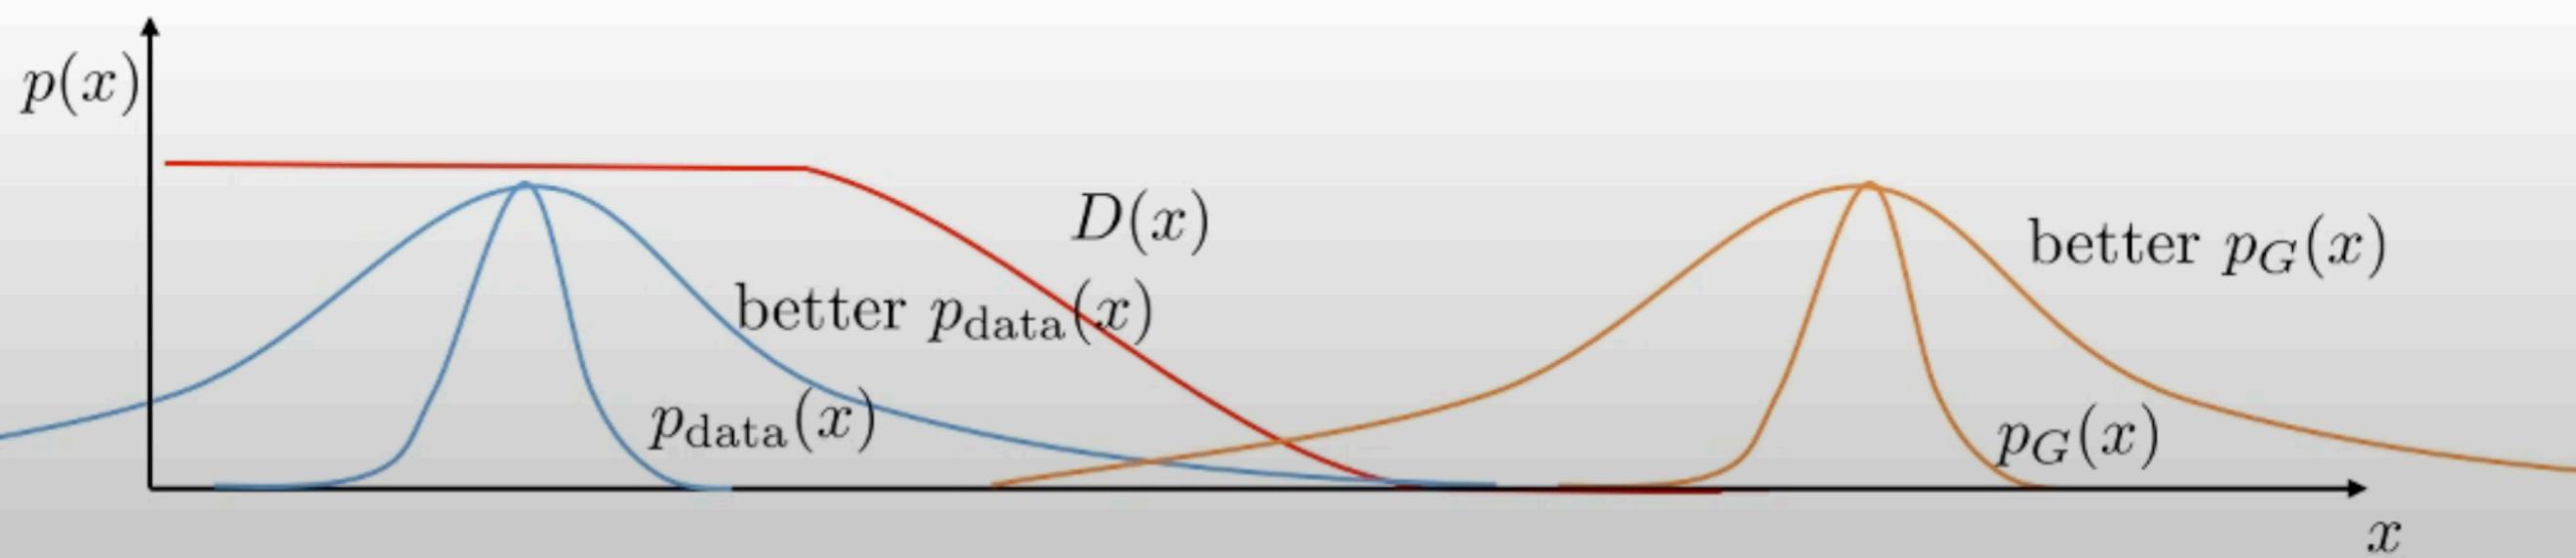
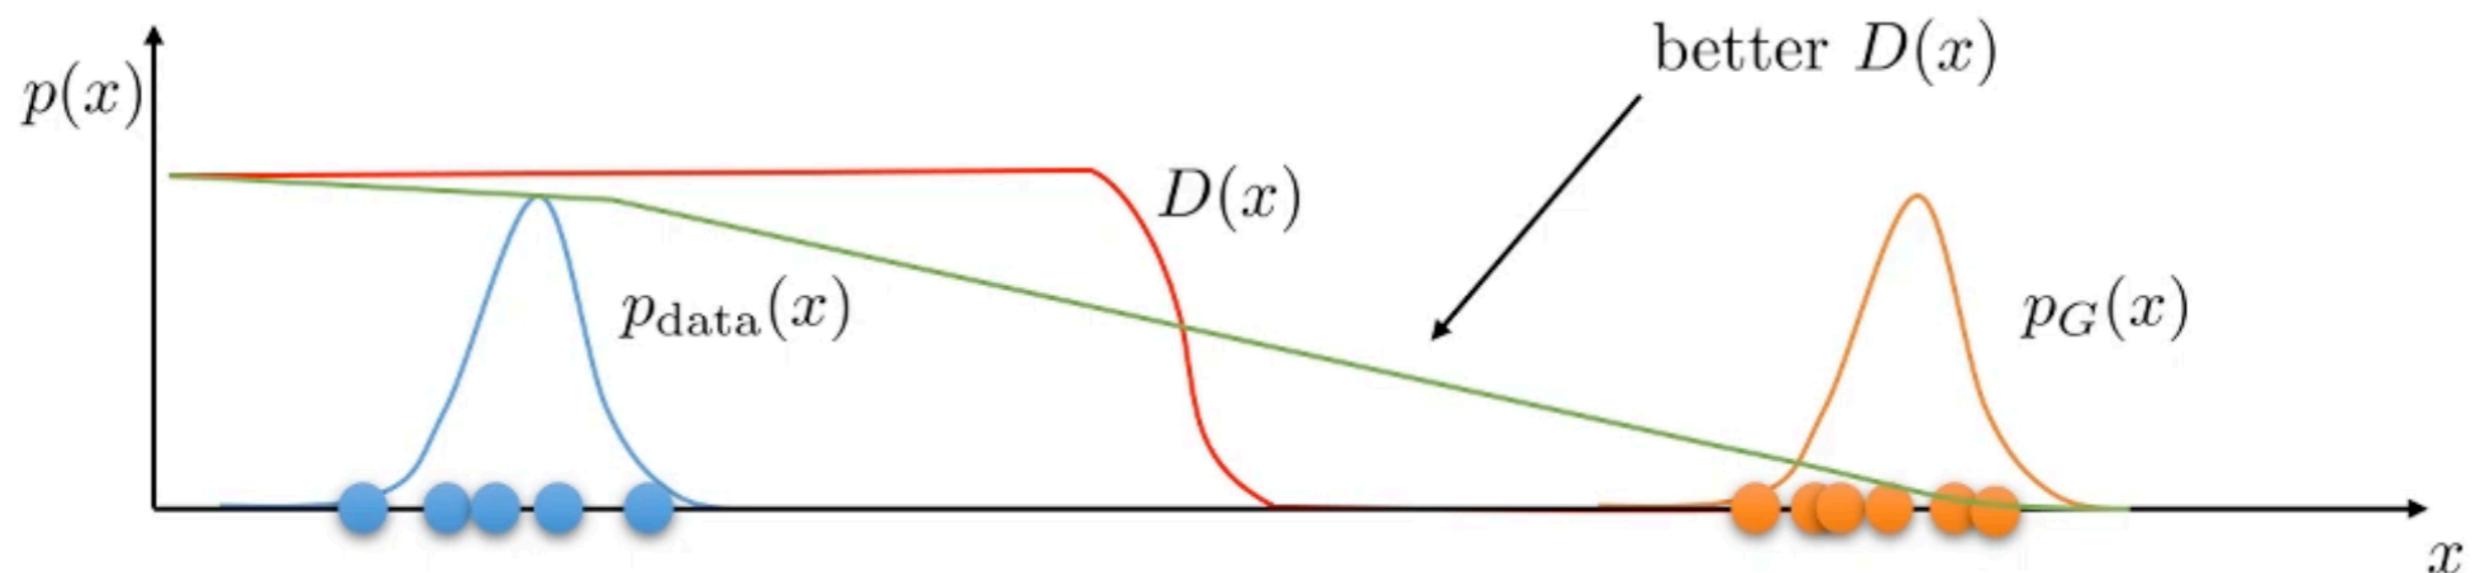
8	8	8	3	8	3	8	8	8	8
3	8	3	8	3	8	8	8	8	3
3	8	3	8	3	8	3	8	8	3
3	8	3	8	3	8	3	8	8	3
3	8	3	3	3	3	3	3	3	3
3	8	3	3	3	3	8	3	3	3
3	8	3	3	3	3	8	3	3	3
3	8	3	3	3	3	8	3	3	3
3	8	3	3	3	3	8	3	3	3
3	8	3	3	3	3	8	3	3	3

2	6	3	4	3	4	4	8	7	3	0	0	9	1	1	2	8	6	9	4	5	1	1	2	8	6	9	4	5	1	1	2		
6	1	4	9	4	2	3	4	4	1	0	9	1	0	5	2	0	3	5	0	6	0	0	3	5	0	6	0	0	3	5	0	6	0
2	3	6	2	6	5	7	3	0	0	9	0	8	0	8	1	5	9	1	0	8	8	8	8	8	9	1	0	8	8	8	8	8	
4	0	2	1	3	3	9	3	6	0	0	8	3	3	7	9	0	6	1	4	1	8	1	8	2	6	1	4	1	8	1	8	2	
7	9	3	7	4	5	6	6	4	5	8	8	0	0	9	0	8	9	0	8	0	8	7	5	2	3	8	8	0	6	7	5	2	
3	3	4	4	1	6	7	6	6	7	9	0	1	0	8	5	7	7	7	0	8	1	8	3	2	7	7	0	8	1	8	3	2	
1	4	2	7	4	7	3	9	9	1	9	4	5	2	8	8	9	3	0	8	7	8	1	0	8	3	0	8	7	8	1	0	8	
0	6	9	7	1	6	1	4	8	9	7	1	3	4	9	8	9	3	5	0	9	5	7	8	9	3	5	0	9	5	7	8	9	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	8	3	
3	3	8	8	3	8	3																											

Why is training GANs hard?



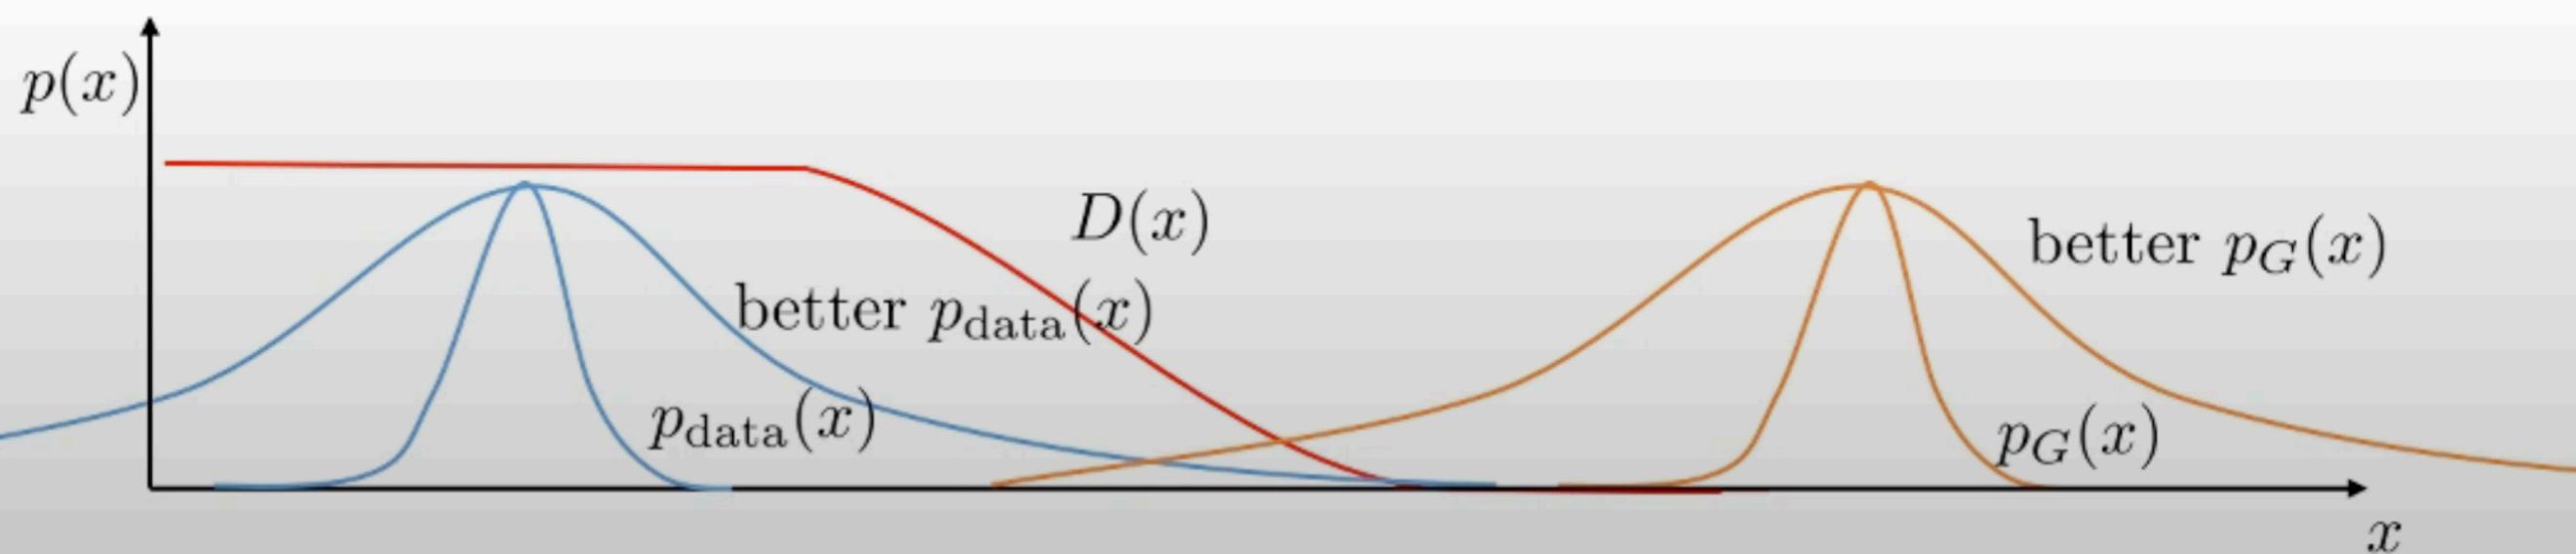
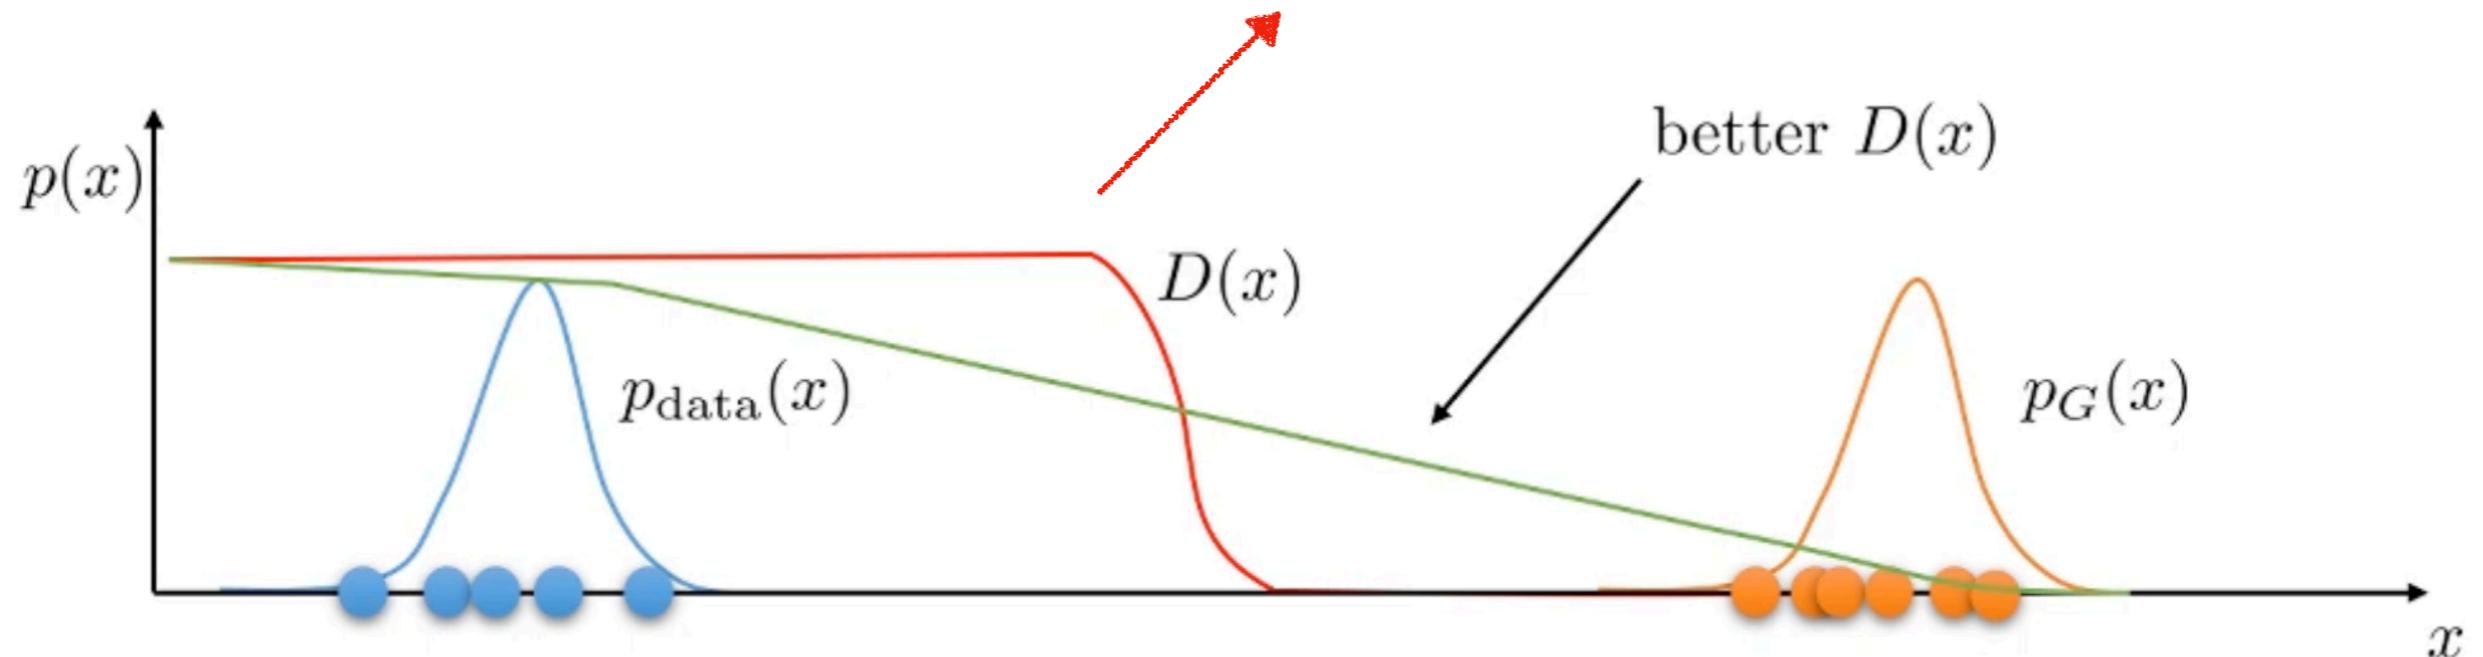
Why is training GANs hard?



Why is training GANs hard?

Least-squares GAN (LSGAN): discriminator outputs real-valued number

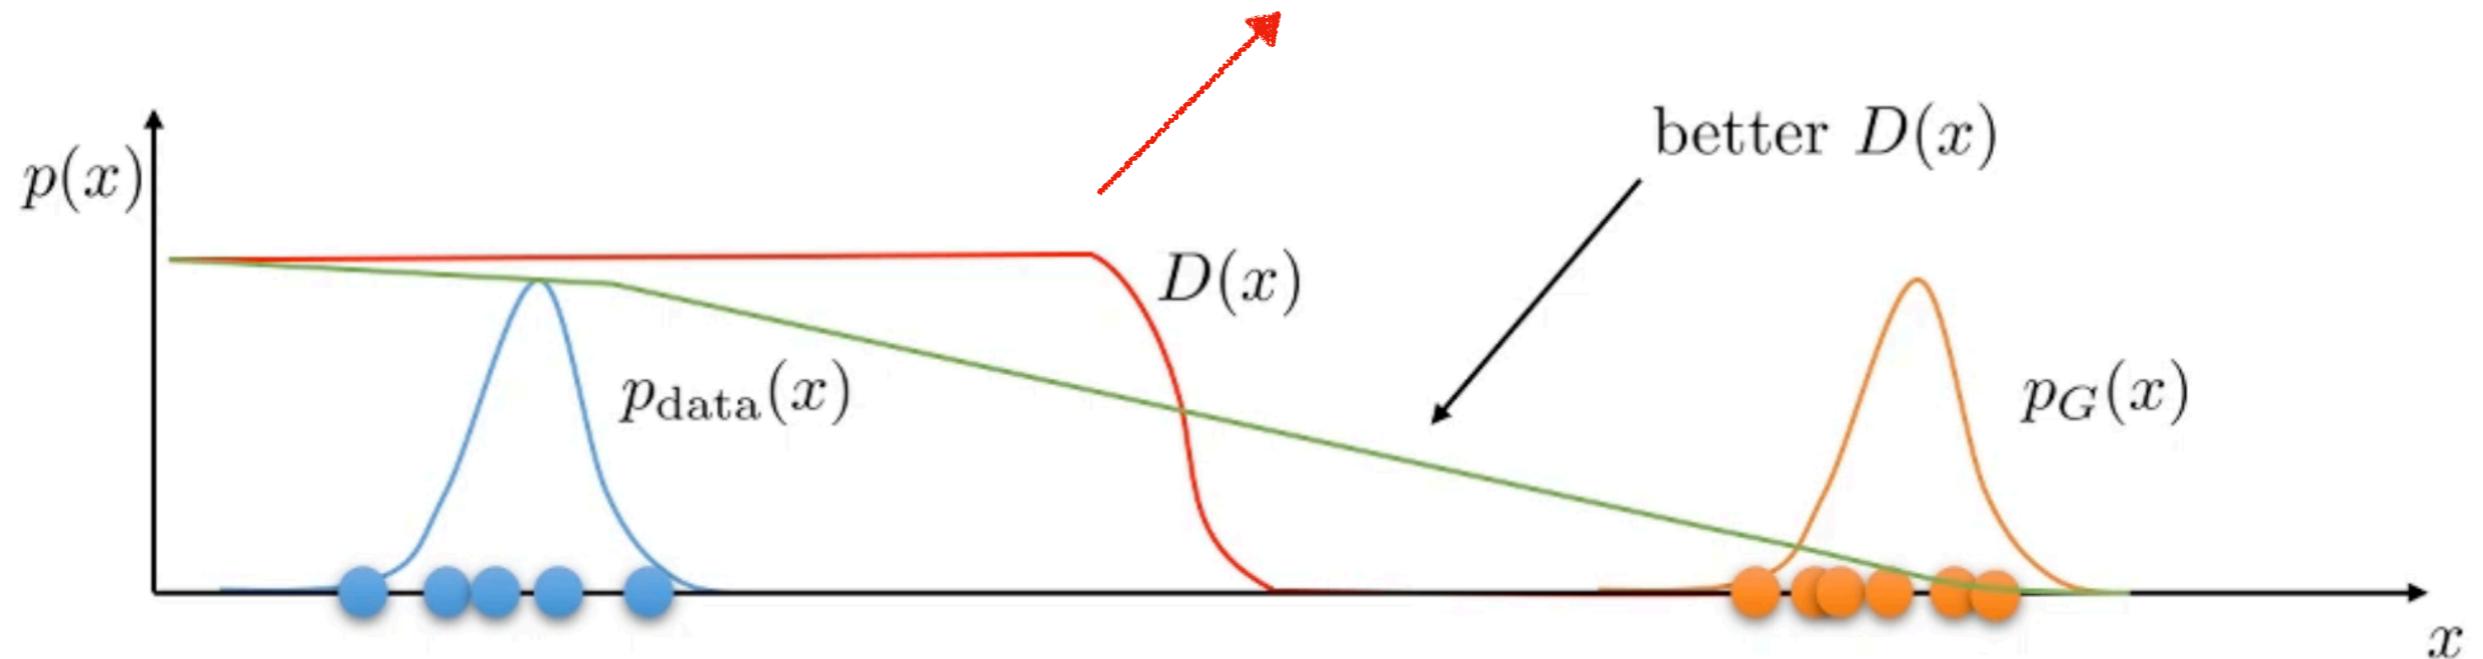
Wasserstein GAN (WGAN): discriminator is constrained to be Lipschitz-continuous
Gradient Penalty, spectral norm, etc



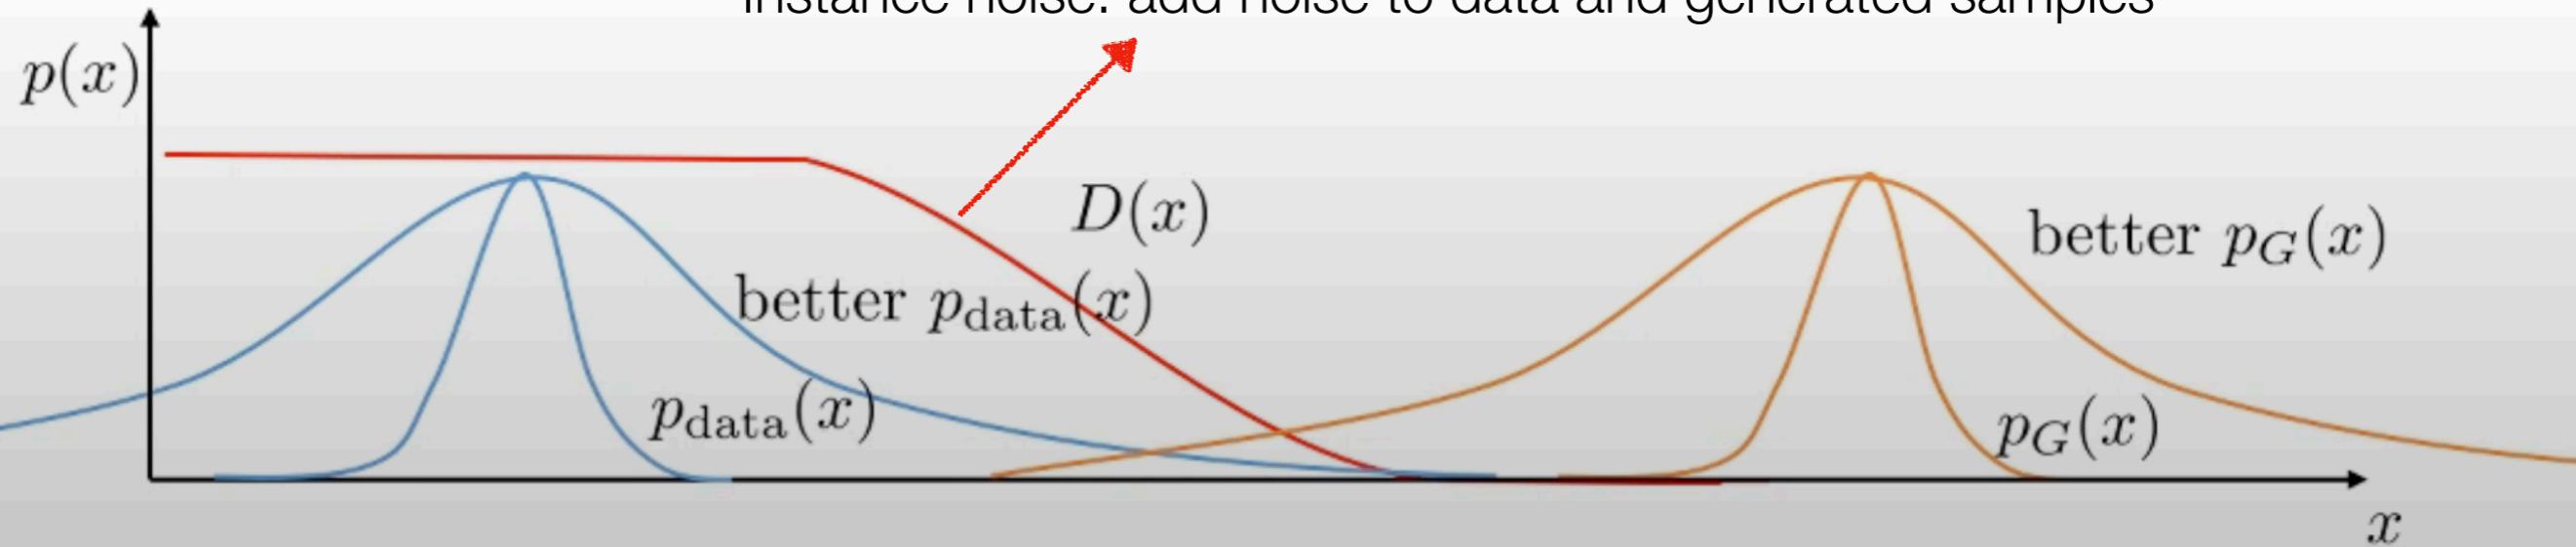
Why is training GANs hard?

Least-squares GAN (LSGAN): discriminator outputs real-valued number

Wasserstein GAN (WGAN): discriminator is constrained to be Lipschitz-continuous
Gradient Penalty, spectral norm, etc

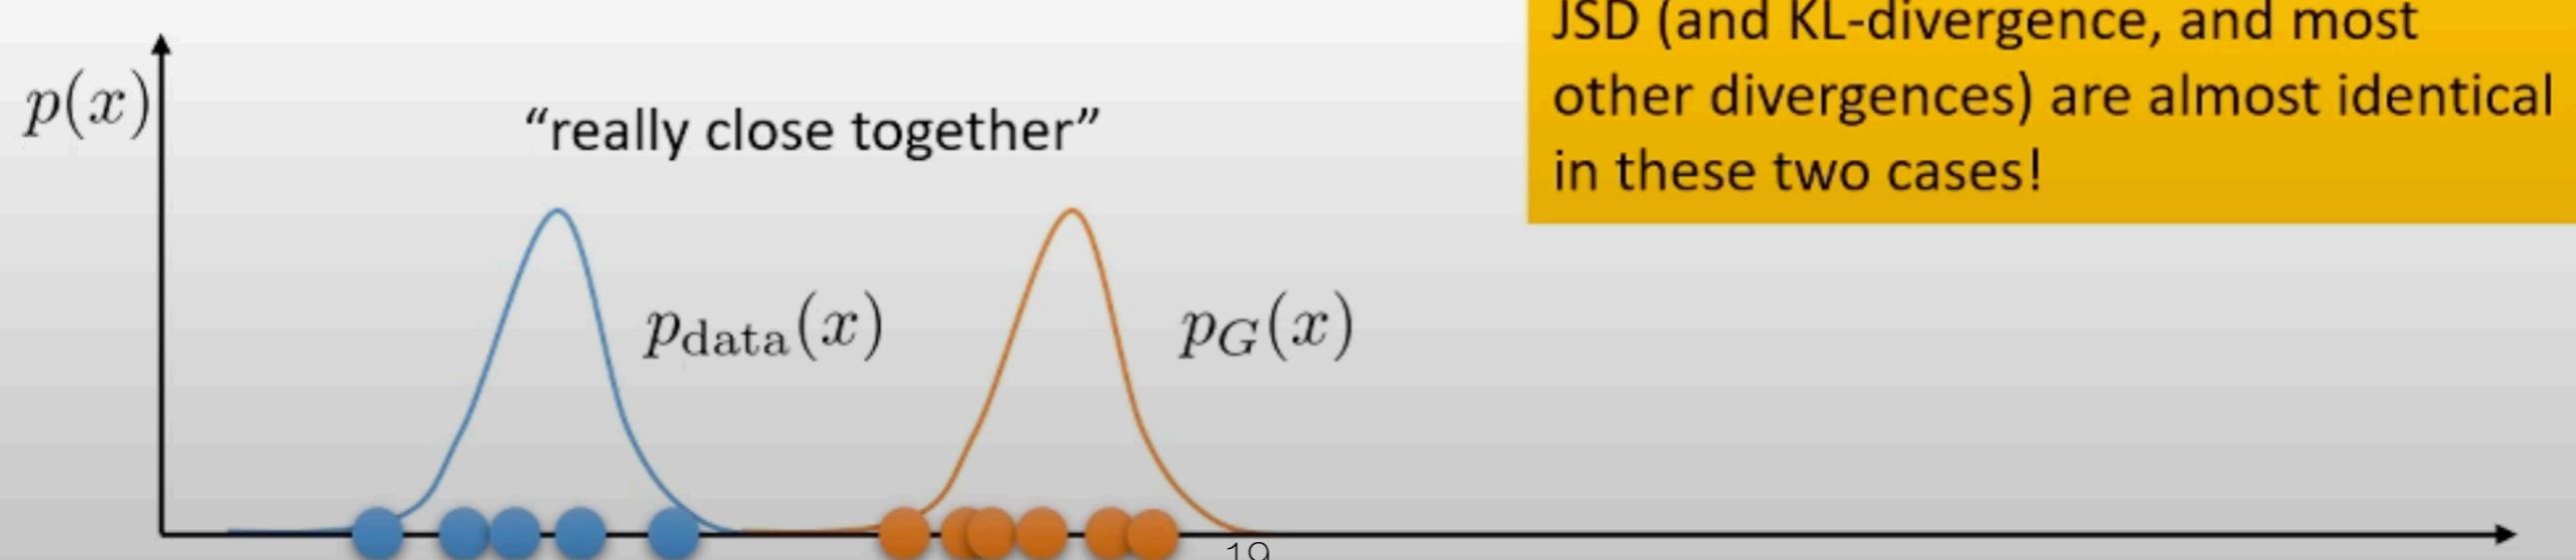
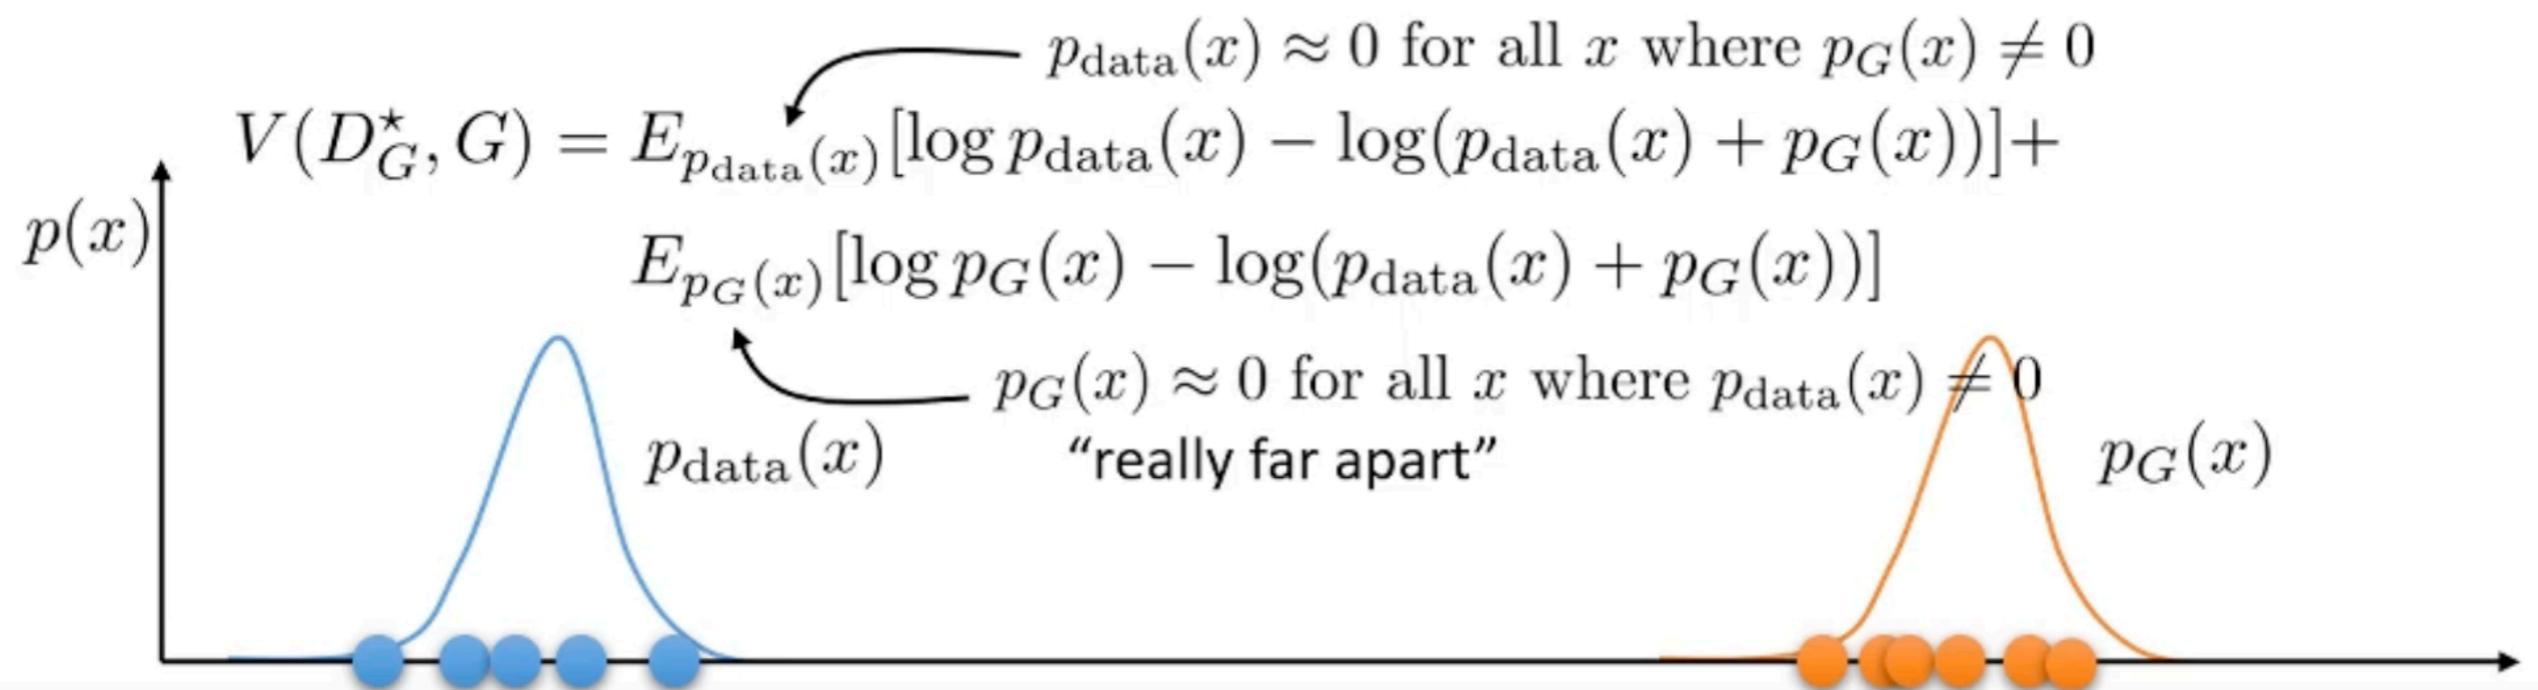


Instance noise: add noise to data and generated samples



Wasserstein GANs (Arjovsky 2017)

High level intuition: the JS divergence used in classic GAN doesn't account for “distance”



Wasserstein GANs (Arjovsky 2017)

- **Intuition:** Distance used to measure the cost of having to transport mass in Euclidean space
 - ▶ In other words, how far do I have to go to “transport” one distribution to another
- Improves the stability when training the model and provides a loss function that correlates with the quality of generated images.
- Wasserstein distance between the real data distribution P_r and the generator’s one P_θ

$$W(P_r, P_\theta) = \sup_{\|f\| \leq 1} \mathbb{E}_{x \sim P_r}[f(x)] - \mathbb{E}_{x \sim P_\theta}[f(x)]$$

where the supremum is taken over all **1-Lipschitz functions.**

1-Lipschitz function: $|f(x) - f(y)| \leq |x - y|$

- ▶ functions can’t be too sharp -> guarantees smoothness of f

Wasserstein GANs (Arjovsky 2017)

- **Intuition:** Distance used to measure the cost of having to transport mass in Euclidean space
 - ▶ In other words, how far do I have to go to “transport” one distribution to another
- Improves the stability when training the model and provides a loss function that correlates with the quality of generated images.
- Wasserstein distance between the real data distribution P_r and the generator’s one P_θ

$$W(P_r, P_\theta) = \sup_{\|f\| \leq 1} \mathbb{E}_{x \sim P_r}[f(x)] - \mathbb{E}_{x \sim P_\theta}[f(x)]$$

where the supremum is taken over all **1-Lipschitz functions.**

1-Lipschitz function: $|f(x) - f(y)| \leq |x - y|$

- ▶ functions can’t be too sharp -> guarantees smoothness of f
- In practice, smoothness is achieved by **weight clipping** ($w \in [-c, c]$)

Wasserstein GANs (Arjovsky 2017)

- **Intuition:** Distance used to measure the cost of having to transport mass in Euclidean space
 - ▶ In other words, how far do I have to go to “transport” one distribution to another
- Improves the stability when training the model and provides a loss function that correlates with the quality of generated images.
- Wasserstein distance between the real data distribution P_r and the generator’s one P_θ

$$W(P_r, P_\theta) = \sup_{\|f\| \leq 1} \mathbb{E}_{x \sim P_r}[f(x)] - \mathbb{E}_{x \sim P_\theta}[f(x)]$$

where the supremum is taken over all **1-Lipschitz functions.**

1-Lipschitz function: $|f(x) - f(y)| \leq |x - y|$

- ▶ functions can’t be too sharp -> guarantees smoothness of f
- In practice, smoothness is achieved by **weight clipping** ($w \in [-c, c]$)

Wasserstein GANs (Arjovsky 2017)

- **Intuition:** Distance used to measure the cost of having to transport mass in Euclidean space
 - ▶ In other words, how far do I have to go to “transport” one distribution to another
- Improves the stability when training the model and provides a loss function that correlates with the quality of generated images.
- Wasserstein distance between the real data distribution P_r and the generator’s one P_θ

$$W(P_r, P_\theta) = \sup_{\|f\| \leq 1} \mathbb{E}_{x \sim P_r}[f(x)] - \mathbb{E}_{x \sim P_\theta}[f(x)]$$

where the supremum is taken over all **1-Lipschitz functions.**

1-Lipschitz function: $|f(x) - f(y)| \leq |x - y|$

Not a probability!!!

- ▶ functions can’t be too sharp -> guarantees smoothness of f

- In practice, smoothness is achieved by **weight clipping** ($w \in [-c, c]$)

Wasserstein GANs (Arjovsky 2017)

$$W(p_{\text{data}}, p_G) = \sup_{\|f\|_L \leq 1} E_{p_{\text{data}}}[f(x)] - E_{p_G(x)}[f(x)]$$

1. update f_θ using gradient of $E_{x \sim p_{\text{data}}}[f_\theta(x)] - E_{z \sim p(z)}[f_\theta(G(z))]$
2. clip all weight matrices in θ to $[-c, c]$
3. update generator to *maximize* $E_{z \sim p(z)}[f_\theta(G(z))]$

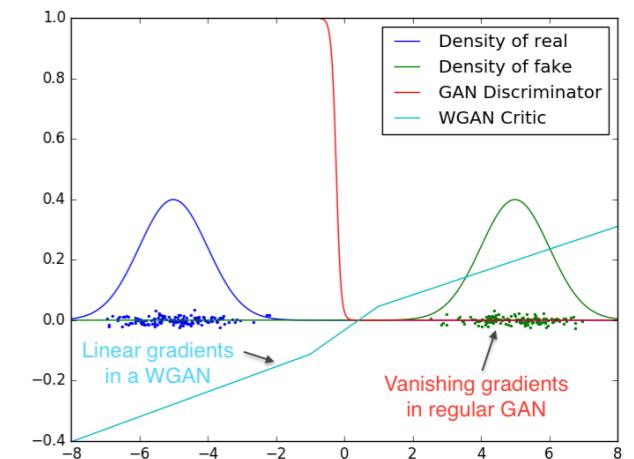
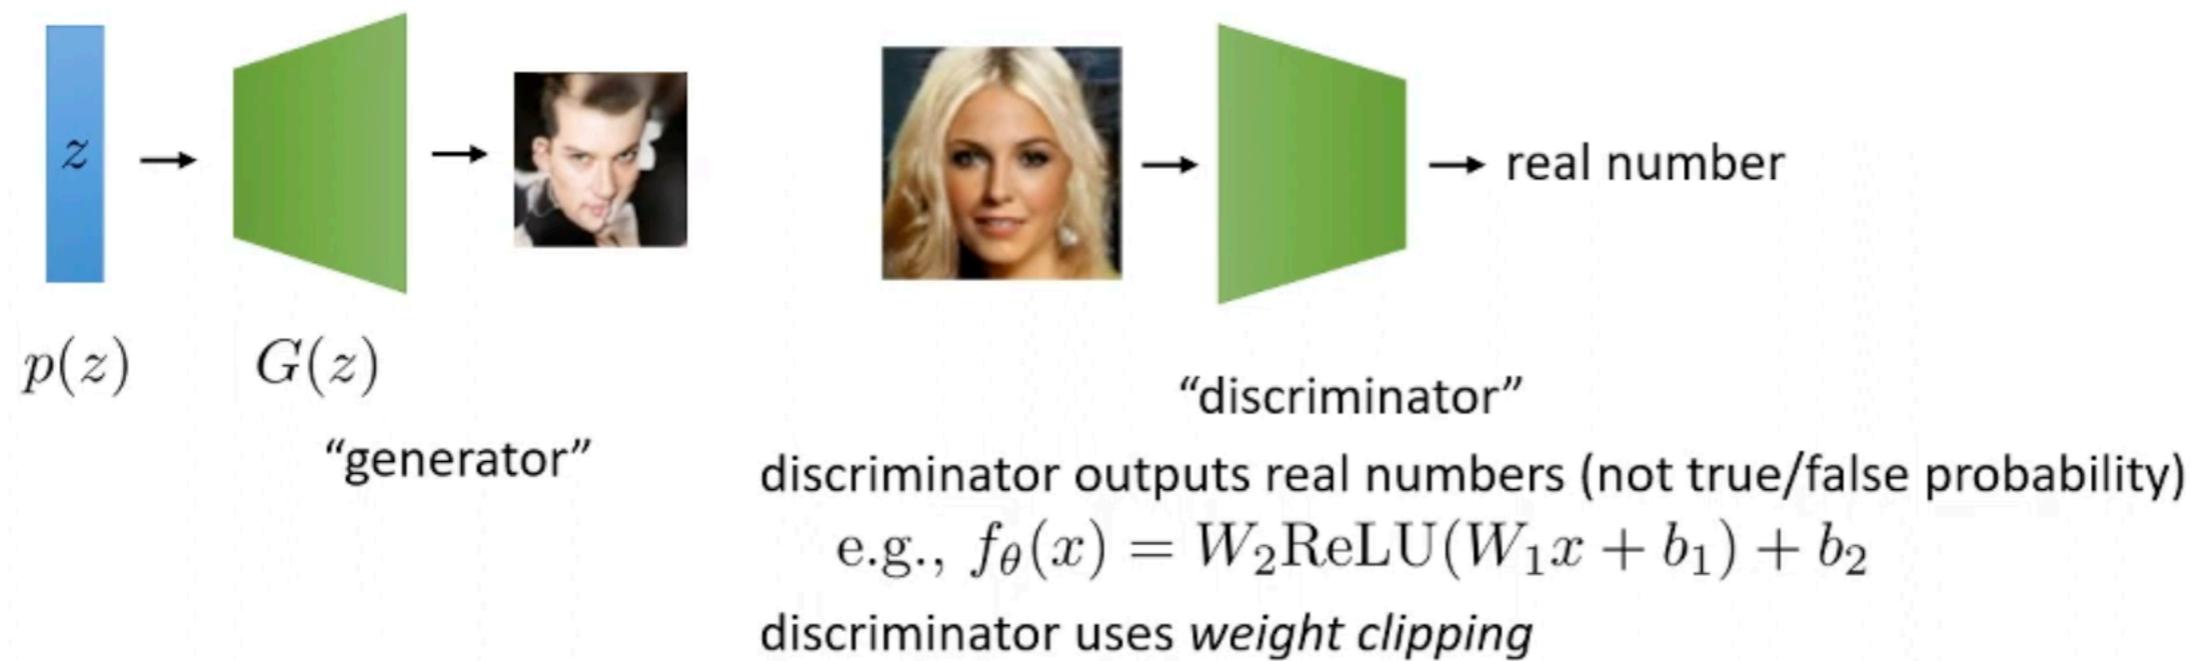


Figure 2: Optimal discriminator and critic when learning to differentiate two Gaussians. As we can see, the discriminator of a minimax GAN saturates and results in vanishing gradients. Our WGAN critic provides very clean gradients on all parts of the space.

Wasserstein GANs (Arjovsky 2017)

Discriminator = Critic (just a change of name in that work)

Algorithm 1 WGAN, our proposed algorithm. All experiments in the paper used the default values $\alpha = 0.00005$, $c = 0.01$, $m = 64$, $n_{\text{critic}} = 5$.

Require: : α , the learning rate. c , the clipping parameter. m , the batch size.
 n_{critic} , the number of iterations of the critic per generator iteration.

Require: : w_0 , initial critic parameters. θ_0 , initial generator's parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{\text{critic}}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w \left[ \frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)})) \right]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSProp}(\theta, g_\theta)$ 
12: end while
```

Improved Training of Wasserstein GANs

Ishaan Gulrajani^{1,*}, Faruk Ahmed¹, Martin Arjovsky², Vincent Dumoulin¹, Aaron Courville^{1,3}

¹ Montreal Institute for Learning Algorithms

² Courant Institute of Mathematical Sciences

³ CIFAR Fellow

igul222@gmail.com

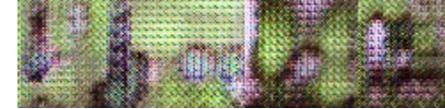
{faruk.ahmed,vincent.dumoulin,aaron.courville}@umontreal.ca

ma4371@nyu.edu

- Discriminator weight clipping can lead to undesired behavior
- A differentiable function is 1-Lipschitz if and only if it has **gradients with norm at most 1 everywhere**.
- Directly **constrain the gradient norm of the critic's output with respect to its input**.

$$L = \underbrace{\mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})] - \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})]}_{\text{Original critic loss}} + \underbrace{\lambda_{\hat{\mathbf{x}} \sim \mathbb{P}_{\hat{\mathbf{x}}}} \left[(\|\nabla_{\hat{\mathbf{x}}} D(\hat{\mathbf{x}})\|_2 - 1)^2 \right]}_{\text{Our gradient penalty}}$$

Improved Wasserstein GANs (Arjovsky 2017)

DCGAN	LSGAN	WGAN (clipping)	WGAN-GP (ours)	
Baseline (G : DCGAN, D : DCGAN)				
G : No BN and a constant number of filters, D : DCGAN				
G : 4-layer 512-dim ReLU MLP, D : DCGAN				
No normalization in either G or D				
Gated multiplicative nonlinearities everywhere in G and D				
tanh nonlinearities everywhere in G and D				
101-layer ResNet G and D				

Improved Techniques for Training GANs

Tim Salimans

tim@openai.com

Ian Goodfellow

ian@openai.com

Wojciech Zaremba

woj@openai.com

Vicki Cheung

vicki@openai.com

Alec Radford

alec.radford@gmail.com

Xi Chen

peter@openai.com

Feature matching addresses the instability of GANs

- The new objective requires the generator to **generate data that matches the statistics of the real data**
- We use the discriminator only to **specify the statistics that we think are worth matching**
- We train the generator to match the expected value of the **features on an intermediate layer of the discriminator**

$$\|\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} \mathbf{f}(\mathbf{x}) - \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}} \mathbf{f}(G(\mathbf{z}))\|_2^2$$

Improved Techniques for Training GANs

Tim Salimans

tim@openai.com

Ian Goodfellow

ian@openai.com

Wojciech Zaremba

woj@openai.com

Vicki Cheung

vicki@openai.com

Alec Radford

alec.radford@gmail.com

Xi Chen

peter@openai.com

Minibatch discrimination alleviates mode collapse

- The concept of minibatch discrimination is quite general: any discriminator model that **looks at multiple examples in combination**, rather than in isolation, could potentially help avoid collapse of the generator.
- Basic idea: the feature representation $\mathbf{f}(\mathbf{x}_i)$ in a certain layer of the discriminator for any point \mathbf{x}_i is augmented with a kernel-based distance w.r.t. other points in the mini-batch

$$\tilde{\mathbf{f}}(\mathbf{x}_1) = \left[\mathbf{f}(\mathbf{x}_1), k(\mathbf{f}(\mathbf{x}_1), \mathbf{f}(\mathbf{x}_2)), k(\mathbf{f}(\mathbf{x}_1), \mathbf{f}(\mathbf{x}_3)), \dots, k(\mathbf{f}(\mathbf{x}_1), \mathbf{f}(\mathbf{x}_M)) \right]$$

Other GAN variants

Style GAN (2018)

One problem with GANs: Latent space is entangled

- This means that the N latent dimensions representing an image contains partly several features

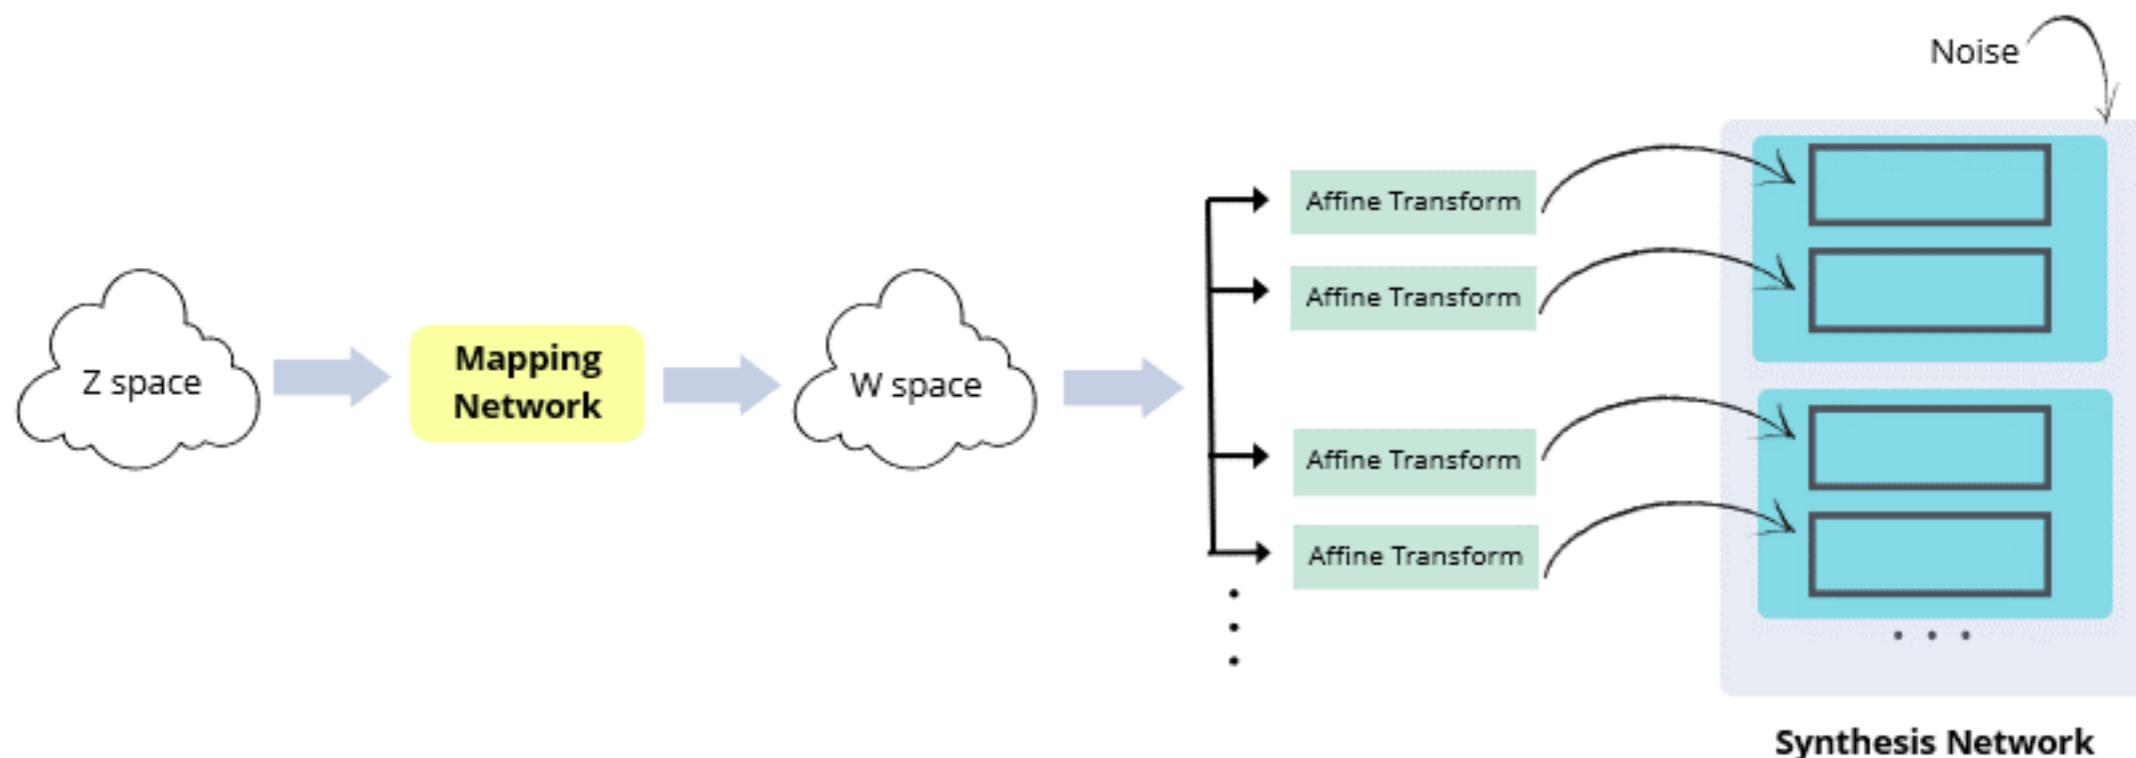
For example: What if I ask how would my cat look like with short hair?



Undesired changes due to entanglement: The short hair dimension could contain dog features as well

Style GAN (2018): key contribution

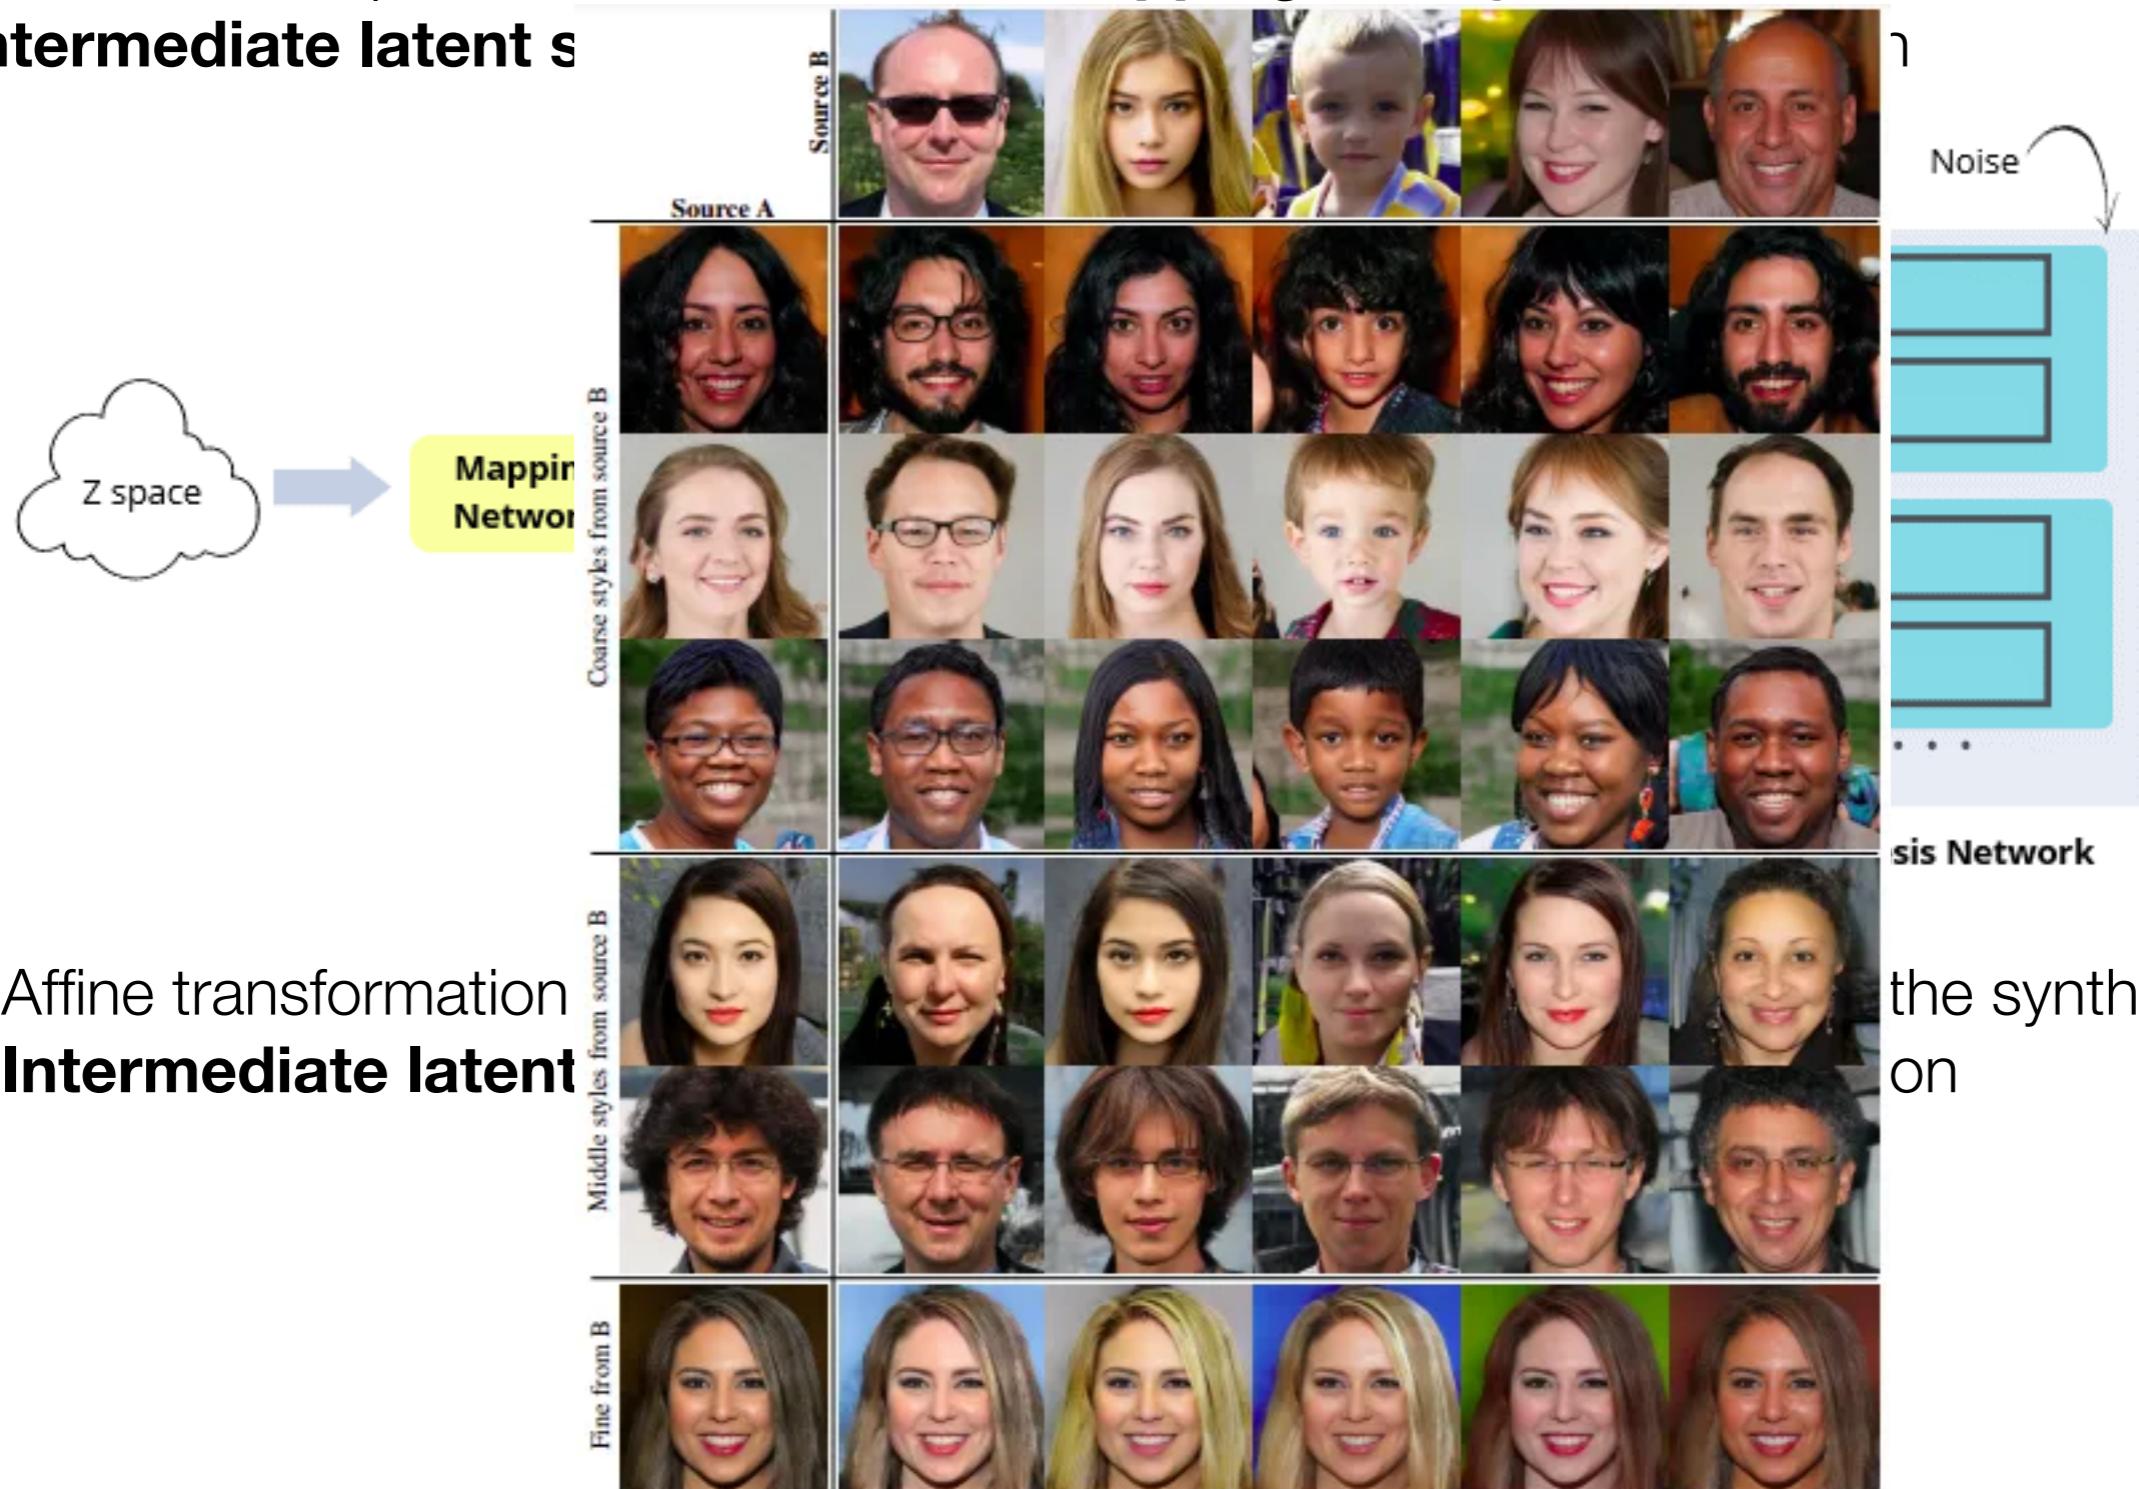
- Generator composed of two networks: **mapping and synthesis**
- **Intermediate latent space (W space)** alongside affine transformation



- Affine transformation: Turns **w** vectors into a “style” that is fed to the synthesis net
- **Intermediate latent space (W space)** alongside affine transformation

Style GAN (2018): key contribution

- Generator composed of two networks: **mapping and synthesis**
- **Intermediate latent s**



- Affine transformation
- **Intermediate latent**

the synthesis net
on

Cycle GAN (2017)

two (conditional) generators:

G : turn X into Y (e.g., horse into zebra)

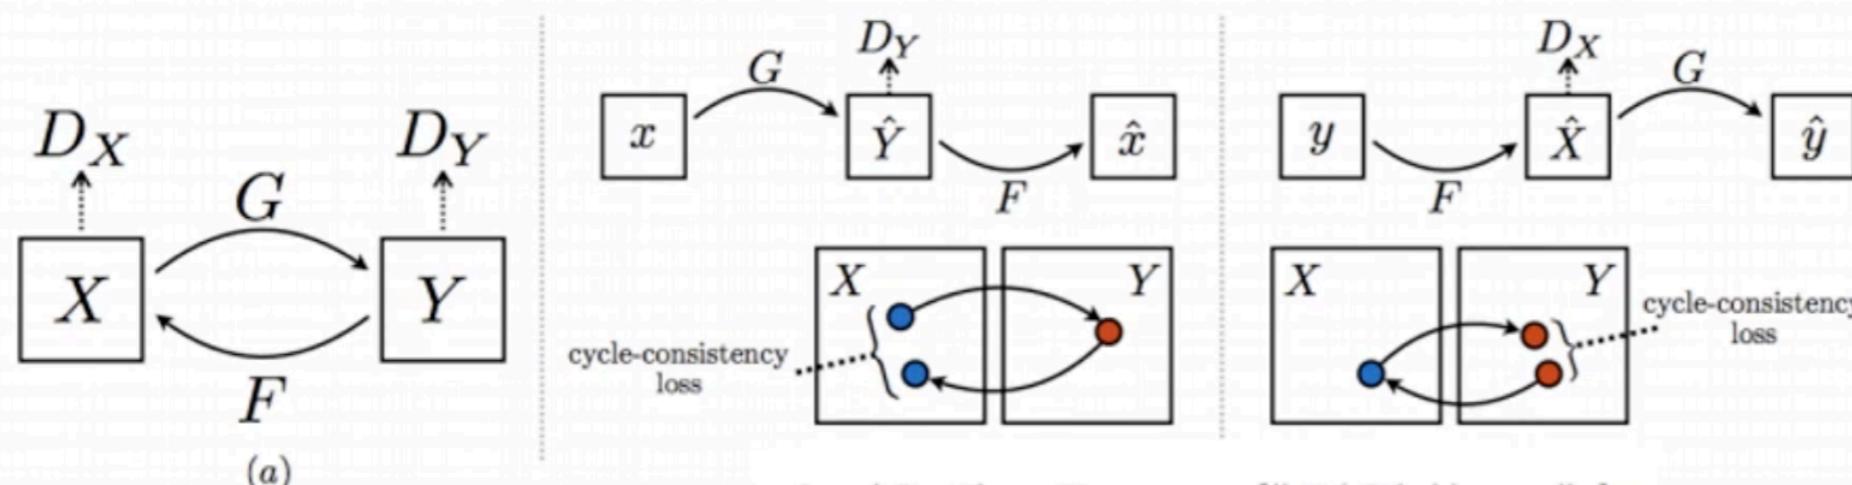
F : turn Y into X (e.g., zebra into horse)

two discriminators:

D_X : is it a realistic horse?

D_Y : is it a realistic zebra?

Problem: why should the “translated” zebra look anything like the original horse?



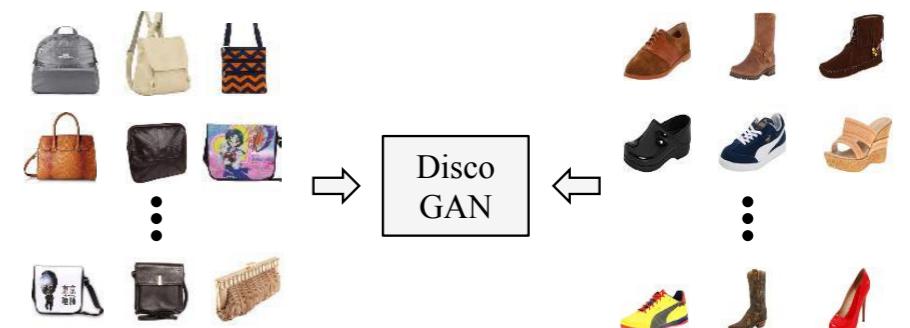
$$\begin{aligned}\mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1].\end{aligned}$$



Zhu, J. Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks.

Learning to Discover Cross-Domain Relations with Generative Adversarial Networks

Taeksoo Kim¹ Moonsu Cha¹ Hyunsoo Kim¹ Jung Kwon Lee¹ Jiwon Kim¹



(b) Handbag images (input) & **Generated** shoe images (output)



(c) Shoe images (input) & **Generated** handbag images (output)

Bidirectional GANs (2017)

ADVERSARIAL FEATURE LEARNING

Jeff Donahue

jdonahue@cs.berkeley.edu
Computer Science Division
University of California, Berkeley

Philipp Krähenbühl

philkr@utexas.edu
Department of Computer Science
University of Texas, Austin

Trevor Darrell

trevor@eecs.berkeley.edu
Computer Science Division
University of California, Berkeley

- Learn both the generator and the inverse mapping
- Discrimination not only in data space, but jointly in data and latent space

Bidirectional GANs (2017)

ADVERSARIAL FEATURE LEARNING

Jeff Donahue

jdonahue@cs.berkeley.edu
Computer Science Division
University of California, Berkeley

Philipp Krähenbühl

philkr@utexas.edu
Department of Computer Science
University of Texas, Austin

Trevor Darrell

trevor@eecs.berkeley.edu
Computer Science Division
University of California, Berkeley

- Learn both the generator and the inverse mapping
- Discrimination not only in data space, but jointly in data and latent space

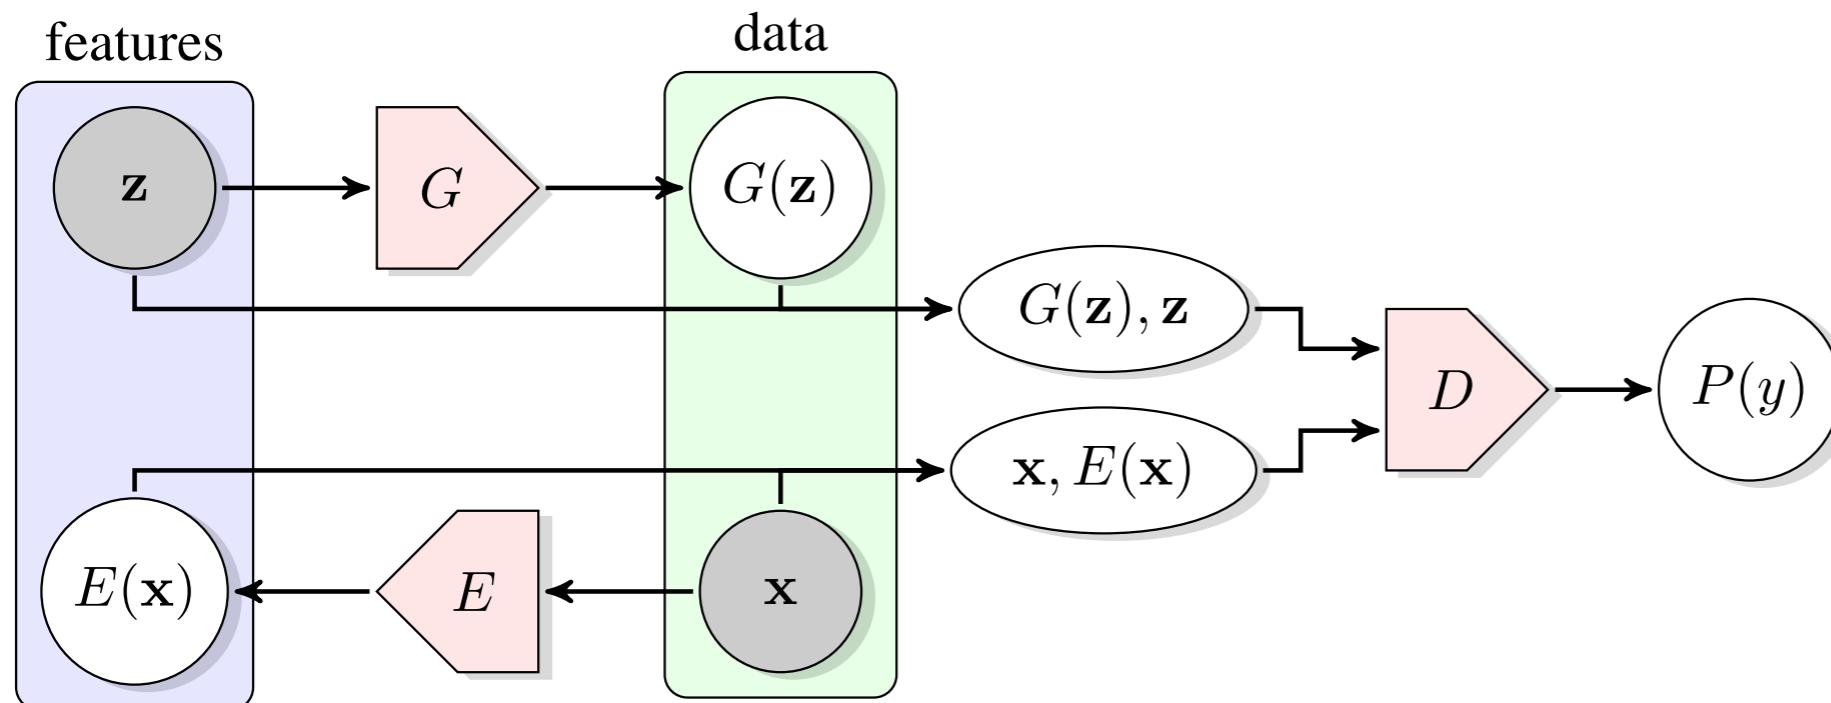


Figure 1: The structure of Bidirectional Generative Adversarial Networks (BiGAN).

Bidirectional GANs (2017)

- Learn both the generator and the inverse mapping
- Discrimination not only in data space, but jointly in data and latent space

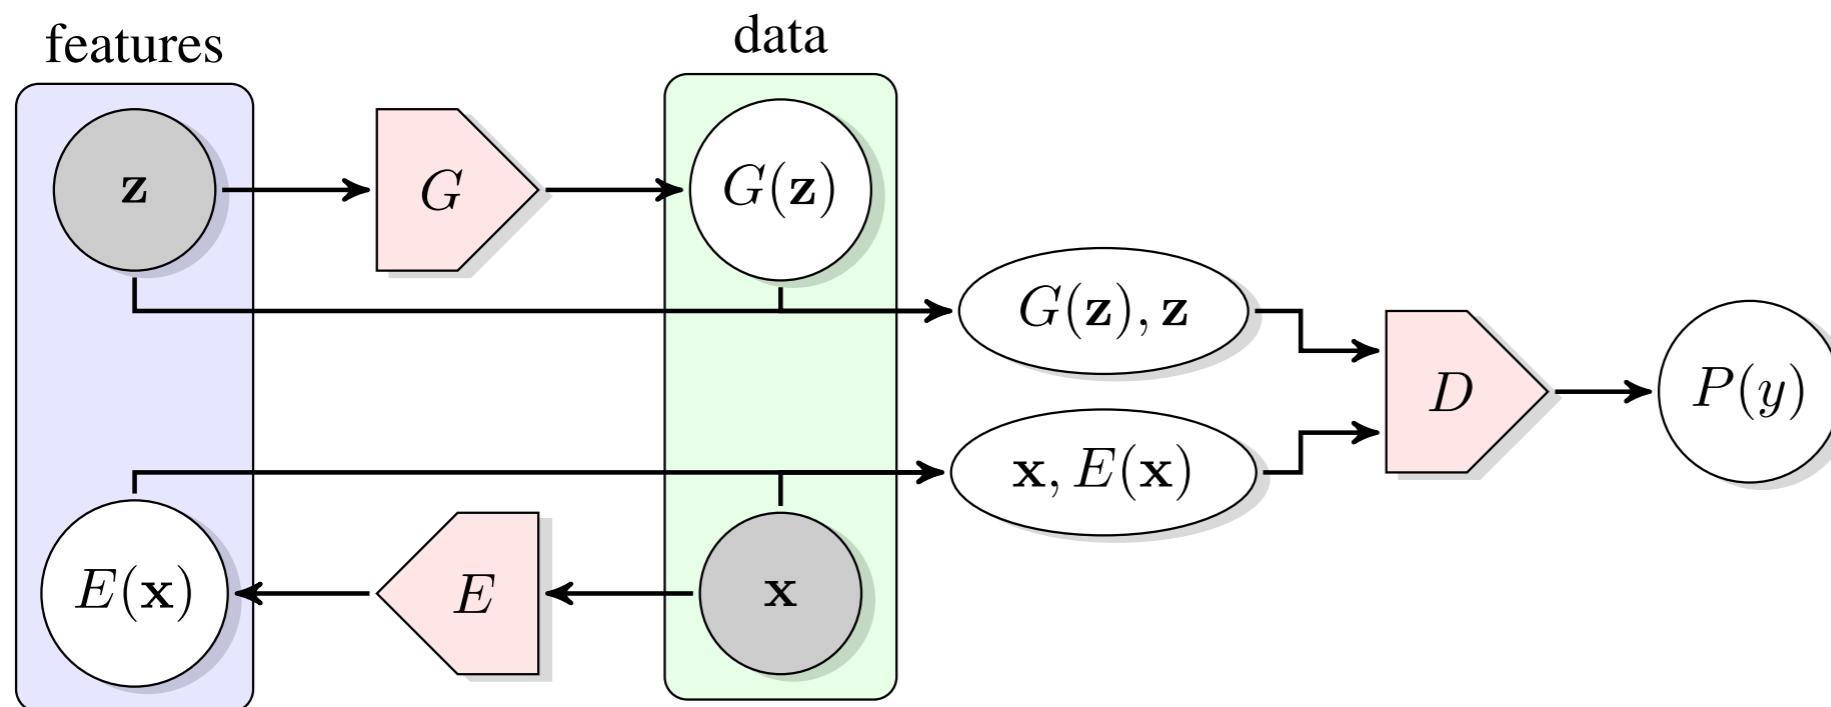


Figure 1: The structure of Bidirectional Generative Adversarial Networks (BiGAN).

Bidirectional GANs (2017)

$$\min_{G,E} \max_D V(D, E, G)$$

where

$$V(D, E, G) := \underbrace{\mathbb{E}_{\mathbf{x} \sim p_{\mathbf{X}}} [\mathbb{E}_{\mathbf{z} \sim p_E(\cdot | \mathbf{x})} [\log D(\mathbf{x}, \mathbf{z})]]}_{\log D(\mathbf{x}, E(\mathbf{x}))} + \underbrace{\mathbb{E}_{\mathbf{z} \sim p_{\mathbf{Z}}} [\mathbb{E}_{\mathbf{x} \sim p_G(\cdot | \mathbf{z})} [\log(1 - D(\mathbf{x}, \mathbf{z}))]]}_{\log(1 - D(G(\mathbf{z}), \mathbf{z}))}$$

- Learn both the generator and the inverse mapping
- Discrimination not only in data space, but jointly in data and latent space

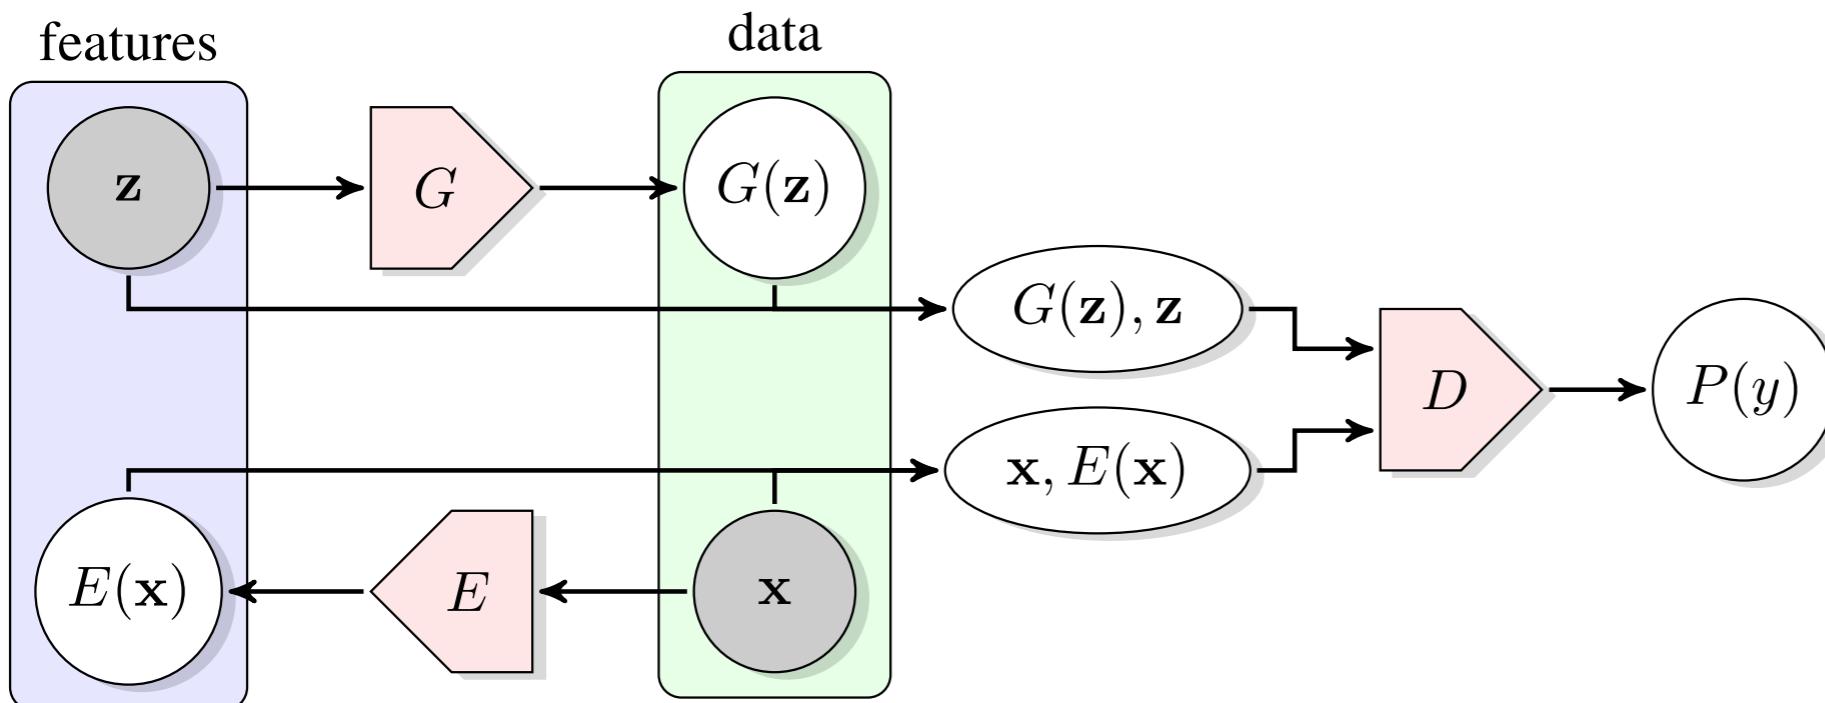


Figure 1: The structure of Bidirectional Generative Adversarial Networks (BiGAN).

UNSUPERVISED CROSS-DOMAIN IMAGE GENERATION

Yaniv Taigman, Adam Polyak & Lior Wolf

Facebook AI Research

Tel-Aviv, Israel

{yaniv, adampolyak, wolf}@fb.com

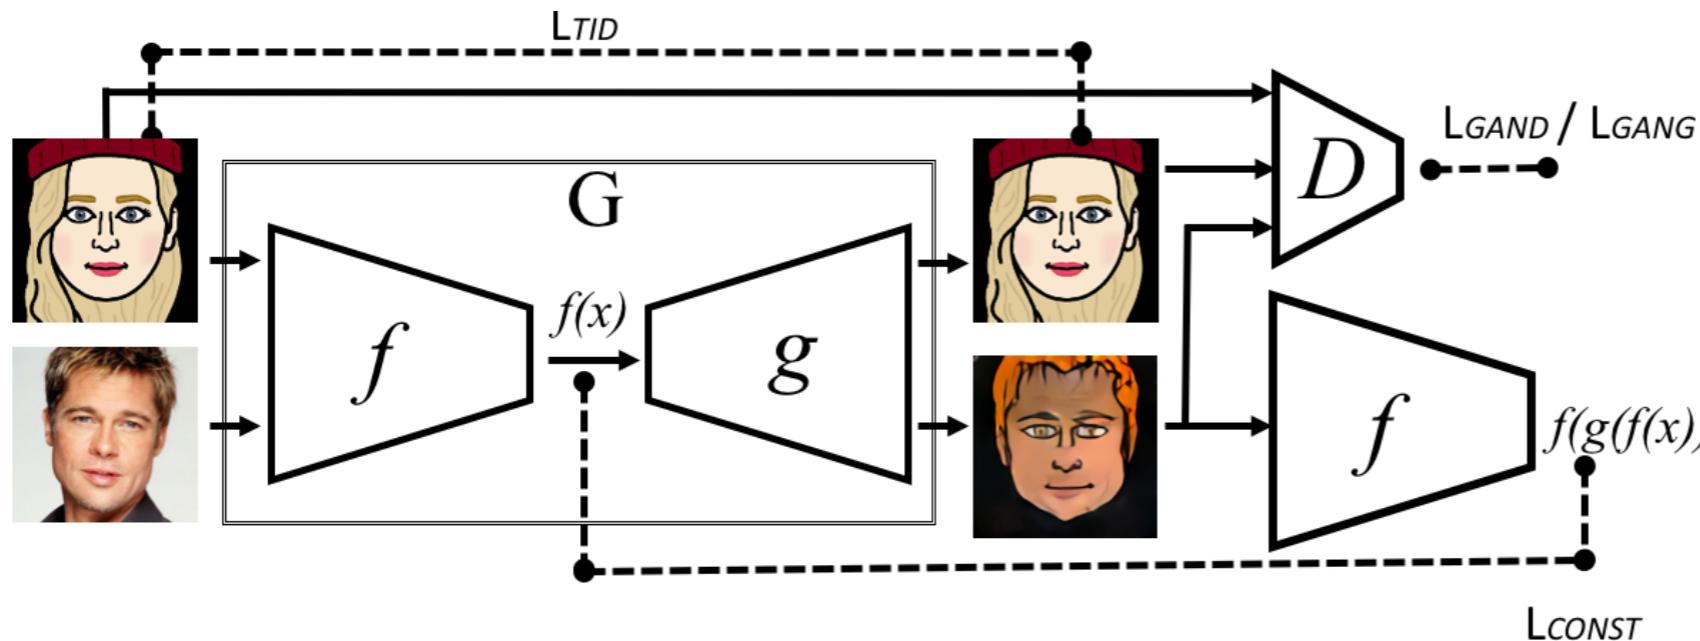


Figure 1: The Domain Transfer Network. Losses are drawn with dashed lines, input/output with solid lines. After training, the forward model G is used for the sample transfer.

UNSUPERVISED CROSS-DOMAIN IMAGE GENERATION

Yaniv Taigman, Adam Polyak & Lior Wolf

Facebook AI Research

Tel-Aviv, Israel

{yaniv, adampolyak, wolf}@fb.com

